




Practical Concurrent and Parallel Programming III

Shared Memory II

Raúl Pardo

Assignment workload

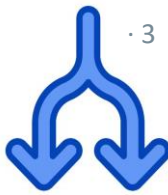


- We would like to get an estimation on the amount of hours you spend on assignments
- Please go to the following mentimeter poll
<https://www.menti.com/aln6fdbutpc>

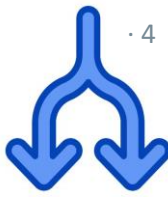
You should indicate the amount of hours that you spent to complete Assignment 1

That is, the amount of hours that you spent on PCPP exercises in the last two weeks combined





- Readers and Writers Problem
- Monitors
- Fairness
- Java Intrinsic Locks (**synchronized**)
- Hardware and Programming Language Concurrency Issues
 - Visibility
 - Reordering
- Volatile variables (**volatile**)



- Definitions of thread-safety
 - Classes
 - Programs
- Safe publication
- Immutability
- Instance confinement
- Synchronization primitives (synchronizers)
 - Semaphores
 - Barriers
- Producer-consumer problem



*A (concurrent) program is correct if and only if
it satisfies its specification*



- A *specification* (or *spec*) is a rigorous statement that describes the expected/desired behaviour of a program
- Examples
 - Many readers can access the shared resource at the same time, but only one can write
 - The output of the program must be `counter*num_threads`
- Specifications can be as precise as formulae in some logic (propositional, temporal, first-order, etc.)
 - We will not cover these details in the course



- A *specification* (or *spec*) is a rigorous statement that describes the expected/desired behaviour
- Examples
 - Many readers can access the shared resource at the same time, but only one can write
 - The output of the program must be `counter*num_threads`
- Specifications can be as precise as formulae in some logic (propositional, temporal, first-order, etc.)
 - We will not cover these details in the course

Is this specification for the readers-writers problem precise?

Reasoning about concurrent programs




- We have already covered the basic concepts to reason about the *correctness* of concurrent programs
- Reasoning about correctness of concurrent programs is tricky
 - You have experienced this already in the assignments where you work with programs consisting in a few lines of code
- Imagine having to reason about applications with hundreds of lines of code and many classes
 - Server applications
 - Operating Systems
 - GUIs
 - ...

Reasoning about concurrent programs



- We have already covered the basic concepts to reason about the *correctness* of concurrent programs
- Reasoning about correctness of concurrent programs is tricky
 - You have experienced this already in the assignments where you work with programs consisting in a few lines of code
- Imagine having to reason about applications with hundreds of lines of code and many classes
 - Server applications
 - Operating Systems
 - GUIs
 - ...



Modular class-based reasoning



- It is more manageable to separately analyse parts of the code and then combine them in safe ways
- In Object Oriented languages (such as Java) we can focus on analysing thread-safety for classes
- This reduces the analysis to concurrent method calls and field accesses

- *A **data race** occurs when two concurrent threads:*
 - *Access a shared memory location*
 - *At least one access is a write*
 - *There is no happens-before relation between the accesses*



New!

- A ***data race*** occurs when two concurrent threads:
 - Access a shared memory location
 - At least one access is a write
 - There is no happens-before relation between the accesses



New!

Inspired by the Java memory model ([JLS](#)): “A program is correctly synchronized if and only if all sequentially consistent executions are free of data races.”

*A class is said to be thread-safe if and only if
no concurrent execution of
method calls or field accesses (read/write)
result in data races on the fields of the class*

PCPP teaching team

Inspired by the Java memory model ([JLS](#)): “A program is correctly synchronized if and only if all sequentially consistent executions are free of data races.”

A class is said to be thread-safe if and only if no concurrent execution of method calls or field accesses (read/write) result in data races on the fields of the class

Note that this definition is independent of class invariants as opposed to Goetz Chapter 4. This definition is more similar to Goetz Chapter 2, page 18.

PCPP teaching team

Thread-safe class

Inspired by the Java memory model ([JLS](#)): “A program is correctly synchronized if and only if all sequentially consistent executions are free of data races.”

IMPORTANT: In this course, *thread-safety* is not an umbrella term for code that seem to behave correctly in concurrent environments.



A class is said to be thread-safe if and only if no concurrent execution of method calls or field accesses (read/write) result in data races on the fields of the class

Note that this definition is independent of class invariants as opposed to Goetz Chapter 4. This definition is more similar to Goetz Chapter 2, page 18.

PCPP teaching team

Thread-safe class

Inspired by the Java memory model ([JLS](#)): “A program is correctly synchronized if and only if all sequentially consistent executions are free of data races.”

IMPORTANT: In this course, *thread-safety* is not an umbrella term for code that seem to behave correctly in concurrent environments.



· 11

What is the specification in this definition?

A class is said to be thread-safe if and only if no concurrent execution of method calls or field accesses (read/write) result in data races on the fields of the class

Note that this definition is independent of class invariants as opposed to Goetz Chapter 4. This definition is more similar to Goetz Chapter 2, page 18.

PCPP teaching team

Thread-safe program



Do not confuse thread-safe classes with thread-safe programs.
Thread-safe programs are not defined in Goetz.

*A concurrent program is said to be thread-safe
if and only if it is race condition free*

Inspired by the Java memory model *correctly synchronized program* (see previous slide), but we impose a different condition by requiring freedom of race conditions

PCPP teaching team

It is very important to note that:

For any program p ,

p only accesses thread-safe classes

\Rightarrow

p is a thread-safe program

It is very important to note that:

For any program p ,

p only accesses thread-safe classes

\Rightarrow

p is a thread-safe program

Programs using thread-safe classes
may contain race conditions.

It is very important

For any program

Does this hold?

p is a thread-safe program

\Rightarrow

p only accesses thread-safe classes

p only accesses thread-safe classes

\Rightarrow

p is a thread-safe program

Programs using thread-safe classes
may contain race conditions.

- To analyse whether a class is thread-safe, we must simply ensure that there is a happens-before relation for any concurrent execution of field access and method calls where at least one of them results in a write access
- In what follows, we list the elements to identify/consider:
 - Class state
 - Escaping
 - (Safe) publication
 - Immutability
 - Mutual exclusion

- To analyse whether a class is thread-safe, we must simply ensure that there is a happens-before relation for any concurrent execution of field access and method calls where at least one of them results in a write access
- In what follows, we list the elements to identify/consider:
 - Class state
 - Escaping
 - (Safe) publication
 - Immutability
 - Mutual exclusion

When asked to reason about the thread-safety of a class, you must always cover these elements



- By definition, (uncontrolled) concurrent access to the shared state (variables) leads to data races
- So, the first thing we need to do is to identify the fields that may be shared by several threads
- The state of a class involves the fields defined in the class
 - In a nutshell, our goal is to ensure that concurrent access to class state is free from data races

```
class C {  
    // class state (variables)  
    T s1;  
    T s2;  
    T s3;  
    T s4;  
    ...  
  
    // class methods  
    T m1 (...) {...}  
    T m2 (...) {...}  
    T m3 (...) {...}  
    ...  
}
```

If a class has no state (variables),
is it thread-safe?



- By definition, (uncontrolled) concurrent access to the shared state (variables) leads to data races
- So, the first thing we need to do is to identify the fields that may be shared by several threads
- The state of a class involves the fields defined in the class
 - In a nutshell, our goal is to ensure that concurrent access to class state is free from data races

```
class C {  
    // class state (variables)  
    T s1;  
    T s2;  
    T s3;  
    T s4;  
    ...  
  
    // class methods  
    T m1 (...) {...}  
    T m2 (...) {...}  
    T m3 (...) {...}  
    ...  
}
```


Only class state (only recommended)



- Methods should only manipulate class state or parameters
 - For instance, avoid the use of variables from parent classes

```
class C {  
    // class state (variables)  
    int i = 0;  
  
    // class methods  
  
    // `x` is not part of the class state  
    // we cannot ensure happens-before  
    public void m(){x = 0;}  
  
    public void n(List<Ingeter> l){l.add(1);}  
}
```

Only class state (only recommended)



- Methods should only manipulate class state or parameters
 - For instance, avoid the use of variables from parent classes
- Methods should avoid using object references as parameters
 - We cannot guarantee happens-before relations with the referenced object

```
class C {  
    // class state (variables)  
    int i = 0;  
  
    // class methods  
  
    // `x` is not part of the class state  
    // we cannot ensure happens-before  
    public void m(){x = 0;}  
  
    public void n(List<Integer> l){l.add(1);}  
}
```

```
// program using Counter  
  
List<Integer> l = new ArrayList<Integer>();  
List<Integer> C = new C();  
new Thread(() -> {  
    l.add(42);  
}).start();  
  
new Thread(() -> {  
    c.m(l); // no happens-before relation  
           // with the thread above  
}).start();
```

Only class state (only recommended)



- Methods should only manipulate class state or parameters
 - For instance, avoid the use of variables from parent classes
- Methods should avoid using object references as parameters
 - We cannot guarantee happens-before relations with the referenced object
- That said, our definition of thread-safe class focuses on data races on the fields of the class
 - Therefore, these problems do not violate the definition

```
class C {  
    // class state (variables)  
    int i = 0;  
  
    // class methods  
  
    // `x` is not part of the class state  
    // we cannot ensure happens-before  
    public void m(){x = 0;}  
  
    public void n(List<Integer> l){l.add(1);}  
}
```

```
// program using Counter  
  
List<Integer> l = new ArrayList<Integer>();  
List<Integer> C = new C();  
new Thread(() -> {  
    l.add(42);  
}).start();  
  
new Thread(() -> {  
    c.m(1); // no happens-before relation  
           // with the thread above  
}).start();
```



```
class Counter {  
    // class state (variables)  
    int i=0;  
  
    // class methods  
    public synchronized void inc(){i++;}  
}
```

Is the class `Counter` thread-safe?



```
class Counter {  
    // class state (variables)  
    int i=0;  
  
    // class methods  
    public synchronized void inc(){i++;}  
}
```

Is the class `Counter` thread-safe?



```
class Counter {  
    // class state (variables)  
    int i=0;  
  
    // class methods  
    public synchronized void inc(){i++;}  
}
```

```
// program using Counter  
  
Counter c = new Counter();  
new Thread(() -> {  
    c.inc();  
}).start();  
  
new Thread(() -> {  
    c.i++; // escaped the lock in inc()  
}).start();
```



- It is important to not expose shared state variables
- Otherwise, threads may use them without ensuring mutual exclusion
 - Thus, we cannot enforce a happens-before relation

```
class Counter {  
    // class state (variables)  
    int i=0;  
  
    // class methods  
    public synchronized void inc(){i++;}  
}
```

```
// program using Counter  
  
Counter c = new Counter();  
new Thread(() -> {  
    c.inc();  
}).start();  
  
new Thread(() -> {  
    c.i++; // escaped the lock in inc()  
}).start();
```

- It is important to not expose shared state variables
- Otherwise, threads may use them without ensuring mutual exclusion
 - Thus, we cannot enforce a happens-before relation
- Defining all (shared) class state (primitive) variables as private ensures that these variables will only be accessed through public methods.
 - Thus, it is easier to control and reason about concurrent access

```
class Counter {  
    // class state (variables)  
    int i=0;  
  
    // class methods  
    public synchronized void inc(){i++;}  
}
```

```
// program using Counter  
  
Counter c = new Counter();  
new Thread(() -> {  
    c.inc();  
}).start();  
  
new Thread(() -> {  
    c.i++; // escaped the lock in inc()  
}).start();
```




```
class IntArrayList {  
    // class state  
    private List<Integer> a = new ArrayList<Integer>();  
  
    public synchronized void set(Integer index, Integer elem)  
    { a.set(index,elem); }  
  
    public synchronized List<Integer> get() { return a; }  
}
```

Is the class `IntArrayList` thread-safe?



```
class IntArrayList {  
    // class state  
    private List<Integer> a = new ArrayList<Integer>();  
  
    public synchronized void set(Integer index, Integer elem)  
    { a.set(index,elem); }  
  
    public synchronized List<Integer> get() { return a; }  
}
```

Is the class `IntArrayList` thread-safe?



```
class IntArrayList {  
    // class state  
    private List<Integer> a = new ArrayList<Integer>();  
  
    public synchronized void set(Integer index, Integer elem)  
    { a.set(index,elem); }  
  
    public synchronized List<Integer> get() { return a; }  
}
```

```
IntArrayList array = new IntArrayList();  
new Thread() -> {  
    array.set(0,1); // access state with lock  
}).start();  
new Thread() -> {  
    array.get().set(0,42); // access state without locks  
}).start();
```

- Remember that when a method returns an object, we get a *reference* to that object
- Therefore, even if obtain the reference using locks, later we can modify the content of the object without locks

```
class IntArrayList {  
    // class state  
    private List<Integer> a = new ArrayList<Integer>();  
  
    public synchronized void set(Integer index, Integer elem)  
    { a.set(index,elem); }  
  
    public synchronized List<Integer> get() { return a; }  
}
```

```
IntArrayList array = new IntArrayList();  
new Thread(() -> {  
    array.set(0,1); // access state with lock  
}).start();  
new Thread(() -> {  
    array.get().set(0,42); // access state without locks  
}).start();
```



- It is important to ensure that initialization *happens-before* publication
 - That is, before making accessible a reference to an object, all its fields must be correctly initialized



- It is important to ensure that initialization *happens-before* publication
 - That is, before making accessible a reference to an object, all its fields must be correctly initialized

```
public class UnsafeLazyInitialization {  
    private static Resource resource;  
  
    public static Resource getInstance() {  
        if (resource == null)  
            resource = new Resource();  
        return resource;  
    }  
}
```



- It is important to ensure that initialization *happens-before* publication
 - That is, before making accessible a reference to an object, all its fields must be correctly initialized

```
public class UnsafeLazyInitialization {  
    private static Resource resource;  
  
    public static Resource getInstance() {  
        if (resource == null)  
            resource = new Resource();  
        return resource;  
    }  
}
```

Is this class thread-safe?

- Visibility issues may appear during initialization of objects

```
public class UnsafeInitialization {  
    private int x;  
    private Object o;  
    public UnsafeInitialization() {  
        x = 42;  
        o = new Object();  
    }  
}
```


Object initialization & visibility



- Visibility issues may appear during initialization of objects



```
public class UnsafeInitialization {  
    private int x;  
    private Object o;  
    public UnsafeInitialization() {  
        x = 42;  
        o = new Object();  
    }  
}
```



- Visibility issues may appear during initialization of objects



```
public class UnsafeInitialization {  
    private int x;  
    private Object o;  
    public UnsafeInitialization() {  
        x = 42;  
        o = new Object();  
    }  
}
```

- For the thread executing the constructor, there are no visibility issues, but if a reference to an instance of UnsafeInitialization object is accessible to another thread, it might not see **x==42** or **o** completely initialized

- We can address visibility issues during initialization as follows

```
public class UnsafeInitialization {  
    private volatile int x;  
    private final Object o;  
    public UnsafeInitialization() {  
        x = 42;  
        o = new Object();  
    }  
}
```

- We can address visibility issues during initialization as follows

For primitive types, we can:

- Declare them as **volatile**
- Declare them as **final** (only works if the content is never modified)
- Initialize as the default value: 0. (only works if the default value is acceptable)
- Use corresponding atomic class from Java standard library: **AtomicInteger**

```
public class UnsafeInitialization {  
    private volatile int x;  
    private final Object o;  
    public UnsafeInitialization() {  
        x = 42;  
        o = new Object();  
    }  
}
```



- We can address visibility issues during initialization as follows

For primitive types, we can:

- Declare them as **volatile**
- Declare them as **final** (only works if the content is never modified)
- Initialize as the default value: 0. (only works if the default value is acceptable)
- Use corresponding atomic class from Java standard library: **AtomicInteger**

```
public class UnsafeInitialization {  
    private volatile int x;  
    private final Object o;  
    public UnsafeInitialization()  
        x = 42;  
        o = new Object();  
    }  
}
```

For complex objects, we can:

- Declare them as **final**
- Initialize as the default value: null. (only works if the default value is acceptable)
- Use the **AtomicReference** class

Object initialization & visibility



- We can address visibility issues during initialization as follows

For primitive types, we can:

- Declare them as **volatile**
- Declare them as **final** (only works if the content is never modified)
- Initialize as the default value: 0. (only works if the default value is acceptable)
- Use corresponding atomic class from Java standard library: **AtomicInteger**

```
public class UnsafeInitialization {  
    private volatile int x;  
    private final Object o;  
    public UnsafeInitialization()  
        x = 42;  
        o = new Object();  
    }  
}
```

For complex objects, we can:

- Declare them as **final**
- Initialize as the default value: null. (only works if the default value is acceptable)
- Use the **AtomicReference** class

Why do these solutions solve visibility issues?



- The previous suggestions ensure safe publication because:
 - They established a *happens-before* relation between initialization and access the object's reference (publication)
 - *A write to a volatile field happens-before every subsequent read of that field.*
 - *The default initialization (zero, false, or null) of any object happens-before any other actions of a program.*
 - *The initialization of a final field happens-before any other actions of a program (after the constructor has finished its execution)*
 - At the JVM level, the reason is that
 - **final** fields cannot be cached or reordered during initialization
 - All fields are initialized with default values during class loading
 - writes on **volatile** are flushed to main memory and reordered (during initialization)



- The previous suggestions ensure safe publication because:
 - They established a *happens-before* relation between initialization and access the object's reference (publication)
 - *A write to a volatile field happens-before every subsequent read of that field.*
 - *The default initialization (zero, false, or null) of any object happens-before any other actions of a program.*
 - *The initialization of a final field happens-before any other actions of a program (after the constructor has finished its execution)*
 - At the JVM level, the reason is that
 - **final** fields cannot be cached or reordered during initialization
 - All fields are initialized with default values during class loading
 - writes on **volatile** are flushed to main memory and reordered (during initialization)

Defined by us from
the JLS explanation.
You can use for
exercises in this
course.

- The previous suggestions ensure safe publication because:
 - They established a *happens-before* relation between initialization and access the object's reference (publication)
 - *A write to a volatile field happens-before every subsequent read of that field.*
 - *The default initialization (zero, false, or null) of any object happens-before any other actions of a program.*
 - *The initialization of a final field happens-before any other actions of a program (after the constructor has finished its execution)*
 - At the JVM level, the reason is that
 - **final** fields cannot be cached or reordered during initialization
 - All fields are initialized with default values during class loading
 - writes on **volatile** are flushed to main memory and reordered (during initialization)

Defined by us from the JLS explanation. You can use for exercises in this course.

If the constructor of the class leaks a reference of the object being constructed before it has completed its execution, then there is no happens-before relation with the accesses to final field

Object initialization & visibility

NOTE: For clarity and simplicity, up to now, we did not take initialization concerns into account. But from now on we will.

29

- The previous suggestions ensure safe publication because:
 - They established a *happens-before* relation between initialization and access the object's reference (publication)
 - *A write to a volatile field happens-before every subsequent read of that field.*
 - *The default initialization (zero, false, or null) of any object happens-before any other actions of a program.*
 - *The initialization of a final field happens-before any other actions of a program (after the constructor has finished its execution)*
 - At the JVM level, the reason is that
 - **final** fields cannot be cached or reordered during initialization
 - All fields are initialized with default values during class loading
 - writes on **volatile** are flushed to main memory and reordered (during initialization)

Defined by us from the JLS explanation. You can use for exercises in this course.

If the constructor of the class leaks a reference of the object being constructed before it has completed its execution, then there is no happens-before relation with the accesses to final field

- An immutable object is one whose state cannot be changed after initialization
 - You can think of it as a constant
 - The **final** keyword in Java prevents modification of fields
 - Remember that variables assigned to an object only hold a reference to the object
- Since immutable objects do not change the state after initialization, data races can only occur during initialization
- An immutable class is one whose instances are immutable objects



- An immutable object is one whose state cannot be changed after initialization
 - You can think of it as a constant
 - The **final** keyword in Java prevents modification of fields
 - Remember that variables assigned to an object only hold a reference to the object
- Since immutable objects do not change the state after initialization, data races can only occur during initialization
- An immutable class is one whose instances are immutable objects

Immutable class & `final`



· 32

Does defining all fields as `final` ensure that the class is immutable?

Does defining all fields as **final** ensure that the class is immutable?

If in a class, no fields are defined as **final**, is it possible to make it immutable?

- To ensure thread-safety of immutable classes you simply need to make sure:
 - No fields can be modified after publication
 - Objects are safely published
 - Access to inner mutable object do not escape



- To ensure thread-safety of immutable classes you simply need to make sure:
 - No fields can be modified after publication
 - Objects are safely published
 - Access to inner mutable object do not escape

```
public final class ThreeStooges {  
    private final Set<String> stooges = new HashSet<String>();  
  
    public ThreeStooges () {  
        stooges.add("Moe");  
        stooges.add("Larry");  
        stooges.add("Curly");  
    }  
  
    public Boolean isStooge(String name) {  
        return stooges.contains(name)  
    }  
}
```

Goetz p. 47



- To ensure thread-safety of immutable classes you simply need to make sure:
 - No fields can be modified after publication
 - Objects are safely published
 - Access to inner mutable object do not escape

```
public final class ThreeStooges {  
    private final Set<String> stooges = new HashSet<String>();  
  
    public ThreeStooges () {  
        stooges.add("Moe");  
        stooges.add("Larry");  
        stooges.add("Curly");  
    }  
  
    public Boolean isStooge(String name) {  
        return stooges.contains(name)  
    }  
}
```

Why is this class thread-safe?
(tip: there are 3 main reasons)

Goetz p. 47

Mutual exclusion



· 36

- Whenever shared mutable state is accessed by several threads, we must ensure mutual exclusion

- Whenever shared mutable state is accessed by several threads, we must ensure mutual exclusion

Are Monitors a thread-safe class?
(when implemented as a class in OO languages)



- Whenever shared mutable state is accessed by several threads, we must ensure mutual exclusion

Are Monitors a thread-safe class?
(when implemented as a class in OO languages)

Is it always necessary to ensure mutual exclusion
in the methods of thread-safe classes?



- To analyse thread-safe in a class, we must identify/consider:
 - Identify the class state
 - Make sure that mutable class state does not escape
 - Ensure safe publication
 - Whenever possible define class state as immutable
 - If class state must be mutable, ensure mutual exclusion

Interesting section (4.5) on documenting synchronization in Goetz. Unfortunately, not widespread.

- *Instance confinement* refers to encapsulating access to a thread-unsafe object into a thread-safe class

- *Instance confinement* refers to encapsulating access to a thread-unsafe object into a thread-safe class

```
public class PersonSet {  
    private final Set<Person> mySet = new HashSet<Person>();  
  
    public synchronized void addPerson (Person p) {  
        mySet.add(p);  
    }  
  
    public synchronized boolean contains(Person p) {  
        return mySet.contains(p);  
    }  
}
```

Goetz p. 59



- *Instance confinement* refers to encapsulating access to a thread-unsafe object into a thread-safe class

```
public class PersonSet {  
    private final Set<Person> mySet = new HashSet<Person>();  
  
    public synchronized void addPerson (Person p) {  
        mySet.add(p);  
    }  
  
    public synchronized boolean contains(Person p) {  
        return mySet.contains(p);  
    }  
}
```

Goetz p. 59

Why is this class thread-safe?



- Java's standard library provides a method to convert ordinary collections in to "synchronized" collections
 - `synchronizedCollection(Collection<T> c)`, `synchronizedList(List<T> l)`, `synchronizedSet(Set<T> s)`, ..., `synchronizedXXX(XXX<T> x)` with **XXX** a Java collection.
 - Internally, these methods turn all the methods in the collection into synchronized
 - That is, they use the instance lock



- Java's standard library provides a method to convert ordinary collections in to “synchronized” collections
 - `synchronizedCollection(Collection<T> c)`, `synchronizedList(List<T> l)`, `synchronizedSet(Set<T> s)`, ..., `synchronizedXXX(XXX<T> x)` with **XXX** a Java collection.
 - Internally, these methods turn all the methods in the collection into synchronized
 - That is, they use the instance lock

Are synchronized collections thread-safe?



- Java's standard library provides a method to convert ordinary collections in to “synchronized” collections
 - `synchronizedCollection(Collection<T> c)`, `synchronizedList(List<T> l)`, `synchronizedSet(Set<T> s)`, ..., `synchronizedXXX(XXX<T> x)` with **XXX** a Java collection.
- Internally, these methods turn all the methods in the collection into synchronized
 - That is, they use the instance lock

Are synchronized collections thread-safe?

Let's look at the Javadoc

(<https://docs.oracle.com/javase/8/docs/api/java/util/Collections.html#synchronizedList-java.util.List->)

p only accesses thread-safe classes $\nRightarrow p$ is a thread-safe program



· 43

p only accesses thread-safe classes \nRightarrow p is a thread-safe program

· 43



```
List<Integer> l = new ArrayList<Integer>();  
List<Integer> lSync = Collections.synchronizedList(l);  
  
...  
  
new Thread(() -> { addIfAbsent(lSync,l); }).start();  
new Thread(() -> { addIfAbsent(lSync,l); }).start();  
  
...  
  
public void addIfAbsent(List l, Integer e) {  
    if (!l.contains(e))  
        l.add(e);  
}
```

p only accesses thread-safe classes \nRightarrow p is a thread-safe program



Is this program thread-safe?

```
List<Integer> l = new ArrayList<Integer>();  
List<Integer> lSync = Collections.synchronizedList(l);  
  
...  
  
new Thread(() -> { addIfAbsent(lSync,l); }).start();  
new Thread(() -> { addIfAbsent(lSync,l); }).start();  
  
...  
  
public void addIfAbsent(List l, Integer e) {  
    if (!l.contains(e))  
        l.add(e);  
}
```



- Thread-safe classes may be extended to include compound actions
 - Intuitively, compound actions can be seen multiple method calls or field accesses within a critical section
 - A common examples are: *check-and-set*, iteration, navigation (*contains*)

```
public void addIfAbsent(List l, Integer e) {  
    synchronized (l) {  
        if (!l.contains(e))  
            l.add(e);  
    }  
}
```

Thread uses the intrinsic lock of a synchronized collection

```
class ThreadSafeList {  
    ...  
    public void synchronized addIfAbsent(T e) {  
        if (!l.contains(e))  
            l.add(e);  
    }  
    ...  
}
```

Thread-safe class is extended with a custom method to perform the action

Other synchronization primitives (synchronizers)



- Semaphores are synchronization primitives that allow at most c number of threads in the critical section where c is called the *capacity*
 - First introduced by Dijkstra
- A semaphore consists of:
 - An integer capacity (c), [permits in Java](#)
 - Initial number of threads allowed in the critical section
 - A method **acquire()**
 - Checks if $c > 0$, if so, it decrements capacity by one ($c--$) and allows the calling thread to make progress, otherwise it blocks the thread
 - It is a blocking call
 - A method **release()**
 - It checks whether there are waiting threads, if so, it wakes up one of them, otherwise it increases the capacity by one ($c++$)
 - It is non-blocking



- Semaphores are synchronization primitives that allow at most c number of threads in the critical section where c is called the *capacity*
 - First introduced by Dijkstra
- A semaphore consists of:
 - An integer capacity (c), [permits in Java](#)
 - Initial number of threads allowed in the critical section
 - A method **acquire()**
 - Checks if $c > 0$, if so, it decrements capacity by one ($c--$) and allows the calling thread to make progress, otherwise it blocks the thread
 - It is a blocking call
 - A method **release()**
 - It checks whether there are waiting threads, if so, it wakes up one of them, otherwise it increases the capacity by one ($c++$)
 - It is non-blocking

Semaphores (1968) appear
before Monitors (1972)



If we set the capacity of a semaphore to 1, does it behave like a lock?

- Semaphores are synchronization primitives that allow at most c number of threads in the critical section where c is called the *capacity*
 - First introduced by Dijkstra
- A semaphore consists of:
 - An integer capacity (c), [permits in Java](#)
 - Initial number of threads allowed in the critical section
 - A method **acquire()**
 - Checks if $c > 0$, if so, it decrements capacity by one ($c--$) and allows the calling thread to make progress, otherwise it blocks the thread
 - It is a blocking call
 - A method **release()**
 - It checks whether there are waiting threads, if so, it wakes up one of them, otherwise it increases the capacity by one ($c++$)
 - It is non-blocking

Semaphores (1968) appear before Monitors (1972)



If we set the capacity of a semaphore to 1, does it behave like a lock?

- Semaphores are synchronization primitives that allow at most c number of threads in the critical section where c is called the

Synchronization primitives that only allow one thread in the critical section are called **mutex** (which is short for mutual exclusion)

Semaphores (1968) appear before Monitors (1972)

- An integer capacity (c), [permits in Java](#)
 - Initial number of threads allowed in the critical section
- A method **acquire()**
 - Checks if $c > 0$, if so, it decrements capacity by one ($c--$) and allows the calling thread to make progress, otherwise it blocks the thread
 - It is a blocking call
- A method **release()**
 - It checks whether there are waiting threads, if so, it wakes up one of them, otherwise it increases the capacity by one ($c++$)
 - It is non-blocking



- You can think of a semaphore as a “bouncer” to enter a critical section or to be allowed to use a shared resource



- You can think of a semaphore as a “bouncer” to enter a critical section or to be allowed to use a shared resource





- You can think of a semaphore as a “bouncer” to enter a critical section or to be allowed to use a shared resource





- You can think of a semaphore as a “bouncer” to enter a critical section or to be allowed to use a shared resource





- You can think of a semaphore as a “bouncer” to enter a critical section or to be allowed to use a shared resource





- Semaphores are typically used to control the number of threads accessing a resource (here we fix a maximum 5 readers and writers)

```
ReadWriteMonitor m = new ReadWriteMonitor();  
Semaphore semReaders = new Semaphore(5,true);  
Semaphore semWriters = new Semaphore(5,true);  
for (int i = 0; i < 10; i++) {  
    // start a reader  
    new Thread(() -> {  
        m.readLock();  
        semReaders.acquire();  
        // read  
        semReaders.release();  
        m.readUnlock();  
    }).start();  
  
    // start a writer  
    new Thread(() -> {  
        m.writeLock();  
        semWriters.acquire();  
        // write  
        semWriters.acquire();  
        m.writeUnlock();  
    }).start();  
}
```

Java semaphores have a fair flag so that their entry queue prioritizes the longest waiting thread



- Semaphores are typically used to control the number of threads accessing a resource (here we fix a maximum 5 readers and writers)

```
ReadWriteMonitor m = new ReadWriteMonitor();  
Semaphore semReaders = new Semaphore(5,true);  
Semaphore semWriters = new Semaphore(5,true);  
for (int i = 0; i < 10; i++) {  
    // start a reader  
    new Thread() -> {  
        m.readLock();  
        semReaders.acquire();  
        // read  
        semReaders.release();  
        m.readUnlock();  
    }.start();  
  
    // start a writer  
    new Thread() -> {  
        m.writeLock();  
        semWriters.acquire();  
        // write  
        semWriters.acquire();  
        m.writeUnlock();  
    }.start();  
}
```

Java semaphores have a fair flag so that their entry queue prioritizes the longest waiting thread

Does the semaphore make any difference for writers?

See `ReadersWritersSemaphore.java`

- Semaphores are typically used to control the number of threads accessing a resource (here we fix a maximum 5 readers and writers)

```
ReadWriteMonitor m = new ReadWriteMonitor();  
Semaphore semReaders = new Semaphore(5,true);  
Semaphore semWriters = new Semaphore(5,true);  
for (int i = 0; i < 10; i++) {  
    // start a reader  
    new Thread() -> {  
        m.readLock();  
        semReaders.acquire();  
        // read  
        semReaders.release();  
        m.readUnlock();  
    }.start();  
  
    // start a writer  
    new Thread() -> {  
        m.writeLock();  
        semWriters.acquire();  
        // write  
        semWriters.acquire();  
        m.writeUnlock();  
    }.start();  
}
```

Java semaphores have a fair flag so that their entry queue prioritizes the longest waiting thread


Do we need a semaphore to impose this constraint, or can we implement it in the monitor?

Does the semaphore make any difference for writers?


See `ReadersWritersSemaphore.java`

- *Barriers* are synchronization primitives used to wait until several threads reach some point in their computation


- *Barriers* are synchronization primitives used to wait until several thread reach some point in their computation



- *Barriers* are synchronization primitives used to wait until several thread reach some point in their computation



- *Barriers* are synchronization primitives used to wait until several thread reach some point in their computation





- *Barriers* are synchronization primitives used to wait until several thread reach some point in their computation
- Barriers consists of
 - A number *parties* to wait for
 - A method **await()**
 - If the number of waiting threads is less than *parties*, then the calling thread blocks, otherwise all waiting threads wake up and the calling thread is allowed to make progress
- Java includes the class **CyclicBarrier**
 - After *parties* called **await()**, then the state is reset and the barrier behaves as initially

Barrier Example | Parallel initialization



· 61

- Several threads are used to initialize an array (each a different position), the barrier is used for threads to know when the initialization is finished
 - This example is a bit artificial, but it illustrates the use of barriers.

```
int parties          = 10;
CyclicBarrier cb     = new CyclicBarrier(parties);
int[] shared_array = new int[parties];
...
for (int i = 0; i < parties; i++) {
    new SetterClass(i).start();
}
...
public class SetterClass extends Thread {
    int index;
    public SetterClass(int index) {this.index = index;}

    public void run() {
        shared_array[index] = index+1;
        cb.await();
        // After this point the array is initialized and it is safe to read it
    }
}
```

Barrier Example | Parallel initialization



· 61

- Several threads are used to initialize an array (each a different position), the barrier is used for threads to know when the initialization is finished
 - This example is a bit artificial, but it illustrates the use of barriers.

```
int parties          = 10;
CyclicBarrier cb     = new CyclicBarrier(parties);
int[] shared_array = new int[parties];
...
for (int i = 0; i < parties; i++) {
    new SetterClass(i).start();
}
...
public class SetterClass extends Thread {
    int index;
    public SetterClass(int index) {this.index = index;}

    public void run() {
        shared_array[index] = index+1;
        cb.await();
        // After this point the array is initialized and it is safe to read it
    }
}
```

See `BarrierExample.java`

Producer-consumer problem



- Consider a shared data structure of fixed size from which threads may add and remove elements
- Producer threads may add elements to the structure as long as it is not full
 - If the structure is full and a producer tries to add an element, it must block until there an element is removed
- Consumer threads remove elements to the structure as long as it is not empty
 - If the structure is empty and a consumer tries to remove an element, then it must block until an element is added
- A good solution to the problem must be deadlock free and (possibly) starvation free

Producer-consumer problem | Intuition



· 63

- Perhaps more intuitive example

Consumers

Producers



Shared data structure of fixed size

- The producer-consumer problem appears in many multi-threaded situations
 - Handling access to a shared bounded data structure
 - Controlling access to limited computational resources
 - E.g., thread pools
 - Asynchronous I/O operations
 - External devices may act as producers providing data to the system (keyboard, mouse, etc...), or consumer obtaining tasks to perform (IoT devices)

- Definitions of thread-safety
 - Classes
 - Programs
- Safe publication
- Immutability
- Instance confinement
- Synchronization primitives (synchronizers)
 - Semaphores
 - Barriers
- Producer-consumer problem