



Practical Concurrent and Parallel Programming VII

Lock-Free Data Structures

Raúl Pardo

Based on slides of PCPP 2019

Troubles running Gradle?




<https://www.menti.com/aljogfb2pj76>




Previously on PCPP...

Reducing locking



IT UNIVERSITY OF COPENHAGEN

© Raúl Pardo Jimenez and Jørgen Staunstrup – F2023



Previously on PCPP...

Reducing locking



This lecture:
Can we get rid of
locks altogether?



- Compare-And-Swap (CAS)
 - Lock-free atomic integer
 - Lock-free number range
 - Atomic libraries
 - CAS based lock implementation
- Lock-free stack
- ABA problem
- Lock-free queue

HW support for atomic compound operations



- Early processors had atomic compound operations to implement mutexes
 - *test-and-set*
 - Not suitable for implementing advanced lock-free data structures
- Modern processors provide special instructions to for managing concurrent access to shared variables
 - *store-conditional*, **compare-and-swap (CAS)**

See Herlihy Sections 5.1 – 5.8
(outside the scope of the course)





- Intel® 64 and IA-32 Architectures Software Developer's Manual
 - “The **CMPXCHG** [...] If the values contained in the destination operand and the EAX register are equal, the destination operand is replaced with the value of the other source operand (the value not in the EAX register).
[...]
For multiple processor systems, **CMPXCHG** can be combined with the **LOCK** prefix to perform the compare and exchange operation atomically.”



- Intel® 64 and IA-32 Architectures Software Developer's Manual
 - “The **CMPXCHG** [...] If the values contained in the destination operand and the EAX register are equal, the destination operand is replaced with the value of the other source operand (the value not in the EAX register).
[...]
For multiple processor systems, **CMPXCHG** can be combined with the **LOCK** prefix to perform the compare and exchange operation atomically.”
 - Intel 64 and IA-32 processors provide a **LOCK#** signal [...] While this output signal is asserted, requests from other processors or bus agents for control of the bus are blocked [...]

Compare-And-Swap (CAS)



- At the programming language level, most languages support an operation: **`val.compareAndSwap(a, b)`**
 - Compares the value of `val` and `a`, and, if they are equal `val` is set to `b`, otherwise it does nothing. In either case, it returns the current value in `val`, i.e., the value when CAS was executed
 - This instruction is translated to `CMPXCHG` (or similar) by the compiler or runtime environment

CAS pseudo-code

.9

```
class MyAtomicInteger {  
    private int value;    // Visibility ensured by locking  
  
    public synchronized int compareAndSwap(int oldValue, int newValue) {  
        int valueInRegister = this.value;  
        if (this.value == oldValue)  
            this.value = newValue;  
        return valueInRegister;  
    }  
  
    public synchronized boolean compareAndSet(int oldValue, int newValue) {  
        return oldValue == this.compareAndSwap(oldValue, newValue);  
    }  
  
    public synchronized int get() { return this.value; }  
    ...  
}
```

Illustrative implementation of CAS

Common alternative.
Also abbreviated as CAS
(optional exercise: implement
natively compareAndSet)

- Only to *illustrate* CAS semantics
 - CAS is implemented via a HW operation (e.g., CMPXCHG in Intel architectures)
 - In fact, locks are implemented using CAS



- Standard AtomicInteger operations can now be implemented using the CAS operations *without blocking*
- This is an example of *optimistic* concurrency (or non-blocking computation)
 - In a nutshell, trying several times until the operation succeeds

```
class MyAtomicInteger {  
...  
    public int addAndGet(int delta) {  
        int oldValue, newValue;  
        do {  
            oldValue = get();  
            newValue = oldValue + delta;  
        } while (!compareAndSet(oldValue, newValue));  
        return newValue;  
    }  
  
    public int getAndAdd(int delta) {  
        int oldValue, newValue;  
        do {  
            oldValue = get();  
            newValue = oldValue + delta;  
        } while (!compareAndSet(oldValue, newValue));  
        return oldValue;  
    }  
...  
}
```

AtomicInteger operations via CAS



- Standard AtomicInteger operations can now be implemented using the CAS operations *without blocking*
- This is an example of *optimistic* concurrency (or non-blocking computation)
 - In a nutshell, trying several times until the operation succeeds

```
class MyAtomicInteger {  
    ...  
    public int addAndGet(int delta) {  
        int oldValue, newValue;  
        do {  
            oldValue = get();  
            newValue = oldValue + delta;  
        } while (!compareAndSet(oldValue, newValue));  
        return newValue;  
    }  
  
    public int getAndAdd(int delta) {  
        int oldValue, newValue;  
        do {  
            oldValue = get();  
            newValue = oldValue + delta;  
        } while (!compareAndSet(oldValue, newValue));  
        return oldValue;  
    }  
    ...  
}
```

Is this the same
as busy-wait?



- *Busy-wait* is an *alternative to blocking* a thread to wait until some condition holds or to enter the critical section
- The main difference with `lock()` or `await()` is that the thread does not transition to the “blocked” state
- Generally, busy-wait is a bad idea,
 - Threads may consume computing resources to check a condition that has not been updated
 - In this course, we will never ask you to use busy-wait
 - Exercise solutions using busy-wait will be rejected
- Very rarely busy-wait may be preferred over blocking the thread
 - When the thread waits for a very short time it might be more efficient to use busy-wait
 - However, as we have discussed, reasoning about how it takes for a thread to do anything is pointless in concurrency

```
...
// state variables
int i = 0;
Lock l = new ReentrantLock();
...

// method example
public void method(...) {
    l.lock()
    try{
        // busy-wait
        while(i>0) {
            // do nothing
        }
    }
    catch (InterruptedException e) {...}
    finally {l.unlock();}
}
```

Progress conditions for non-blocking DS

· 12



There exist 3 main notions of progress in non-blocking data structures/computation

1. Wait-free: A method of an object implementation is wait-free if every call finishes its execution in a finite number of steps
 - My operations are guaranteed to complete in a bounded number of steps (no matter what other threads do)

Progress conditions for non-blocking DS

· 12



There exist 3 main notions of progress in non-blocking data structures/computation

1. Wait-free: A method of an object implementation is wait-free if every call finishes its execution in a finite number of steps
 - My operations are guaranteed to complete in a bounded number of steps (no matter what other threads do)
2. Lock-free: A method of an object implementation is lock-free if executing the method guarantees that some method call (including concurrent) finishes in a finite number of steps
 - Somebody's operations are guaranteed to complete in a bounded number of my steps
 - Most non-blocking data structures are lock-free, e.g., Treap's stack (see next slides)



There exist 3 main notions of progress in non-blocking data structures/computation

1. Wait-free: A method of an object implementation is wait-free if every call finishes its execution in a finite number of steps
 - My operations are guaranteed to complete in a bounded number of steps (no matter what other threads do)
2. Lock-free: A method of an object implementation is lock-free if executing the method guarantees that some method call (including concurrent) finishes in a finite number of steps
 - Somebody's operations are guaranteed to complete in a bounded number of my steps
 - Most non-blocking data structures are lock-free, e.g., Treabor's stack (see next slides)
3. Obstruction-free: A method of an object implementation is obstruction-free if, from any point after which it executes in isolation, it finishes in a finite number of steps;
 - My operations are guaranteed to complete in a bounded of steps (if I get to execute them)



There exist 3 main notions of progress in non-blocking data structures/computation

1. Wait-free: A method of an object implementation is wait-free if every call finishes its execution in a finite number of steps
 - My operations are guaranteed to complete in a bounded number of steps (no matter what other threads do)
2. Lock-free: A method of an object implementation is lock-free if executing the method guarantees that some method call (including concurrent) finishes in a finite number of steps
 - Somebody's operations are guaranteed to complete in a bounded number of my steps
 - Most non-blocking data structures are lock-free, e.g., Treap's stack (see next slides)
3. Obstruction-free: A method of an object implementation is obstruction-free if, from any point after which it executes in isolation, it finishes in a finite number of steps;
 - My operations are guaranteed to complete in a bounded number of steps (if I get to execute them)

$$wait_{free} \Rightarrow lock_{free} \Rightarrow obstruction_{free}$$

Why non-blocking \neq busy-wait?



- Threads cannot wait forever because other thread finished incorrectly
 - Even in obstruction-free, completion must be guaranteed when the thread runs in isolation
- All progress definitions for non-blocking programs are part of an algorithm composed by multiple threads
 - Threads collaborate with each other



- Standard AtomicInteger operations can now be implemented using the CAS

What type of non-blocking progress applies to this AtomicInteger?

- In a nutshell, trying several times until the operation succeeds

```
class MyAtomicInteger {
...
    public int addAndGet(int delta) {
        int oldValue, newValue;
        do {
            oldValue = get();
            newValue = oldValue + delta;
        } while (!compareAndSet(oldValue, newValue));
        return newValue;
    }

    public int getAndAdd(int delta) {
        int oldValue, newValue;
        do {
            oldValue = get();
            newValue = oldValue + delta;
        } while (!compareAndSet(oldValue, newValue));
        return oldValue;
    }
...
}
```



- In Java, CAS and similar operations are accessed through **AtomicXX** objects
 - These classes support atomic lock-free operations for single variables
 - See package [java.util.concurrent.atomic](https://docs.oracle.com/javase/8/docs/api/java/util/concurrent/atomic/package-summary.html)
- **AtomicXX** CAS operations have the same memory semantics as volatile read/write
 - “The memory effects for accesses and updates of atomics generally follow the rules for volatiles” [[java.util.concurrent.atomic](https://docs.oracle.com/javase/8/docs/api/java/util/concurrent/atomic/package-summary.html)]
 - Unlike in `AtomicXX`, recall that volatile reads/writes are not atomic!

A note on correctness and thread-safety



- The volatile visibility semantics of **AtomicXX** objects allow us to reason about safe-publication
- However, operations on **AtomicXX** are not explicitly part of the Java memory model
 - We cannot use them to reason about correctness or thread-safety
- In this lecture (and exercises), we use specifications to establish correctness (*à la* week 4)
- In two weeks, we will see a finer-grained definition of correctness for concurrent objects, Linearizability

What if we have to update multiple fields?

· 17



```
public class NumberRange {

    private final AtomicInteger lower = new AtomicInteger(0);
    private final AtomicInteger upper = new AtomicInteger(0);

    public int getLower() { return lower.get(); }
    public int getUpper() { return upper.get(); }

    public void setLower(int i) {
        if (i > getUpper()) {
            throw new IllegalArgumentException(
                "Can't set lower to " + i + " > upper");
        }
        lower.set(i);
    }

    public void setUpper(int i) {
        if (i < getLower()) {
            throw new IllegalArgumentException(
                ("Can't set upper to " + i + " < lower"));
        }
        upper.set(i);
    }
}
```

Specification: For any concurrent execution of `setLower()` and `setUpper()`, it must always hold that **$\text{lower} \leq \text{upper}$**

What if we have to update multiple fields?

· 17



```
public class NumberRange {

    private final AtomicInteger lower = new AtomicInteger(0);
    private final AtomicInteger upper = new AtomicInteger(0);

    public int getLower() { return lower.get(); }
    public int getUpper() { return upper.get(); }

    public void setLower(int i) {
        if (i > getUpper()) {
            throw new IllegalArgumentException(
                "Can't set lower to " + i + " > upper");
        }
        lower.set(i);
    }

    public void setUpper(int i) {
        if (i < getLower()) {
            throw new IllegalArgumentException(
                ("Can't set upper to " + i + " < lower"));
        }
        upper.set(i);
    }
}
```

Specification: For any concurrent execution of `setLower()` and `setUpper()`, it must always hold that $\text{lower} \leq \text{upper}$

Does this implementation of `setLower` and `setUpper` satisfy the specification?

What if we have to update multiple fields?

· 18



```
public class NumberRange {  
  
    private final AtomicInteger lower = new AtomicInteger(0);  
    private final AtomicInteger upper = new AtomicInteger(0);  
  
    public int getLower() { return lower.get(); }  
    public int getUpper() { return upper.get(); }  
  
    public void setLower(int i) {  
        if (i > getUpper()) {  
            throw new IllegalArgumentException(  
                "Can't set lower to " + i + " > upper");  
        }  
        lower.set(i);  
    }  
  
    public void setUpper(int i) {  
        if (i < getLower()) {  
            throw new IllegalArgumentException(  
                ("Can't set upper to " + i + " < lower");  
        }  
        upper.set(i);  
    }  
}
```

Specification: For any concurrent execution of `setLower()` and `setUpper()`, it must always hold that **lower ≤ upper**

Goetz p. 67, Don't do this



What if we have to update multiple fields?

· 19



```
public class CasNumberRange {  
  
    private final AtomicReference<IntPair> values =  
        new AtomicReference<IntPair>(new IntPair(0,0));  
  
    public int getLower() { return values.get().lower;}  
    public int getUpper() { return values.get().upper;}  
  
    public void setLower(int i) {  
        while(true) {  
            IntPair oldv = values.get();  
            if (i > oldv.upper) {  
                throw new IllegalArgumentException(  
                    "Can't set lower to " + i + " > upper");  
            }  
            IntPair newv = new IntPair(i, oldv.upper);  
            if(values.compareAndSet(oldv,newv))  
                return;  
        }  
    }  
  
    ... // Similarly for setUpper();  
}
```

```
// Immutable class  
private static class IntPair {  
    private final int lower;  
    private final int upper;  
  
    public IntPair(int lower,  
                   int upper) {  
        this.lower = lower;  
        this.upper = upper;  
    }  
}
```

Specification: For any concurrent execution of `setLower()` and `setUpper()`, it must always hold that $\text{lower} \leq \text{upper}$

What if we have to update multiple fields?

· 19



```
public class CasNumberRange {  
  
    private final AtomicReference<IntPair> values =  
        new AtomicReference<IntPair>(new IntPair(0,0));  
  
    public int getLower() { return values.get().lower;}  
    public int getUpper() { return values.get().upper;}  
  
    public void setLower(int i) {  
        while(true) {  
            IntPair oldv = values.get();  
            if (i > oldv.upper) {  
                throw new IllegalArgumentException(  
                    "Can't set lower to " + i + " > upper");  
            }  
            IntPair newv = new IntPair(i, oldv.upper);  
            if(values.compareAndSet(oldv,newv))  
                return;  
        }  
    }  
  
    ... // Similarly for setUpper();  
}
```

```
// Immutable class  
private static class IntPair {  
    private final int lower;  
    private final int upper;  
  
    public IntPair(int lower,  
                   int upper) {  
        this.lower = lower;  
        this.upper = upper;  
    }  
}
```

Why is it important that
IntPair is immutable?

Specification: For any concurrent
execution of setLower() and
setUpper(), it must always hold
that lower ≤ upper

CAS based Lock implementation

Implementing Locks using CAS



1. Simple TryLock

- Non-blocking tryLock, the lock may be acquired only once
- Regular unlock

2. Reentrant TryLock

- Non-blocking tryLock, the lock may be acquired repeatedly by the same thread
- Regular (reentrant) unlock

3. Simple Lock

- Blocking lock, the lock may be acquired only once
- Regular unlock

4. Reentrant Lock

- Blocking lock, the lock may be acquired repeatedly by the same thread
- Regular (reentrant) unlock

See `TestCasLocks.java`

1. SimpleTryLock, non-blocking



```
class SimpleTryLock {  
  
    // Refers to holding thread, null iff unheld  
    private final AtomicReference<Thread> holder = new AtomicReference<Thread>();  
  
    public boolean tryLock() {  
        final Thread current = Thread.currentThread();  
        return holder.compareAndSet(null, current);  
    }  
  
    public void unlock() {  
        final Thread current = Thread.currentThread();  
        if (!holder.compareAndSet(current, null))  
            throw new RuntimeException("Not lock holder");  
    }  
}
```

If the lock is free (holder *value* equals null), takes it and return true. Otherwise, holder is unmodified and returns false.

Sets holder *value* to null. If CAS returns false throws an exception indicating that this thread is not holding the lock.

1. SimpleTryLock, non-blocking




```
class SimpleTryLock {  
  
    // Refers to holding thread, null iff unheld  
    private final AtomicReference<Thread> holder = new AtomicReference<Thread>();  
  
    public boolean tryLock() {  
        final Thread current = Thread.currentThread();  
        return holder.compareAndSet(null, current);  
    }  
  
    public void unlock() {  
        final Thread current = Thread.currentThread();  
        if (!holder.compareAndSet(current, null))  
            throw new RuntimeException("Not lock holder");  
    }  
}
```

If the lock is free (holder *value* equals null), takes it and return true. Otherwise, holder is unmodified and returns false.

How is CAS used here to determine that the lock is not held by the thread executing unlock?

Sets holder *value* to null. If CAS returns false throws an exception indicating that this thread is not holding the lock.

2. ReentrantTryLock, non-blocking

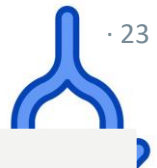


```
class ReentrantTryLock {
    private final AtomicReference<Thread> holder = new AtomicReference<Thread>();
    private volatile int holdCount = 0;

    public boolean tryLock() {
        final Thread current = Thread.currentThread();
        if (holder.get() == current) {
            holdCount++;
            return true;
        } else if (holder.compareAndSet(null, current)) {
            holdCount = 1;
            return true;
        }
        return false;
    }

    public void unlock() {
        final Thread current = Thread.currentThread();
        if (holder.get() == current) {
            holdCount--;
            if (holdCount == 0)
                holder.compareAndSet(current, null);
            return;
        } else
            throw new RuntimeException("Not lock holder");
    }
}
```

2. ReentrantTryLock, non-blocking



```
class ReentrantTryLock {  
    private final AtomicReference<Thread> holder = new AtomicReference<Thread>();  
    private volatile int holdCount = 0;
```


```
    public boolean tryLock() {  
        final Thread current = Thread.currentThread();  
        if (holder.get() == current) {  
            holdCount++;  
            return true;  
        } else if (holder.compareAndSet(null, current)) {  
            holdCount = 1;  
            return true;  
        }  
        return false;  
    }  
}
```

If the calling thread already holds the lock, we increase the counter

If not, we try to acquire the lock and set the count to 1

```
    public void unlock() {  
        final Thread current = Thread.currentThread();  
        if (holder.get() == current) {  
            holdCount--;  
            if (holdCount == 0)  
                holder.compareAndSet(current, null);  
            return;  
        } else  
            throw new RuntimeException("Not lock holder");  
    }  
}
```


2. ReentrantTryLock, non-blocking



```
class ReentrantTryLock {
    private static final AtomicReference<Thread> holder = new AtomicReference<Thread>();
    private int holdCount = 0;

    public boolean tryLock() {
        final Thread current = Thread.currentThread();
        if (holder.get() == current) {
            holdCount++;
            return true;
        } else if (holder.compareAndSet(null, current)) {
            holdCount = 1;
            return true;
        }
        return false;
    }

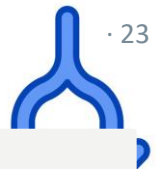
    public void unlock() {
        final Thread current = Thread.currentThread();
        if (holder.get() == current) {
            holdCount--;
            if (holdCount == 0)
                holder.compareAndSet(current, null);
            return;
        } else
            throw new RuntimeException("Not lock holder");
    }
}
```

Do we have a race condition in this branch of the if-statement?

If the calling thread already holds the lock, we increase the counter

If not, we try to acquire the lock and set the count to 1

2. ReentrantTryLock, non-blocking



```
class ReentrantTryLock {  
    private final AtomicReference<Thread> holder = new AtomicReference<Thread>();  
    private volatile int holdCount = 0;
```

```
    public boolean tryLock() {  
        final Thread current = Thread.currentThread();  
        if (holder.get() == current) {  
            holdCount++;  
            return true;  
        } else if (holder.compareAndSet(null, current)) {  
            holdCount = 1;  
            return true;  
        }  
        return false;  
    }  
}
```

If the calling thread already holds the lock, we increase the counter

If not, we try to acquire the lock and set the count to 1

```
    public void unlock() {  
        final Thread current = Thread.currentThread();  
        if (holder.get() == current) {  
            holdCount--;  
            if (holdCount == 0)  
                holder.compareAndSet(current, null);  
            return;  
        } else  
            throw new RuntimeException("Not lock holder");  
    }  
}
```

If the calling thread holds the hold we decrement the counter

If the counter is 0 then we release the lock

If the calling thread does not hold the lock we throw an exception

3. SimpleLock, blocking



```
class SimpleLock {
    private final AtomicReference<Thread> holder = new AtomicReference<Thread>();
    // The FIFO queue of threads waiting for this lock
    private final Queue<Thread> waiters = new ConcurrentLinkedQueue<Thread>();

    public void lock() {
        final Thread current = Thread.currentThread();
        waiters.add(current);
        while (waiters.peek() != current || !holder.compareAndSet(null, current)) {
            LockSupport.park(this);
        }
        waiters.remove();
    }

    public void unlock() {
        final Thread current = Thread.currentThread();
        if (holder.compareAndSet(current, null))
            LockSupport.unpark(waiters.peek());
        else
            throw new RuntimeException("Not lock holder");
    }
}
```

3. SimpleLock, blocking



```
class SimpleLock {  
    private final AtomicReference<Thread> holder = new AtomicReference<Thread>();  
    // The FIFO queue of threads waiting for this lock  
    private final Queue<Thread> waiters = new ConcurrentLinkedQueue<Thread>();
```

```
    public void lock() {  
        final Thread current = Thread.currentThread();  
        waiters.add(current);  
        while (waiters.peek() != current || !holder.compareAndSet(null, current)) {  
            LockSupport.park(this);  
        }  
        waiters.remove();  
    }
```

First we add the thread to a waiting queue

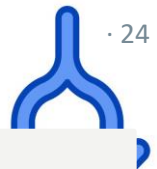
We check whether the current thread is at the head of the queue and try to take the lock

Finally, we remove the thread from the waiting list

If not successful, we put the thread to wait (see Javadoc LockSupport)

```
    public void unlock() {  
        final Thread current = Thread.currentThread();  
        if (holder.compareAndSet(current, null))  
            LockSupport.unpark(waiters.peek());  
        else  
            throw new RuntimeException("Not lock holder");  
    }  
}
```

3. SimpleLock, blocking



```
class SimpleLock {  
    private final AtomicReference<Thread> holder = new AtomicReference<Thread>();  
    // The FIFO queue of threads waiting for this lock  
    private final Queue<Thread> waiters = new ConcurrentLinkedQueue<Thread>();
```

```
    public void lock() {  
        final Thread current = Thread.currentThread();  
        waiters.add(current);  
        while (waiters.peek() != current || !holder.compareAndSet(null, current)) {  
            LockSupport.park(this);  
        }  
        waiters.remove();  
    }
```

First we add the thread to a waiting queue

We check whether the current thread is at the head of the queue and try to take the lock

Finally, we remove the thread from the waiting list

If not successful, we put the thread to wait (see Javadoc LockSupport)

```
    public void unlock() {  
        final Thread current = Thread.currentThread();  
        if (holder.compareAndSet(current, null))  
            LockSupport.unpark(waiters.peek());  
        else  
            throw new RuntimeException("Not lock holder");  
    }
```

We simply try to release the lock

If successful, we wake up one of the waiting threads, if any

As before, we throw an exception if the calling thread is not the thread holding the lock

3. SimpleLock, blocking (interruptions)



```
class SimpleLock {
    private final AtomicReference<Thread> holder = new AtomicReference<Thread>();
    // The FIFO queue of threads waiting for this lock
    private final Queue<Thread> waiters = new ConcurrentLinkedQueue<Thread>();

    public void lock() {
        final Thread current = Thread.currentThread();
        boolean wasInterrupted = false;
        waiters.add(current);
        while (waiters.peek() != current || !holder.compareAndSet(null, current)) {
            LockSupport.park(this);
            if (Thread.interrupted())
                wasInterrupted = true;
        }
        waiters.remove();
        if (wasInterrupted)
            current.interrupt();
    }
    ...
}
```

We use a flag to record whether the thread has been interrupted during execution

Immediately after taking the lock, we check if the thread was interrupted

If the thread was interrupted, we call interrupt to propagate the interruption.

4. SimpleReentrantLock, blocking

· 26



```
class MyReentrantLock {
    private final AtomicReference<Thread> holder = new AtomicReference<Thread>();
    private final Queue<Thread> waiters = new ConcurrentLinkedQueue<Thread>();
    private volatile int holdCount = 0;

    public void lock() {
        final Thread current = Thread.currentThread();
        if (holder.get() == current)
            holdCount++;
        else {
            boolean wasInterrupted = false;
            waiters.add(current);
            while (waiters.peek() != current || !holder.compareAndSet(null, current)) {
                LockSupport.park(this);
                if (Thread.interrupted())
                    wasInterrupted = true;
            }
            holdCount = 1;
            waiters.remove();
            if (wasInterrupted)
                current.interrupt();
        }
    }
    ...
}
```

Here we simply combine what we have seen in threads 2 and 3

Scalability of CAS

- Pros
 - A CAS operation is faster than acquiring a lock
 - An unsuccessful CAS operation does not cause thread de-scheduling (blocking)
- Cons
 - CAS operations result in high memory overhead

Will CAS based implementations always scale well?
(that is, better than lock based)



```
// Lock-based implementation
class LockingRandom implements MyRandom {
    private long seed;
    ...
    // Recipe from java.util.Random.next(int bits) documentation
    public synchronized int nextInt() {
        seed = (seed * 0x5DEECE66DL + 0xBL) & ((1L << 48) - 1);
        return (int)(seed >>> 16);
    }
}

// CAS based implementation
class CasRandom implements MyRandom {
    private final AtomicLong seed;
    ...
    public int nextInt() {
        long oldSeed, newSeed;
        do {
            oldSeed = seed.get();
            newSeed = (oldSeed * 0x5DEECE66DL + 0xBL) & ((1L << 48) - 1);
        } while (!seed.compareAndSet(oldSeed, newSeed));
        return (int)(newSeed >>> 16);
    }
}
```

- We use different implementations of a simple Pseudo-Random Number Generator (PRNG) to test the scalability of CAS vs Locks vs Thread Local (see next slide)
- Example execution with seed = 10
 - 3847489, 1334288366,
1486862010, 711662464,
-1453296530, -775316920,
1157481928, 294681619, ...



```
// Lock-based implementation
class LockingRandom implements MyRandom {
    private long seed;
    ...
    // Recipe from java.util.Random.next(int bits) documentation
    public synchronized int nextInt() {
        seed = (seed * 0x5DEECE66DL + 0xBL) & ((1L << 48) - 1);
        return (int)(seed >>> 16);
    }
}

// CAS based implementation
class CasRandom implements MyRandom {
    private final AtomicLong seed;
    ...
    public int nextInt() {
        long oldSeed, newSeed;
        do {
            oldSeed = seed.get();
            newSeed = (oldSeed * 0x5DEECE66DL + 0xBL) & ((1L << 48) - 1);
        } while (!seed.compareAndSet(oldSeed, newSeed));
        return (int)(newSeed >>> 16);
    }
}
```


- We use different implementations of a simple Pseudo-Random Number Generator (PRNG) to test the scalability of CAS vs Locks vs Thread Local (see next slide)
- Example execution with seed = 10
 - 3847489, 1334288366, 1486862010, 711662464, -1453296530, -775316920, 1157481928, 294681619, ...

Could we have an implementation using volatile?

Scalability of PRNGs (unrealistic contention)



· 30



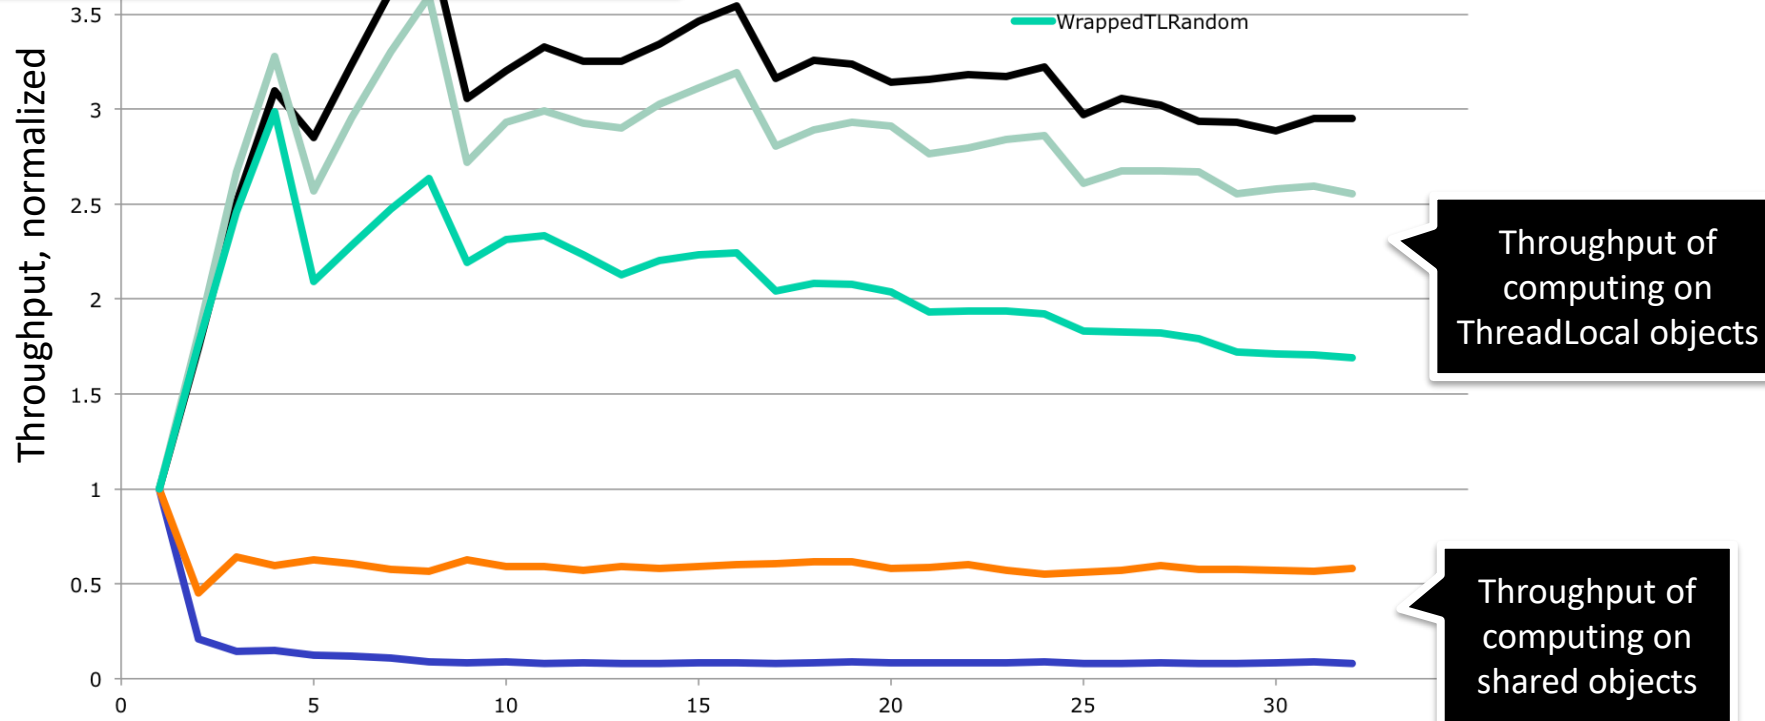
Scalability of PRNGs (unrealistic contention)



· 30


Locality has larger impact on scalability than “lock-freedom”

See `TestPseudoRandom.java`



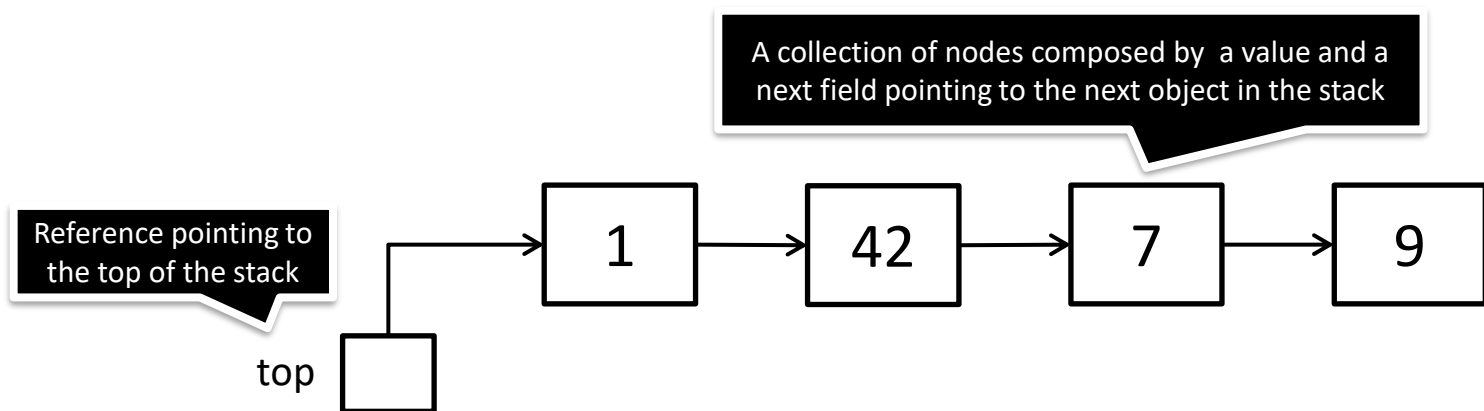
Lock-free data structures: Stack

- A stack is a data structure following a LIFO (*last-in-first-out*) policy
 - `push()` – adds an element to the top of the stack
 - `pop()` – removes the top of the stack
- It is typically implemented as a linked list





- A stack is a data structure following a LIFO (*last-in-first-out*) policy
 - `push()` – adds an element to the top of the stack
 - `pop()` – removes the top of the stack
- It is typically implemented as a linked list



Lock-free Stack

Introduced by
Treiber in 1986

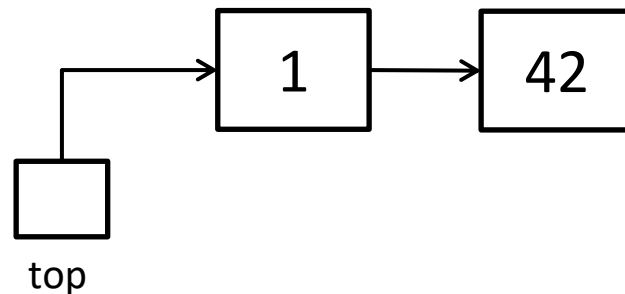
· 33



```
class LockFreeStack<T> {  
    AtomicReference<Node<T>> top =  
        new AtomicReference<Node<T>>();  
  
    public void push(T value) {  
        Node<T> newHead = new Node<T>(value);  
        Node<T> oldHead;  
        do {  
            oldHead = top.get();  
            newHead.next = oldHead;  
        } while (!top.compareAndSet(oldHead, newHead));  
    }  
  
    public T pop() {  
        Node<T> newHead;  
        Node<T> oldHead;  
        do {  
            oldHead = top.get();  
            if(oldHead == null) { return null; }  
            newHead = oldHead.next;  
        } while (!top.compareAndSet(oldHead, newHead));  
        return oldHead.value;  
    }  
}
```

See `TestLockFreeStack.java`

```
private static class Node {  
    private final T item;  
    private Node<T> next;  
  
    public Node(T item) {  
        this.item = item;  
        this.next = null;  
    }  
}
```

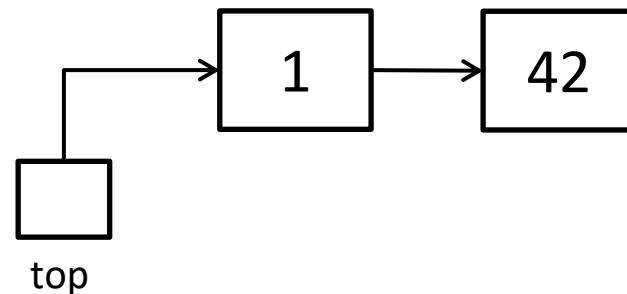


Lock-free Heap | push()

· 34



```
class LockFreeStack<T> {  
    AtomicReference<Node<T>> top =  
        new AtomicReference<Node<T>>();  
  
    public void push(T value) {  
        Node<T> newHead = new Node<T>(value);  
        Node<T> oldHead;  
        do {  
            oldHead = top.get();  
            newHead.next = oldHead;  
        } while (!top.compareAndSet(oldHead, newHead));  
    }  
    ...  
}
```




Lock-free Heap | push()

· 34



```
class LockFreeStack<T> {  
    AtomicReference<Node<T>> top =  
        new AtomicReference<Node<T>>();  
  
    public void push(T value) {  
        Node<T> newHead = new Node<T>(value);  
        Node<T> oldHead;  
        do {  
            oldHead = top.get();  
            newHead.next = oldHead;  
        } while (!top.compareAndSet(oldHead, newHead));  
    }  
    ...  
}
```

newHead




Lock-free Heap | push()

· 34



```
class LockFreeStack<T> {  
    AtomicReference<Node<T>> top =  
        new AtomicReference<Node<T>>();  
  
    public void push(T value) {  
        Node<T> newHead = new Node<T>(value);  
        Node<T> oldHead;  
        do {  
            oldHead = top.get();  
            newHead.next = oldHead;  
        } while (!top.compareAndSet(oldHead, newHead));  
    }  
    ...  
}
```

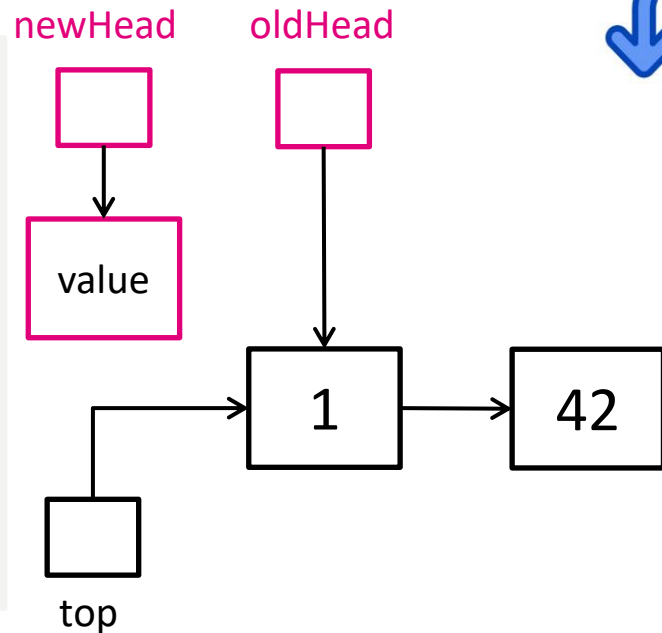


Lock-free Heap | push()

· 34



```
class LockFreeStack<T> {  
    AtomicReference<Node<T>> top =  
        new AtomicReference<Node<T>>();  
  
    public void push(T value) {  
        Node<T> newHead = new Node<T>(value);  
        Node<T> oldHead;  
        do {  
            oldHead = top.get();  
            newHead.next = oldHead;  
        } while (!top.compareAndSet(oldHead, newHead));  
    }  
    ...  
}
```

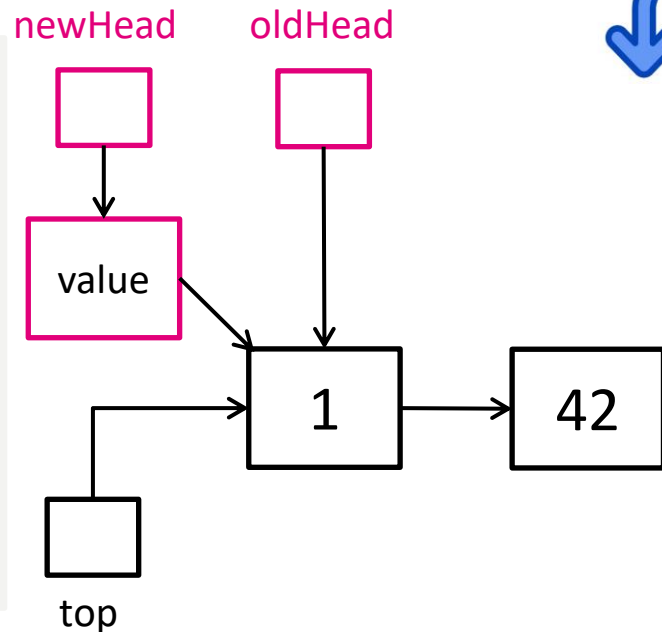


Lock-free Heap | push()

· 34



```
class LockFreeStack<T> {  
    AtomicReference<Node<T>> top =  
        new AtomicReference<Node<T>>();  
  
    public void push(T value) {  
        Node<T> newHead = new Node<T>(value);  
        Node<T> oldHead;  
        do {  
            oldHead = top.get();  
            newHead.next = oldHead;  
        } while (!top.compareAndSet(oldHead, newHead));  
    }  
    ...  
}
```

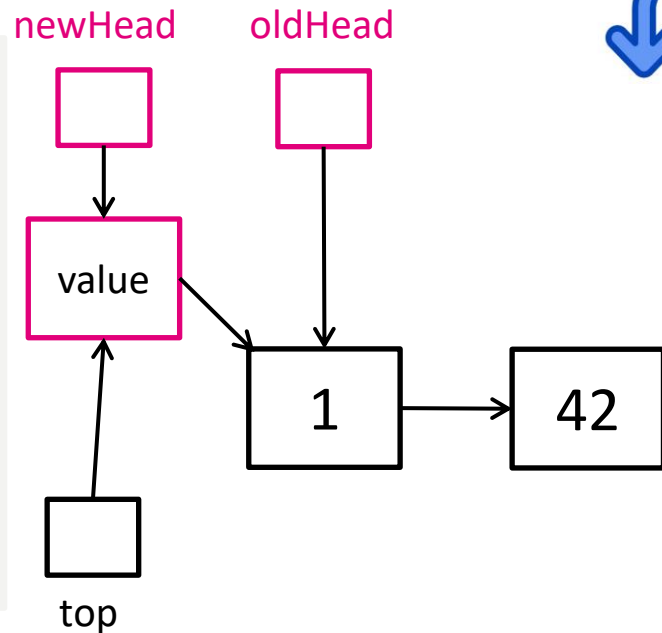


Lock-free Heap | push()

· 34



```
class LockFreeStack<T> {  
    AtomicReference<Node<T>> top =  
        new AtomicReference<Node<T>>();  
  
    public void push(T value) {  
        Node<T> newHead = new Node<T>(value);  
        Node<T> oldHead;  
        do {  
            oldHead = top.get();  
            newHead.next = oldHead;  
        } while (!top.compareAndSet(oldHead, newHead));  
    }  
    ...  
}
```



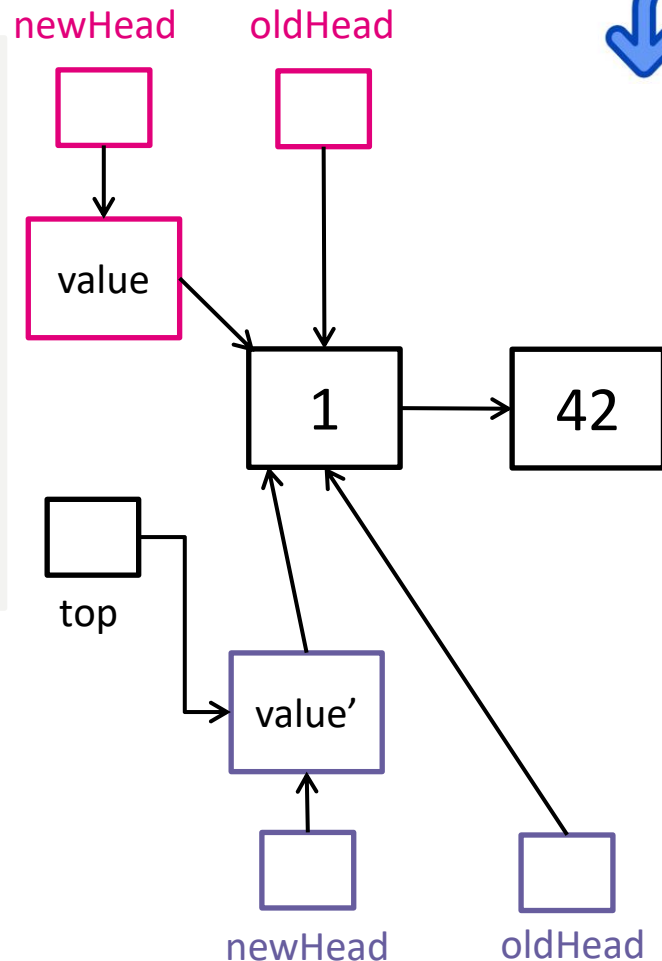
Lock-free Heap | push()

· 34



```
class LockFreeStack<T> {  
    AtomicReference<Node<T>> top =  
        new AtomicReference<Node<T>>();  
  
    public void push(T value) {  
        Node<T> newHead = new Node<T>(value);  
        Node<T> oldHead;  
        do {  
            oldHead = top.get();  
            newHead.next = oldHead;  
        } while (!top.compareAndSet(oldHead, newHead));  
    }  
    ...  
}
```

If before the pink thread executes CAS another thread had pushed concurrently, then CAS fails and push restarts




Lock-free Heap | push()

· 34



```
class LockFreeStack<T> {  
    AtomicReference<Node<T>> top =  
        new AtomicReference<Node<T>>();  
  
    public void push(T value) {  
        Node<T> newHead = new Node<T>(value);  
        Node<T> oldHead;  
        do {  
            oldHead = top.get();  
            newHead.next = oldHead;  
        } while (!top.compareAndSet(oldHead, newHead));  
    }  
    ...  
}
```

If before the pink thread executes CAS another thread had pushed concurrently, then CAS fails and push restarts



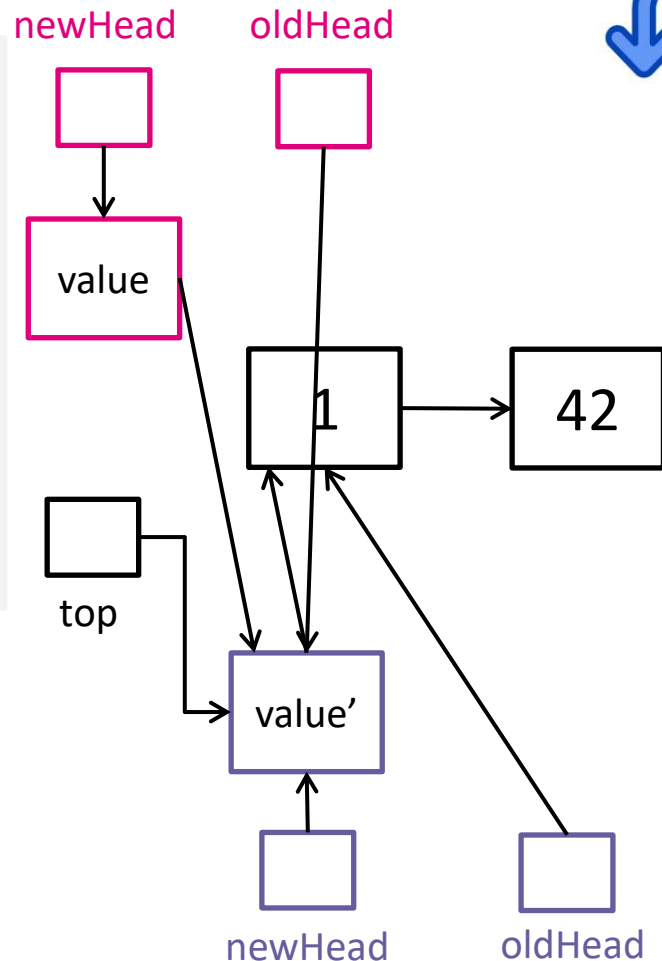
Lock-free Heap | push()

· 34



```
class LockFreeStack<T> {  
    AtomicReference<Node<T>> top =  
        new AtomicReference<Node<T>>();  
  
    public void push(T value) {  
        Node<T> newHead = new Node<T>(value);  
        Node<T> oldHead;  
        do {  
            oldHead = top.get();  
            newHead.next = oldHead;  
        } while (!top.compareAndSet(oldHead, newHead));  
    }  
    ...  
}
```

If before the pink thread executes CAS another thread had pushed concurrently, then CAS fails and push restarts



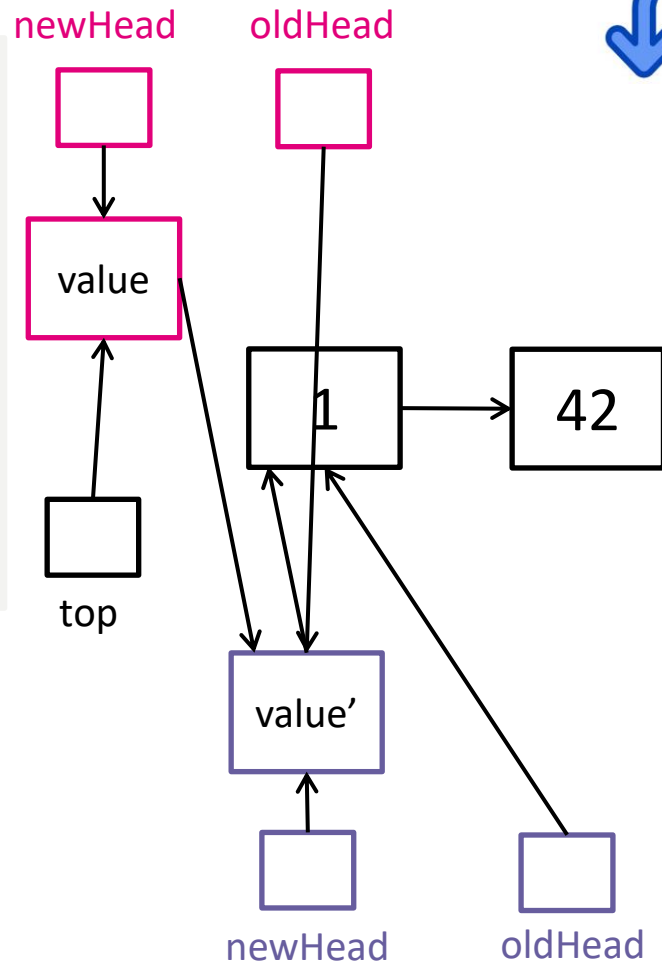
Lock-free Heap | push()

· 34



```
class LockFreeStack<T> {  
    AtomicReference<Node<T>> top =  
        new AtomicReference<Node<T>>();  
  
    public void push(T value) {  
        Node<T> newHead = new Node<T>(value);  
        Node<T> oldHead;  
        do {  
            oldHead = top.get();  
            newHead.next = oldHead;  
        } while (!top.compareAndSet(oldHead, newHead));  
    }  
    ...  
}
```

If before the pink thread executes CAS another thread had pushed concurrently, then CAS fails and push restarts



Lock-free Heap | pop()

· 35

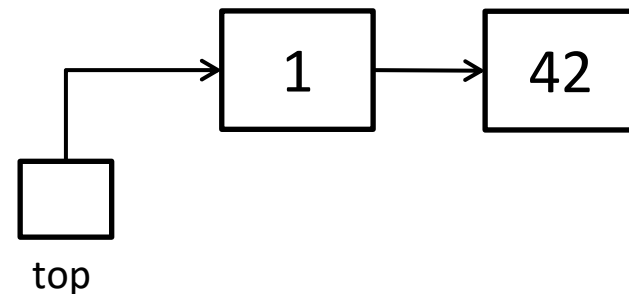


```
class LockFreeStack<T> {  
...  
    public T pop() {  
        Node<T> newHead;  
        Node<T> oldHead;  
        do {  
            oldHead = top.get();  
            if(oldHead == null) { return null; }  
            newHead = oldHead.next;  
        } while (!top.compareAndSet(oldHead,newHead));  
        return oldHead.value;  
    }  
}
```

newHead



oldHead

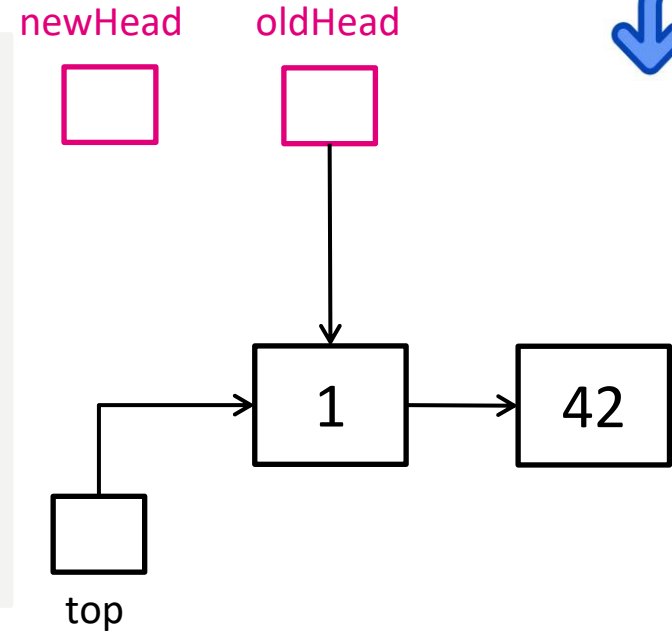


Lock-free Heap | pop()

· 35



```
class LockFreeStack<T> {  
...  
    public T pop() {  
        Node<T> newHead;  
        Node<T> oldHead;  
        do {  
            oldHead = top.get();  
            if(oldHead == null) { return null; }  
            newHead = oldHead.next;  
        } while (!top.compareAndSet(oldHead, newHead));  
        return oldHead.value;  
    }  
}
```




Lock-free Heap | pop()

· 35



```
class LockFreeStack<T> {  
...  
    public T pop() {  
        Node<T> newHead;  
        Node<T> oldHead;  
        do {  
            oldHead = top.get();  
            if(oldHead == null) { return null; }  
            newHead = oldHead.next;  
        } while (!top.compareAndSet(oldHead,newHead));  
        return oldHead.value;  
    }  
}
```

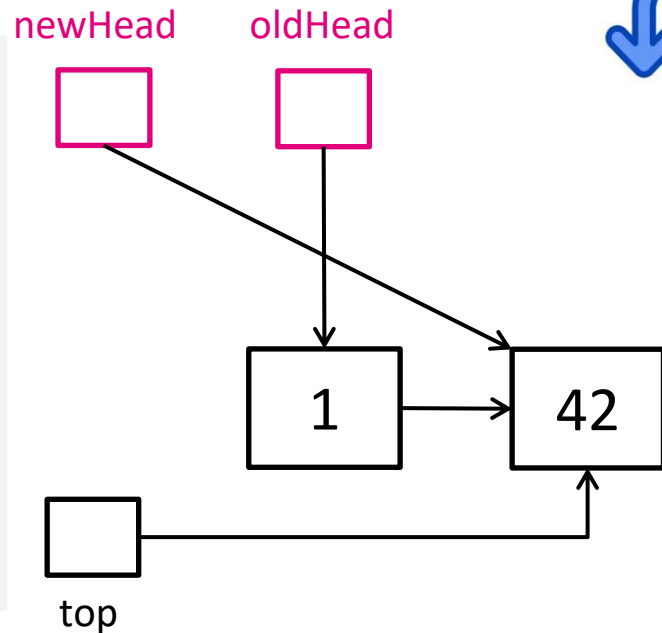


Lock-free Heap | pop()

· 35



```
class LockFreeStack<T> {  
...  
    public T pop() {  
        Node<T> newHead;  
        Node<T> oldHead;  
        do {  
            oldHead = top.get();  
            if(oldHead == null) { return null; }  
            newHead = oldHead.next;  
        } while (!top.compareAndSet(oldHead, newHead));  
        return oldHead.value;  
    }  
}
```

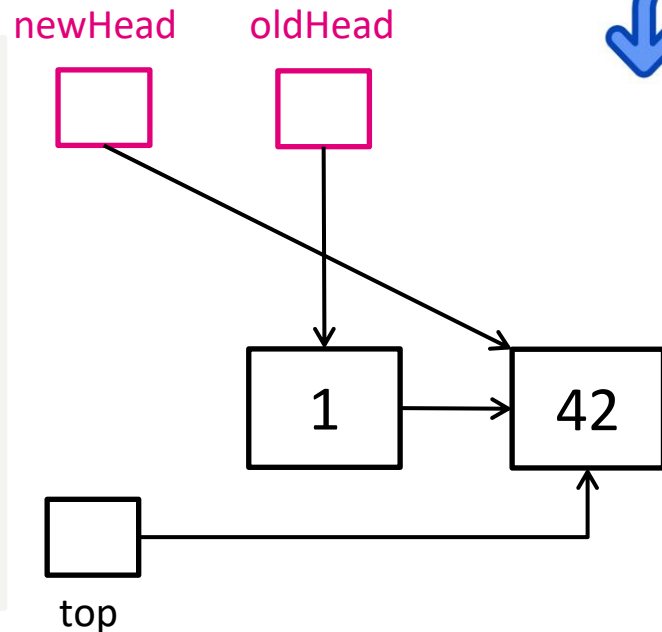


Lock-free Heap | pop()

· 35



```
class LockFreeStack<T> {  
...  
    public T pop() {  
        Node<T> newHead;  
        Node<T> oldHead;  
        do {  
            oldHead = top.get();  
            if(oldHead == null) { return null; }  
            newHead = oldHead.next;  
        } while (!top.compareAndSet(oldHead, newHead));  
        return oldHead.value;  
    }  
}
```



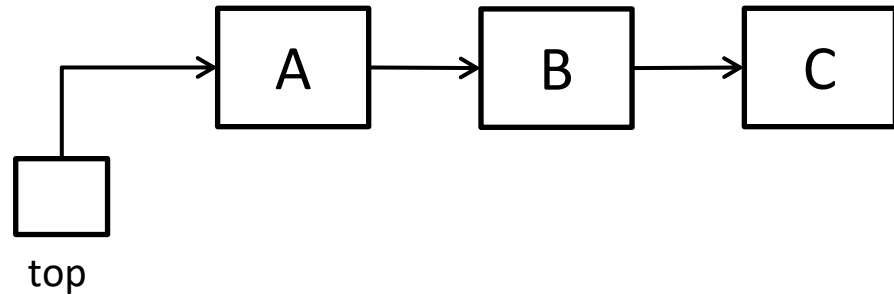
Would there be any problem if another thread is executing `pop()` concurrently?

The ABA Problem

ABA Problem



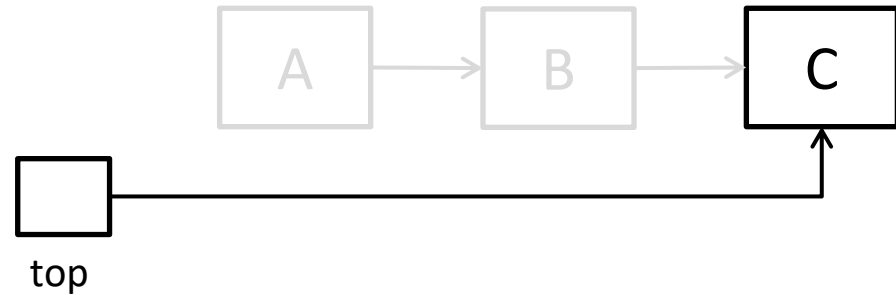
1. Thread 1 starts popping A



ABA Problem



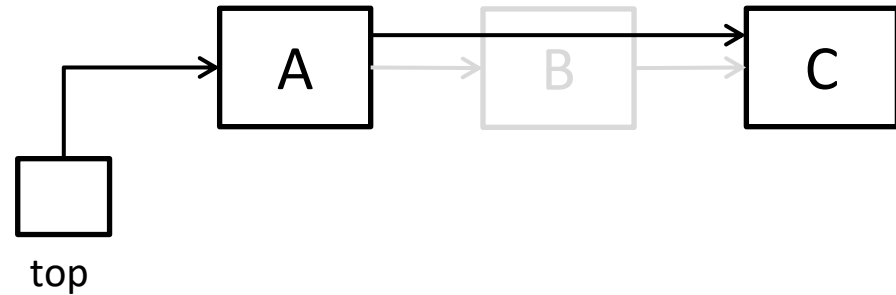
1. Thread 1 starts popping A
2. Before thread 1 finishes, thread 2 pops A and B



ABA Problem



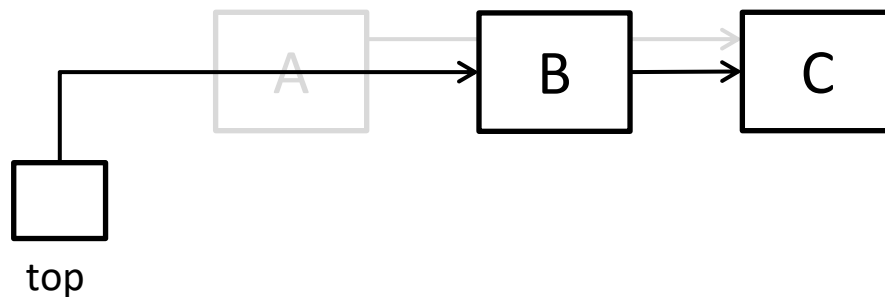
1. Thread 1 starts popping A
2. Before thread 1 finishes, thread 2 pops A and B
3. Thread 2 pushes A back *// recovered from memory (same as thread 1 was popping)*



ABA Problem



1. Thread 1 starts popping A
2. Before thread 1 finishes, thread 2 pops A and B
3. Thread 2 pushes A back *// recovered from memory (same as thread 1 was popping)*
4. Thread 1 finishes popping A
 1. Incorrectly, as it thinks that A is the same as before thread 2 operations

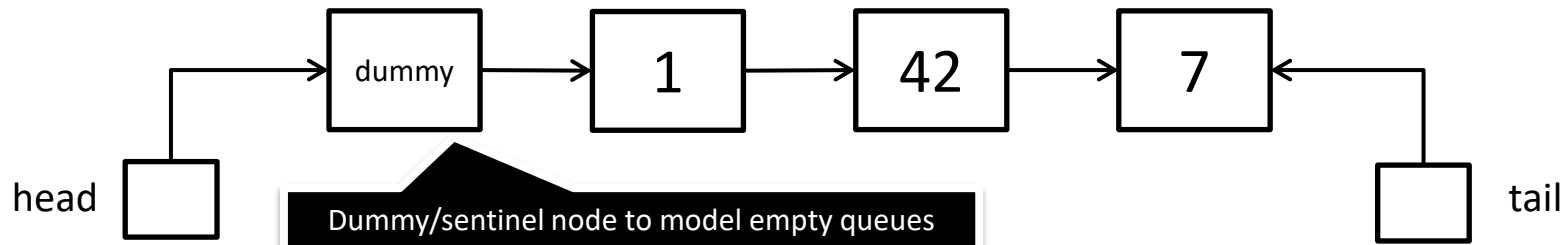




- It is a memory allocation issue which affects mainly languages without garbage collection (e.g., C and C++)
- Languages such as Java do not have this problem because garbage collection ensures that newly created objects are fresh
 - Step 3 in the previous slides would have created a new A object

Lock-free data structures: Queue

- A queue is a data structure following a FIFO (*first-in-first-out*) policy
 - enqueue() – adds an element to the tail of the queue
 - dequeue() – removes an element from the head of the queue
- It is typically implemented as a linked list



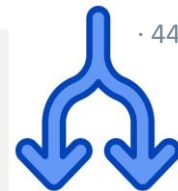
Lock-free queue

- Michael-Scott lock free queue, introduced in 1996 (see optional readings)
- Implemented in ConcurrentLinkedQueue in java.concurrent.* by Doug Lea et. al. (see [here](#))
 - The version on the right is not the JDK implementation

```
private static class Node<T> {  
    final T item;  
    final AtomicReference<Node<T>> next;  
  
    public Node(T item, Node<T> next) {  
        this.item = item;  
        this.next = new AtomicReference<Node<T>>(next);  
    }  
}
```

```
class MSQueue<T> implements UnboundedQueue<T> {  
    private final AtomicReference<Node<T>> head, tail;  
  
    public MSQueue() {  
        Node<T> dummy = new Node<T>(null, null);  
        head = new AtomicReference<Node<T>>(dummy);  
        tail = new AtomicReference<Node<T>>(dummy);  
    }  
  
    public void enqueue(T item) {  
        Node<T> node = new Node<T>(item, null);  
        while (true) {  
            Node<T> last = tail.get(), next = last.next.get();  
            if (last == tail.get()) {  
                if (next == null) {  
                    // In quiescent state, try inserting new node  
                    if (last.next.compareAndSet(next, node)) {  
                        // Insertion succeeded, try advancing tail  
                        tail.compareAndSet(last, node);  
                        return;  
                    }  
                } else  
                    // Queue in intermediate state, advance tail  
                    tail.compareAndSet(last, next);  
            }  
        }  
    }  
  
    public T dequeue() {  
        while (true) {  
            Node<T> first = head.get(), last = tail.get(), next = first.next.get();  
            if (first == head.get()) {  
                if (first == last) {  
                    if (next == null)  
                        return null;  
                    else  
                        tail.compareAndSet(last, next);  
                } else {  
                    T result = next.item;  
                    if (head.compareAndSet(first, next))  
                        return result;  
                }  
            }  
        }  
    }  
}
```


See `TestMSQueue.java`



Lock-free queue | initialization




```
class MSQueue<T> implements UnboundedQueue<T> {  
    private final AtomicReference<Node<T>> head, tail;  
  
    public MSQueue() {  
        Node<T> dummy = new Node<T>(null, null);  
        head = new AtomicReference<Node<T>>(dummy);  
        tail = new AtomicReference<Node<T>>(dummy);  
    }  
    ...  
}
```



Lock-free queue | enqueue



```
class MSQueue<T> implements UnboundedQueue<T> {  
...  
    public void enqueue(T item) {  
        Node<T> node = new Node<T>(item, null);  
        while (true) {  
            Node<T> last = tail.get();  
            Node<T> next = last.next.get();  
            if (last == tail.get()) {  
                if (next == null) {  
                    // In quiescent state, try inserting new node  
                    if (last.next.compareAndSet(next, node)) {  
                        // Insertion succeeded, try advancing tail  
                        tail.compareAndSet(last, node);  
                        return;  
                    }  
                } else  
                    // Queue in intermediate state, advance tail  
                    tail.compareAndSet(last, next);  
            }  
        }  
    }  
...  
}
```

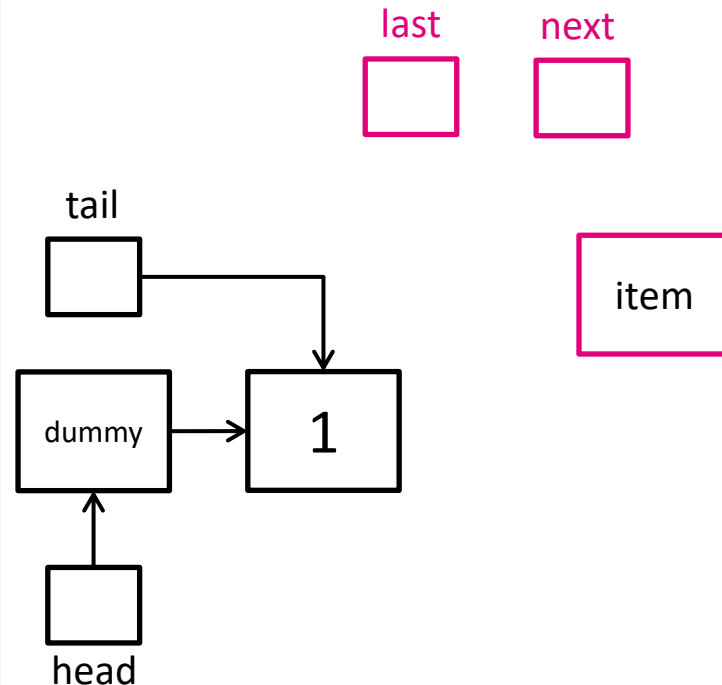


Lock-free queue | enqueue

· 46



```
class MSQueue<T> implements UnboundedQueue<T> {  
...  
    public void enqueue(T item) {  
        Node<T> node = new Node<T>(item, null);  
        while (true) {  
            Node<T> last = tail.get();  
            Node<T> next = last.next.get();  
            if (last == tail.get()) {  
                if (next == null) {  
                    // In quiescent state, try inserting new node  
                    if (last.next.compareAndSet(next, node)) {  
                        // Insertion succeeded, try advancing tail  
                        tail.compareAndSet(last, node);  
                        return;  
                    }  
                } else  
                    // Queue in intermediate state, advance tail  
                    tail.compareAndSet(last, next);  
            }  
        }  
    }  
...  
}
```

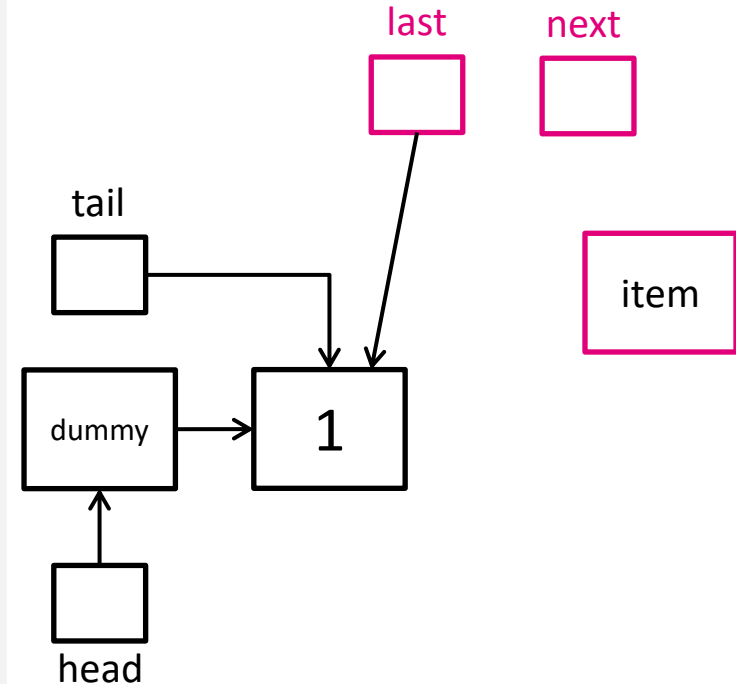


Lock-free queue | enqueue

· 46



```
class MSQueue<T> implements UnboundedQueue<T> {  
...  
    public void enqueue(T item) {  
        Node<T> node = new Node<T>(item, null);  
        while (true) {  
            Node<T> last = tail.get();  
            Node<T> next = last.next.get();  
            if (last == tail.get()) {  
                if (next == null) {  
                    // In quiescent state, try inserting new node  
                    if (last.next.compareAndSet(next, node)) {  
                        // Insertion succeeded, try advancing tail  
                        tail.compareAndSet(last, node);  
                        return;  
                    }  
                } else  
                    // Queue in intermediate state, advance tail  
                    tail.compareAndSet(last, next);  
            }  
        }  
    }  
}
```

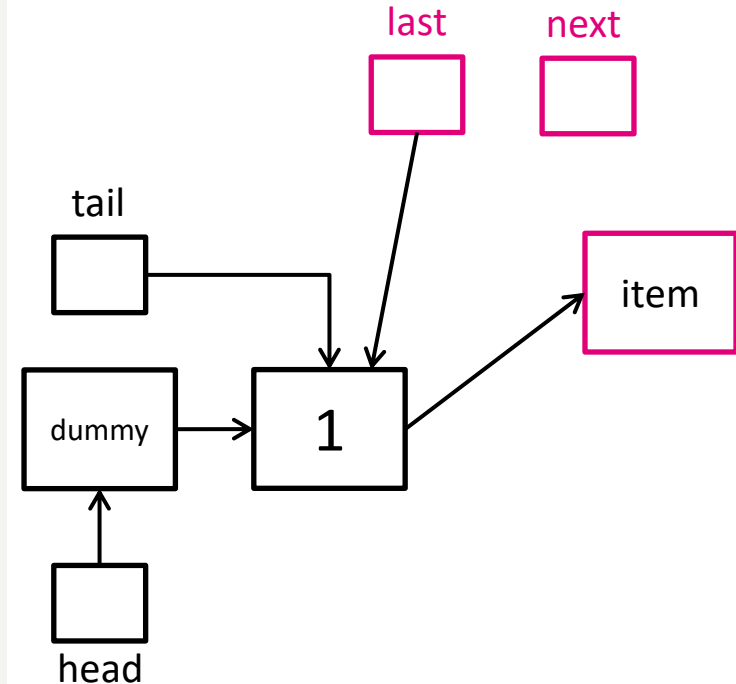


Lock-free queue | enqueue

· 46



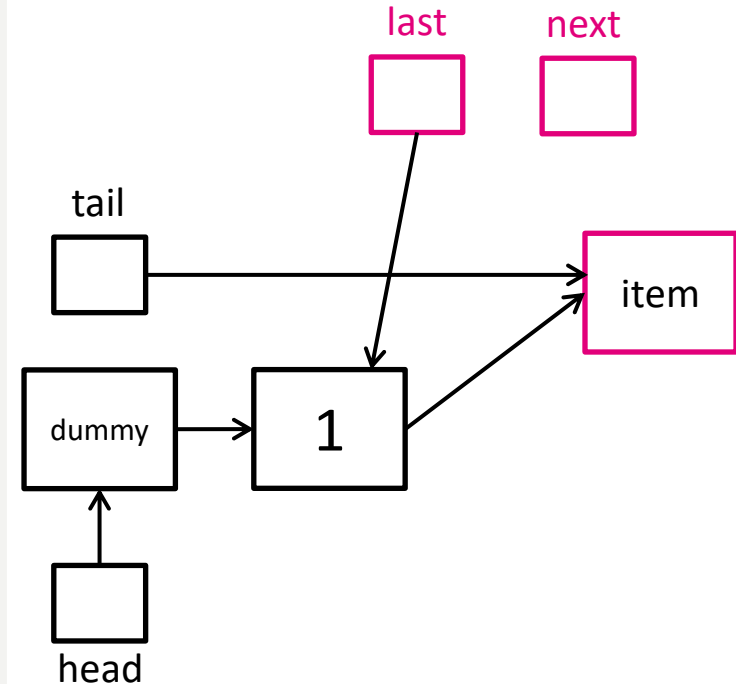
```
class MSQueue<T> implements UnboundedQueue<T> {  
...  
    public void enqueue(T item) {  
        Node<T> node = new Node<T>(item, null);  
        while (true) {  
            Node<T> last = tail.get();  
            Node<T> next = last.next.get();  
            if (last == tail.get()) {  
                if (next == null) {  
                    // In quiescent state, try inserting new node  
                    if (last.next.compareAndSet(next, node)) {  
                        // Insertion succeeded, try advancing tail  
                        tail.compareAndSet(last, node);  
                        return;  
                    }  
                } else  
                    // Queue in intermediate state, advance tail  
                    tail.compareAndSet(last, next);  
            }  
        }  
    }  
}
```



Lock-free queue | enqueue



```
class MSQueue<T> implements UnboundedQueue<T> {  
...  
    public void enqueue(T item) {  
        Node<T> node = new Node<T>(item, null);  
        while (true) {  
            Node<T> last = tail.get();  
            Node<T> next = last.next.get();  
            if (last == tail.get()) {  
                if (next == null) {  
                    // In quiescent state, try inserting new node  
                    if (last.next.compareAndSet(next, node)) {  
                        // Insertion succeeded, try advancing tail  
                        tail.compareAndSet(last, node);  
                        return;  
                    }  
                } else  
                    // Queue in intermediate state, advance tail  
                    tail.compareAndSet(last, next);  
            }  
        }  
    }  
}
```




Lock-free queue | enqueue

· 47



```
class MSQueue<T> implements UnboundedQueue<T> {  
...  
    public void enqueue(T item) {  
        Node<T> node = new Node<T>(item, null);  
        while (true) {  
            Node<T> last = tail.get();  
            Node<T> next = last.next.get();  
            if (last == tail.get()) {  
                if (next == null) {  
                    // In quiescent state, try inserting new node  
                    if (last.next.compareAndSet(next, node)) {  
                        // Insertion succeeded, try advancing tail  
                        tail.compareAndSet(last, node);  
                        return;  
                    }  
                } else {  
                    // Queue in intermediate state, advance tail  
                    tail.compareAndSet(last, next);  
                }  
            }  
        }  
    }  
...  
}
```



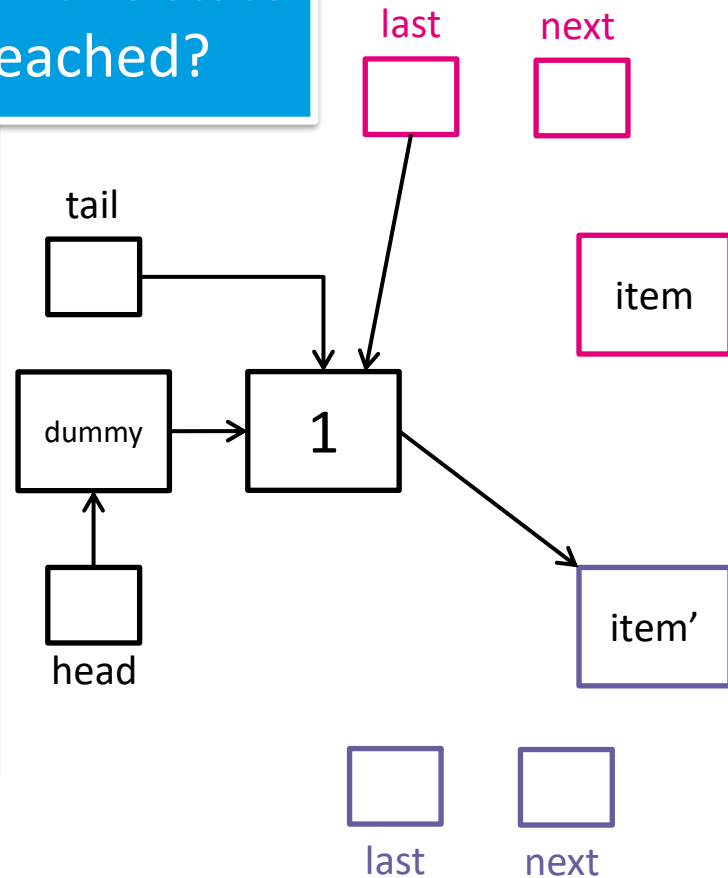
Lock-free queue | enqueue

· 47



How can this state be reached?

```
class MSQueue<T> implements UnboundedQueue<T> {  
...  
    public void enqueue(T item) {  
        Node<T> node = new Node<T>(item, null);  
        while (true) {  
            Node<T> last = tail.get();  
            Node<T> next = last.next.get();  
            if (last == tail.get()) {  
                if (next == null) {  
                    // In quiescent state, try inserting new node  
                    if (last.next.compareAndSet(next, node)) {  
                        // Insertion succeeded, try advancing tail  
                        tail.compareAndSet(last, node);  
                        return;  
                    }  
                } else {  
                    // Queue in intermediate state, advance tail  
                    tail.compareAndSet(last, next);  
                }  
            }  
        }  
    }  
}
```

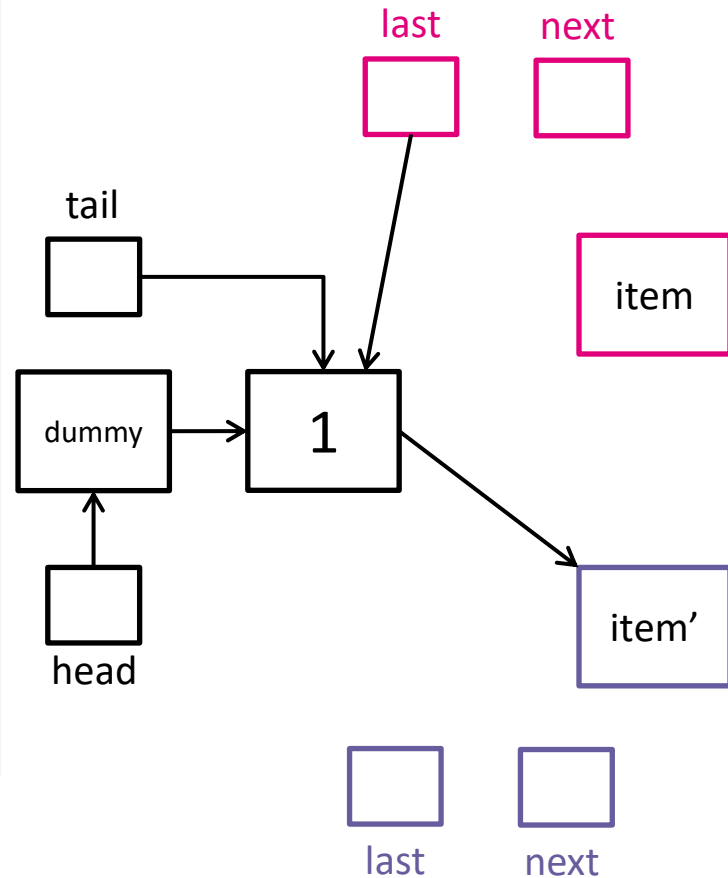


Lock-free queue | enqueue

· 47



```
class MSQueue<T> implements UnboundedQueue<T> {  
...  
    public void enqueue(T item) {  
        Node<T> node = new Node<T>(item, null);  
        while (true) {  
            Node<T> last = tail.get();  
            Node<T> next = last.next.get();  
            if (last == tail.get()) {  
                if (next == null) {  
                    // In quiescent state, try inserting new node  
                    if (last.next.compareAndSet(next, node)) {  
                        // Insertion succeeded, try advancing tail  
                        tail.compareAndSet(last, node);  
                        return;  
                    }  
                } else {  
                    // Queue in intermediate state, advance tail  
                    tail.compareAndSet(last, next);  
                }  
            }  
        }  
    }  
}
```

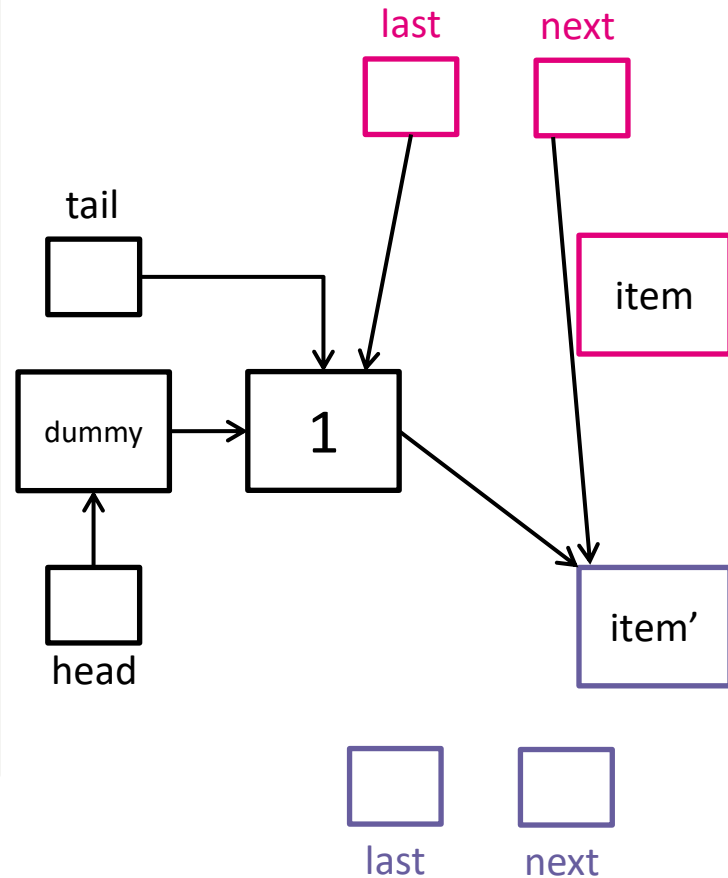


Lock-free queue | enqueue

· 47



```
class MSQueue<T> implements UnboundedQueue<T> {  
...  
public void enqueue(T item) {  
    Node<T> node = new Node<T>(item, null);  
    while (true) {  
        Node<T> last = tail.get();  
        Node<T> next = last.next.get();  
        if (last == tail.get()) {  
            if (next == null) {  
                // In quiescent state, try inserting new node  
                if (last.next.compareAndSet(next, node)) {  
                    // Insertion succeeded, try advancing tail  
                    tail.compareAndSet(last, node);  
                    return;  
                }  
            } else {  
                // Queue in intermediate state, advance tail  
                tail.compareAndSet(last, next);  
            }  
        }  
    }  
}  
...  
}
```

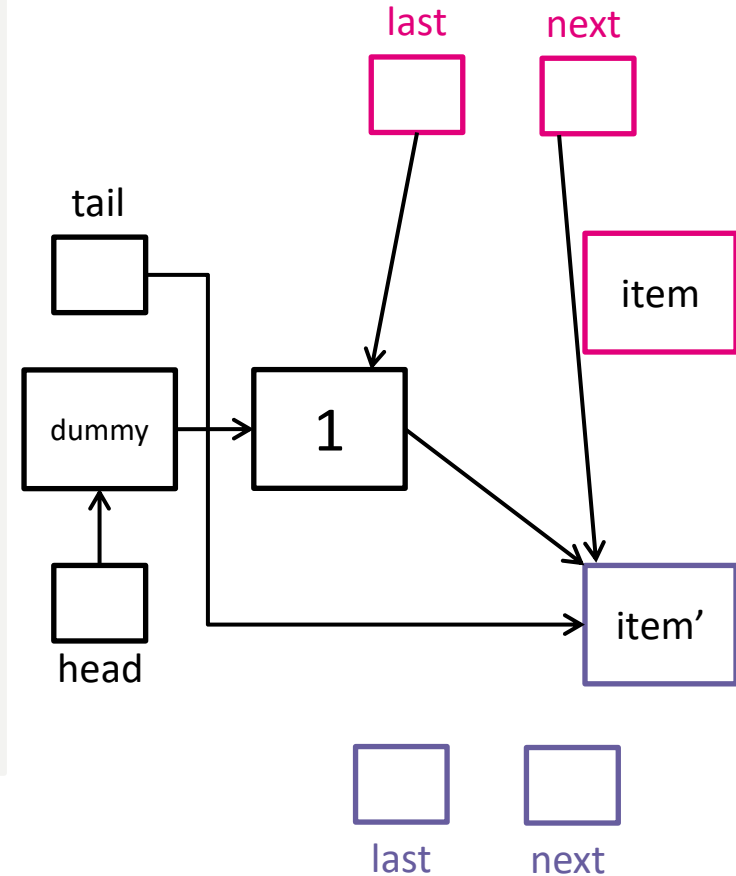


Lock-free queue | enqueue

· 47



```
class MSQueue<T> implements UnboundedQueue<T> {  
...  
    public void enqueue(T item) {  
        Node<T> node = new Node<T>(item, null);  
        while (true) {  
            Node<T> last = tail.get();  
            Node<T> next = last.next.get();  
            if (last == tail.get()) {  
                if (next == null) {  
                    // In quiescent state, try inserting new node  
                    if (last.next.compareAndSet(next, node)) {  
                        // Insertion succeeded, try advancing tail  
                        tail.compareAndSet(last, node);  
                        return;  
                    }  
                } else {  
                    // Queue in intermediate state, advance tail  
                    tail.compareAndSet(last, next);  
                }  
            }  
        }  
    }  
}
```



• 47



• 47




Lock-free queue | enqueue

· 47



```
class MSQueue<T> implements UnboundedQueue<T> {  
...  
public void enqueue(T item) {  
    Node<T> node = new Node<T>(item, null);  
    while (true) {  
        Node<T> last = tail.get();  
        Node<T> next = last.next.get();  
        if (last == tail.get()) {  
            if (next == null) {  
                // In quiescent state, try inserting new node  
                if (last.next.compareAndSet(next, node)) {  
                    // Insertion succeeded, try advancing tail  
                    tail.compareAndSet(last, node);  
                    return;  
                }  
            } else {  
                // Queue in intermediate state, advance tail  
                tail.compareAndSet(last, next);  
            }  
        }  
    }  
}
```

In case another thread is enqueueing, and didn't update the tail, the current thread helps by advancing the tail




Lock-free queue | enqueue

· 48



```
class MSQueue<T> implements UnboundedQueue<T> {  
...  
    public void enqueue(T item) {  
        Node<T> node = new Node<T>(item, null);  
        while (true) {  
            Node<T> last = tail.get();  
            Node<T> next = last.next.get();  
            if (last == tail.get()) {  
                if (next == null) {  
                    // In quiescent state, try inserting new node  
                    if (last.next.compareAndSet(next, node)) {  
                        // Insertion succeeded, try advancing tail  
                        tail.compareAndSet(last, node);  
                        return;  
                    }  
                } else  
                    // Queue in intermediate state, advance tail  
                    tail.compareAndSet(last, next);  
            }  
        }  
    }  
}
```

If the tail has been changed, then the thread restarts right away



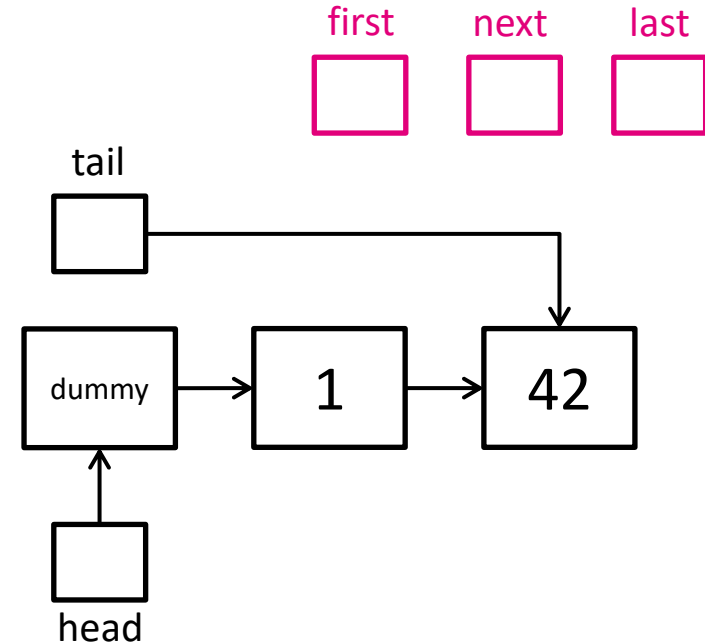
Lock-free queue | dequeue

· 49



```
class MSQueue<T> implements UnboundedQueue<T> {
...

    public T dequeue() {
        while (true) {
            Node<T> first = head.get();
            Node<T> last = tail.get();
            Node<T> next = first.next.get();
            if (first == head.get()) {
                if (first == last) {
                    if (next == null)
                        return null;
                    else
                        tail.compareAndSet(last, next);
                } else {
                    T result = next.item;
                    if (head.compareAndSet(first, next))
                        return result;
                }
            }
        }
    }
...
}
```




Lock-free queue | dequeue

· 49



```
class MSQueue<T> implements UnboundedQueue<T> {
...

    public T dequeue() {
        while (true) {
            Node<T> first = head.get();
            Node<T> last = tail.get();
            Node<T> next = first.next.get();
            if (first == head.get()) {
                if (first == last) {
                    if (next == null)
                        return null;
                    else
                        tail.compareAndSet(last, next);
                } else {
                    T result = next.item;
                    if (head.compareAndSet(first, next))
                        return result;
                }
            }
        }
    }
...
}
```

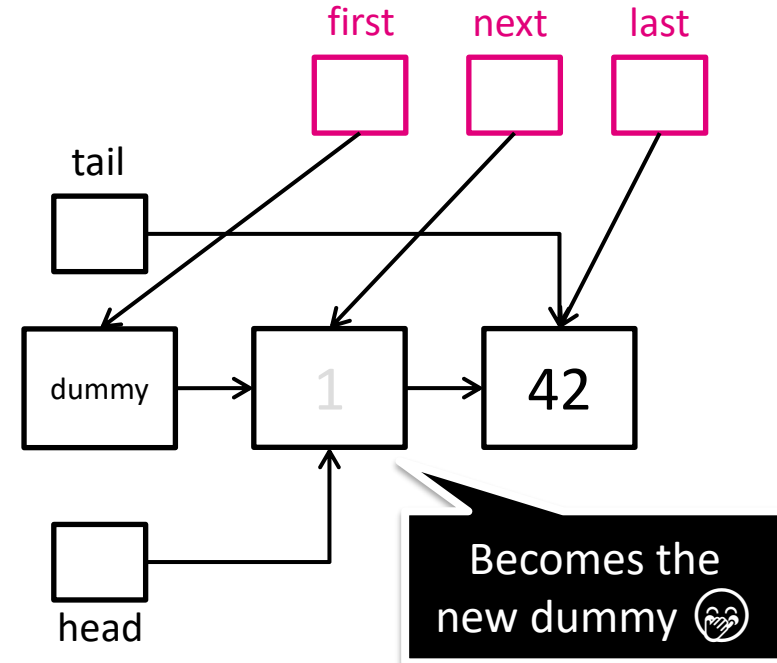


Lock-free queue | dequeue

· 49



```
class MSQueue<T> implements UnboundedQueue<T> {  
    ...  
  
    public T dequeue() {  
        while (true) {  
            Node<T> first = head.get();  
            Node<T> last = tail.get();  
            Node<T> next = first.next.get();  
            if (first == head.get()) {  
                if (first == last) {  
                    if (next == null)  
                        return null;  
                    else  
                        tail.compareAndSet(last, next);  
                } else {  
                    T result = next.item;  
                    if (head.compareAndSet(first, next))  
                        return result;  
                }  
            }  
        }  
    }  
    ...  
}
```




Lock-free queue | dequeue

· 50



```
class MSQueue<T> implements UnboundedQueue<T> {
...

    public T dequeue() {
        while (true) {
            Node<T> first = head.get();
            Node<T> last = tail.get();
            Node<T> next = first.next.get();
            if (first == head.get()) {
                if (first == last) {
                    if (next == null)
                        return null;
                    else
                        tail.compareAndSet(last, next);
                } else {
                    T result = next.item;
                    if (head.compareAndSet(first, next))
                        return result;
                }
            }
        }
    }
...
}
```

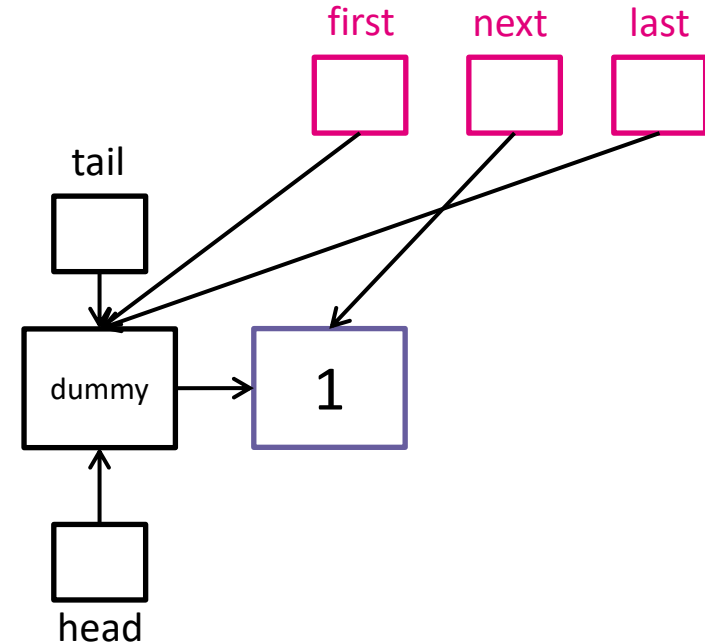


Lock-free queue | dequeue

· 50



```
class MSQueue<T> implements UnboundedQueue<T> {  
    ...  
  
    public T dequeue() {  
        while (true) {  
            Node<T> first = head.get();  
            Node<T> last = tail.get();  
            Node<T> next = first.next.get();  
            if (first == head.get()) {  
                if (first == last) {  
                    if (next == null)  
                        return null;  
                    else  
                        tail.compareAndSet(last, next);  
                } else {  
                    T result = next.item;  
                    if (head.compareAndSet(first, next))  
                        return result;  
                }  
            }  
        }  
    }  
    ...  
}
```



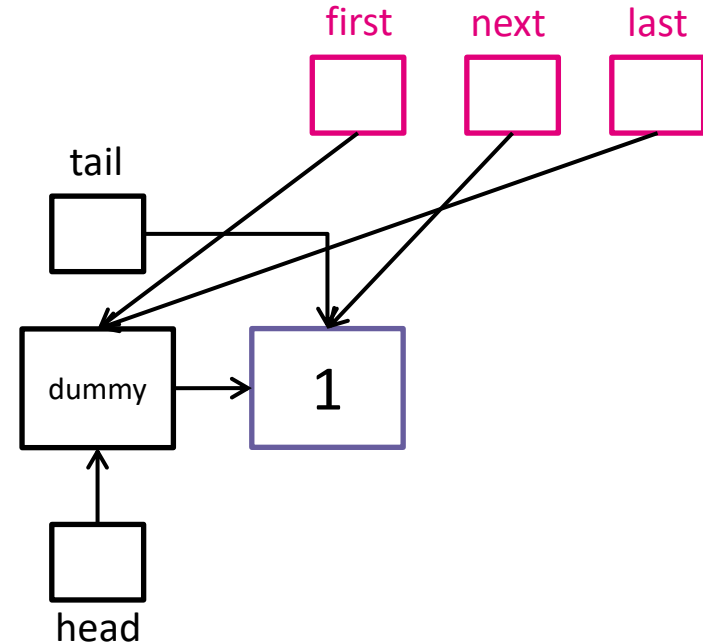
Lock-free queue | dequeue

· 50



```
class MSQueue<T> implements UnboundedQueue<T> {  
    ...  
  
    public T dequeue() {  
        while (true) {  
            Node<T> first = head.get();  
            Node<T> last = tail.get();  
            Node<T> next = first.next.get();  
            if (first == head.get()) {  
                if (first == last) {  
                    if (next == null)  
                        return null;  
                    else  
                        tail.compareAndSet(last, next);  
                } else {  
                    T result = next.item;  
                    if (head.compareAndSet(first, next))  
                        return result;  
                }  
            }  
        }  
    }  
}
```

If the next field of the head is not null (because another thread push in the meantime), then the calling thread helps advancing the tail and tries to pop again.

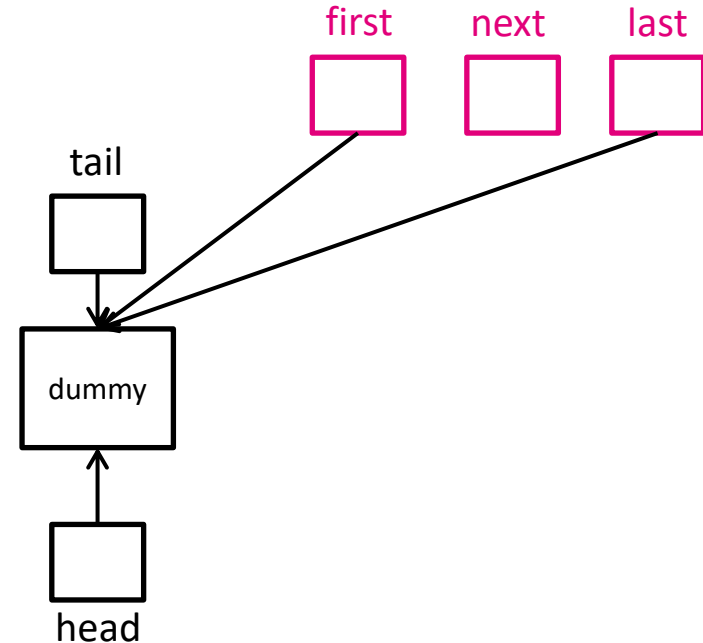


Lock-free queue | dequeue

· 51



```
class MSQueue<T> implements UnboundedQueue<T> {  
    ...  
  
    public T dequeue() {  
        while (true) {  
            Node<T> first = head.get();  
            Node<T> last = tail.get();  
            Node<T> next = first.next.get();  
            if (first == head.get()) {  
                if (first == last) {  
                    if (next == null)  
                        return null;  
                    else  
                        tail.compareAndSet(last, next);  
                } else {  
                    T result = next.item;  
                    if (head.compareAndSet(first, next))  
                        return result;  
                }  
            }  
        }  
    }  
    ...  
}
```



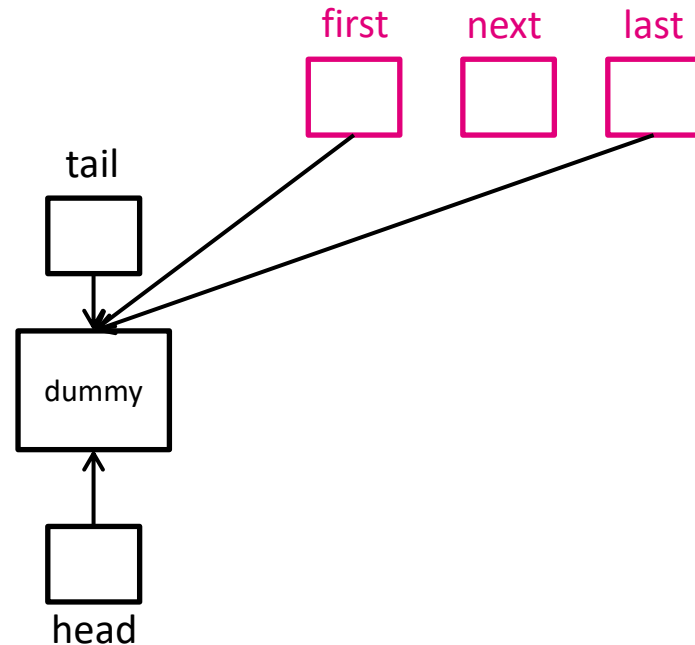
Lock-free queue | dequeue

· 51



```
class MSQueue<T> implements UnboundedQueue<T> {  
    ...  
  
    public T dequeue() {  
        while (true) {  
            Node<T> first = head.get();  
            Node<T> last = tail.get();  
            Node<T> next = first.next.get();  
            if (first == head.get()) {  
                if (first == last) {  
                    if (next == null)  
                        return null;  
                    else  
                        tail.compareAndSet(last, next);  
                } else {  
                    T result = next.item;  
                    if (head.compareAndSet(first, next))  
                        return result;  
                }  
            }  
        }  
    }  
}
```

If the next field of the head is pointing to null, then we return null. Why is this correct?




Lock-free queue | dequeue

· 52



```
class MSQueue<T> implements UnboundedQueue<T> {  
    ...  
  
    public T dequeue() {  
        while (true) {  
            Node<T> first = head.get();  
            Node<T> last = tail.get();  
            Node<T> next = first.next.get();  
            if (first == head.get()) {  
                if (first == last) {  
                    if (next == null)  
                        return null;  
                    else  
                        tail.compareAndSet(last, next);  
                } else {  
                    T result = next.item;  
                    if (head.compareAndSet(first, next))  
                        return result;  
                }  
            }  
        }  
    }  
}
```

How can a state where $\text{first} \neq \text{head.get()}$ be reached?



What about correctness?



- We have seen several implementations of lock-free data structures today
- However, we have not seen any techniques to reason about their correctness
 - In two weeks -> Linearizability

No class next week

· 54



Enjoy a well-deserved Autumn break!

- Compare-And-Swap (CAS)
 - Lock-free atomic integer
 - Lock-free number range
 - Atomic libraries
 - CAS based lock implementation
- Lock-free stack
- ABA problem
- Lock-free queue