

#### **Practical Concurrent and Parallel Programming XII**

#### Message Passing II

Raúl Pardo

## Agenda

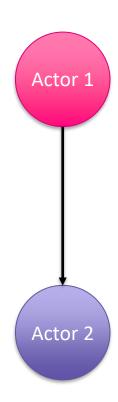


- Actors model (revisited)
  - Primer
- Dynamic topology
- Fault-tolerance
  - Supervision
- Adaptive load balancing
  - Scatter-Gather
- Changing behaviour

# What is an Actor? (Bird's eye, revisited)



- An actor can be seen as a sequential unit of computation
  - Although, formally, the model allows for parallelism within the actor, one can safely assume that there are not concurrency issues within the actor.
  - You can think of an actor as a thread
- Actors can send messages to other actors



# Actor – Specification (revisited)



- An actor is an abstraction of a thread (intuitively)
- An actors can only execute any of these 4 actions
  - 1. Receive messages from other actors
  - 2. <u>Send asynchronous messages</u> to other actors
  - 3. Create new actors
  - 4. Change its behaviour (local state and/or message handlers)
- Actors <u>do not share memory</u>
  - They only have access to:
    - Their *local state* (local memory)
    - Their mailbox (multiset of fixed size with received messages)
    - By default, the mailbox is of unbounded size



#### Akka actors

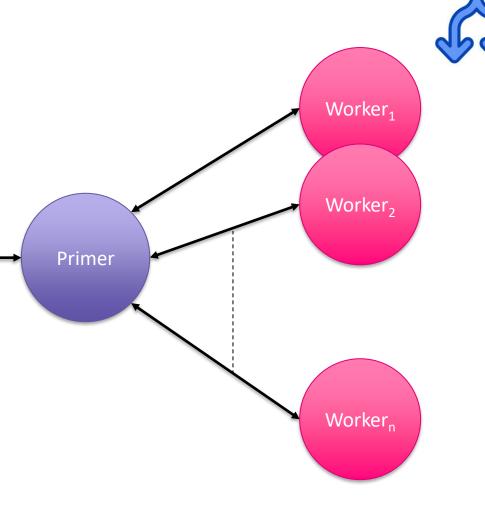


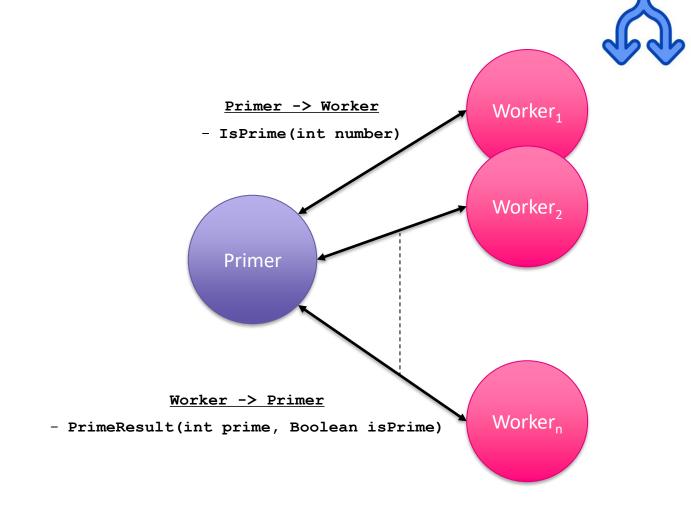
 There is a one-to-one correspondence of the basic actor operations and the Akka API

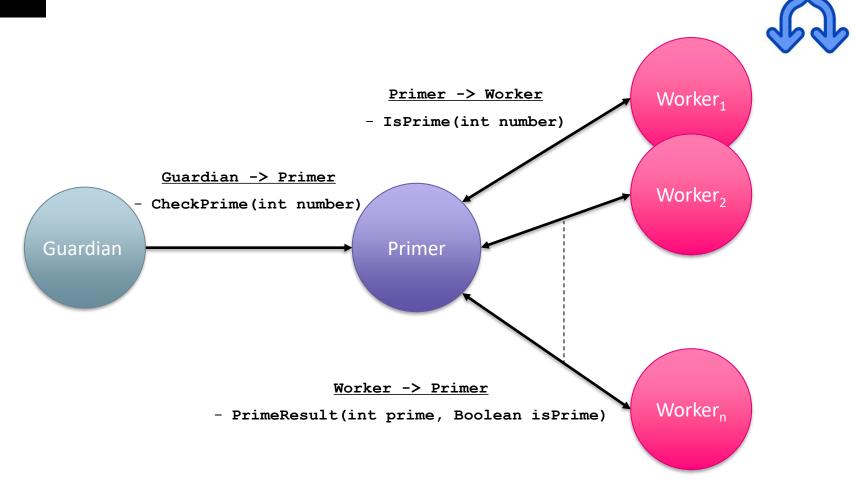
Actors Model	Akka
Actor	Actor class (AbstractBehaviour)
Mailbox Address	Reference to Actor class
Message	Message static final class
State	Actor class local attributes
Behaviour	Handler functions in the Actor class
Create actor	API function
Send message	API function
Receive message	Message handler builder (from API)

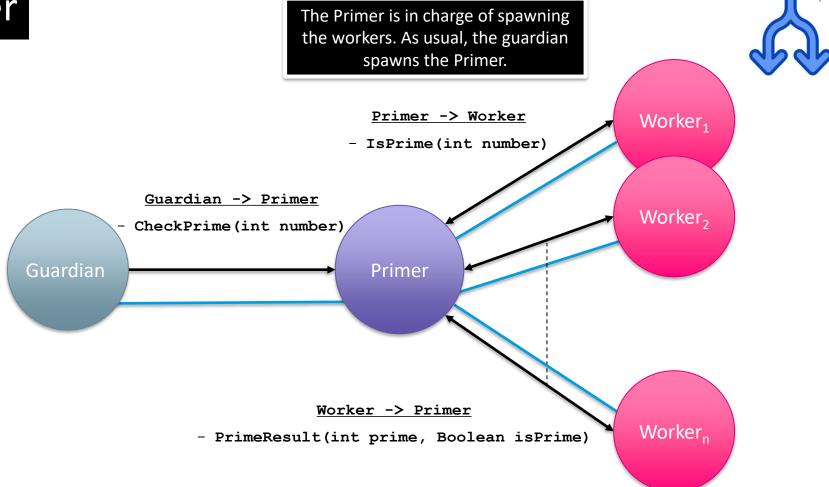
 Consider a Primer actor that receives numbers and checks whether they are prime

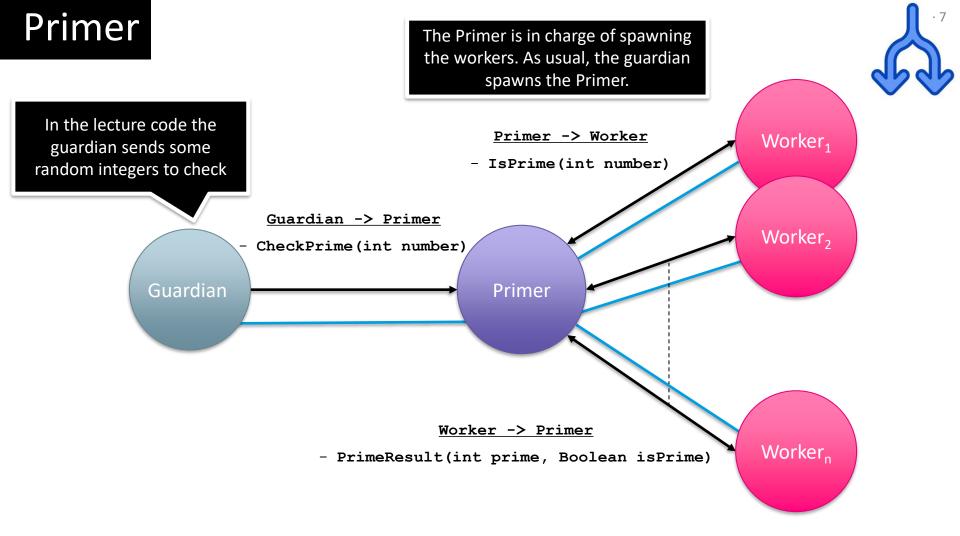
 The actor uses a (fixed) set of worker actors to which it forwards the numbers so that several primes are checked in parallel

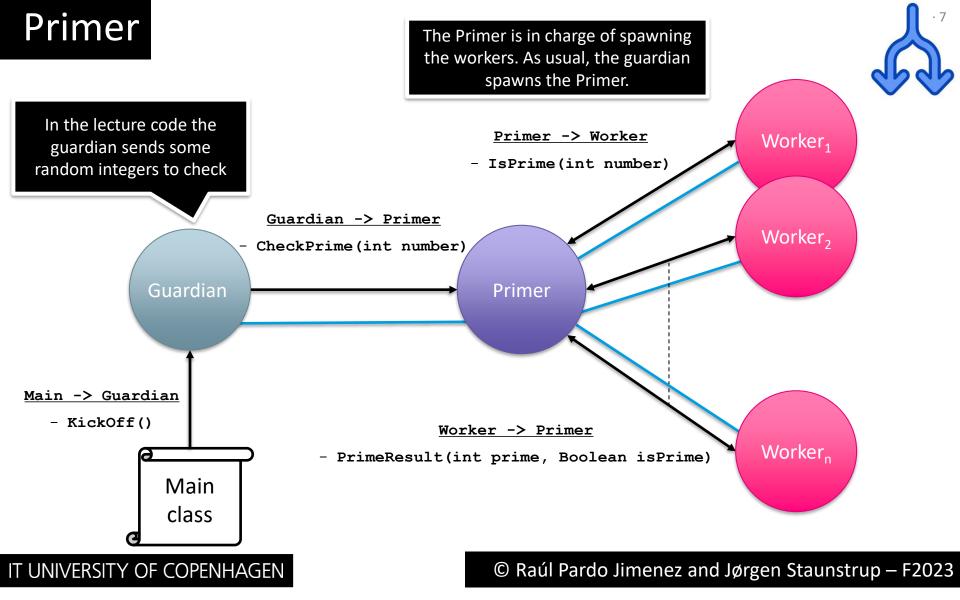




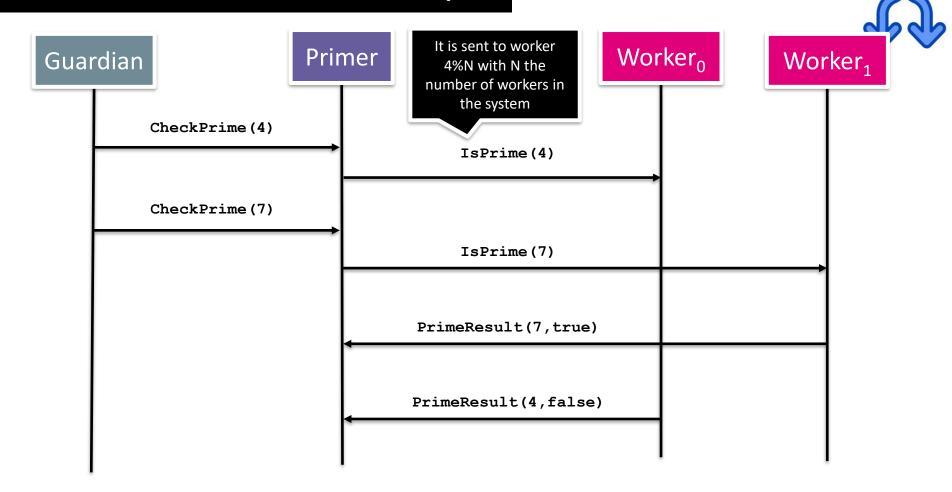




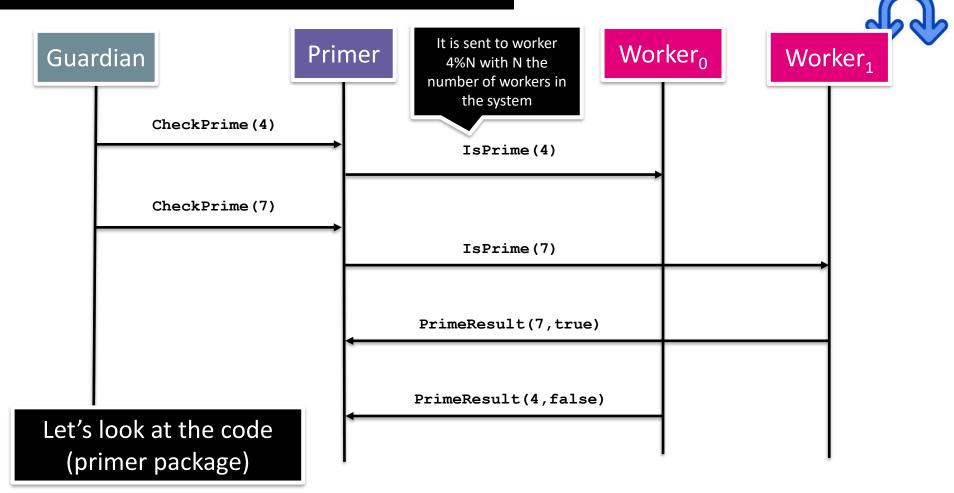




### Primer – execution example







#### Primer – Printing order



 Note that the printing order of the results does not correspond to the order of sending the requests

```
> Task :app:run
[primer_system-akka.actor.default-dispatcher-3] INFO akka.event.slf4j.Slf4jLogger - Slf4jLogger started
>>> Press ENTER to exit <<<
[primer_system-akka.actor.default-dispatcher-3] INFO primer.Primer - primer_server: Server and workers started
[primer_system-akka.actor.default-dispatcher-3] INFO primer.Primer - primer_server: Cheking whether 21098598 is prime by worker worker_19
[primer_system-akka.actor.default-dispatcher-3] INFO primer.Primer - primer_server: Cheking whether 1001439 is prime by worker worker_20
[primer_system-akka.actor.default-dispatcher-3] INFO primer.Primer - primer_server: Cheking whether 47257026 is prime by worker worker_7
[primer_system-akka.actor.default-dispatcher-3] INFO primer.Primer - primer_server: Cheking whether 40857223 is prime by worker worker_4
[primer_system-akka.actor.default-dispatcher-3] INFO primer.Primer - primer_server: Cheking whether 10667083 is prime by worker worker_4
[primer_system-akka.actor.default-dispatcher-3] INFO primer.Primer - primer_server: Number 1001439 is not prime. [1/5]
[primer_system-akka.actor.default-dispatcher-3] INFO primer.Primer - primer_server: Number 47257026 is not prime. [2/5]
[primer_system-akka.actor.default-dispatcher-3] INFO primer.Primer - primer_server: Number 47257026 is not prime. [3/5]
[primer_system-akka.actor.default-dispatcher-3] INFO primer.Primer - primer_server: Number 40857223 is not prime. [4/5]
[primer_system-akka.actor.default-dispatcher-3] INFO primer.Primer - primer_server: Number 40857223 is not prime. [5/5]
```

#### Primer – Printing order



 Note that the printing order of the results does not correspond to the order of sending the requests

```
> Task :app:run
[primer_system-akka.actor.default-dispatcher-3] INFO akka.event.slf4j.Slf4jLogger - Slf4jLogger started

>>> Press ENTER to exit <<

[primer_system-akka.actor.default-dispatcher-3] INFO primer.Primer - primer_server: Server and workers started

[primer_system-akka.actor.default-dispatcher-3] INFO primer.Primer - primer_server: Cheking whether 21098598 is prime by worker worker_19

[primer_system-akka.actor.default-dispatcher-3] INFO primer.Primer - primer_server: Cheking whether 1001439 is prime by worker worker_20

[primer_system-akka.actor.default-dispatcher-3] INFO primer.Primer - primer_server: Cheking whether 47257026 is prime by worker worker_7

[primer_system-akka.actor.default-dispatcher-3] INFO primer.Primer - primer_server: Cheking whether 40857223 is prime by worker worker_4

[primer_system-akka.actor.default-dispatcher-3] INFO primer.Primer - primer_server: Cheking whether 10667083 is prime by worker worker_4

[primer_system-akka.actor.default-dispatcher-3] INFO primer.Primer - primer_server: Number 1001439 is not prime. [2/5]

[primer_system-akka.actor.default-dispatcher-3] INFO primer.Primer - primer_server: Number 21098598 is not prime. [2/5]

[primer_system-akka.actor.default-dispatcher-3] INFO primer.Primer - primer_server: Number 47257026 is not prime. [2/5]

[primer_system-akka.actor.default-dispatcher-3] INFO primer.Primer - primer_server: Number 47257026 is not prime. [2/5]

[primer_system-akka.actor.default-dispatcher-3] INFO primer.Primer - primer_server: Number 47257026 is not prime. [3/5]

[primer_system-akka.actor.default-dispatcher-3] INFO primer.Primer - primer_server: Number 40857223 is not prime. [4/5]

[primer_system-akka.actor.default-dispatcher-3] INFO primer.Primer - primer_server: Number 40857223 is not prime. [5/5]
```

How can this ordering happen?

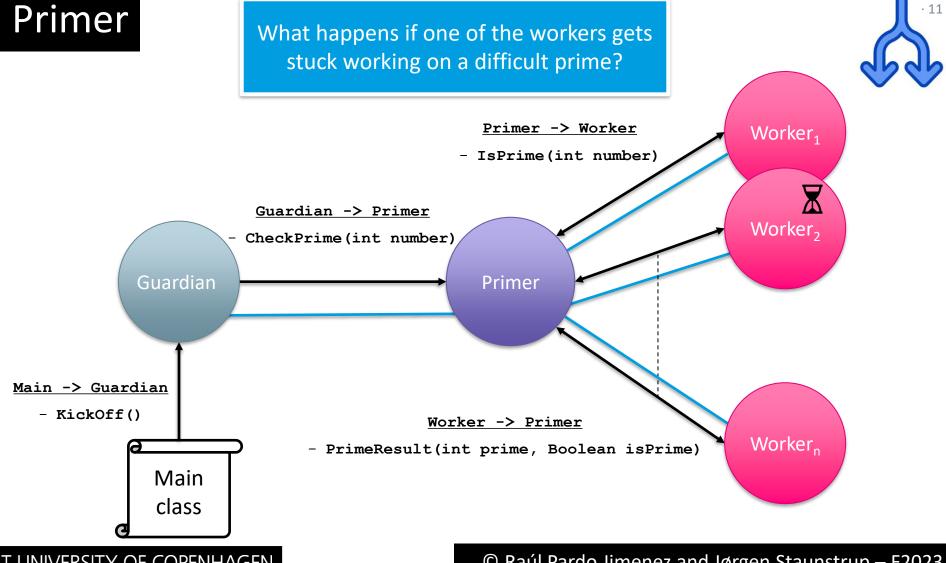
#### Primer – Printing order



 Note that the printing order of the results does not correspond to the order of sending the requests

```
> Task :app:run
[primer_system-akka.actor.default-dispatcher-3] INFO akka.event.slf4j.Slf4jLogger - Slf4jLogger started
>>> Press ENTER to exit <<
[primer_system-akka.actor.default-dispatcher-3] INFO primer.Primer - primer_server: Server and workers started
[primer_system-akka.actor.default-dispatcher-3] INFO primer.Primer - primer_server: Cheking whether 21098598 is prime by worker worker_19
[primer_system-akka.actor.default-dispatcher-3] INFO primer.Primer - primer_server: Cheking whether 1001439 is prime by worker worker_20
[primer_system-akka.actor.default-dispatcher-3] INFO primer.Primer - primer_server: Cheking whether 47257026 is prime by worker worker_7
[primer_system-akka.actor.default-dispatcher-3] INFO primer.Primer - primer_server: Cheking whether 40857223 is prime by worker worker_4
[primer_system-akka.actor.default-dispatcher-3] INFO primer.Primer - primer_server: Cheking whether 10667083 is prime by worker worker_4
[primer_system-akka.actor.default-dispatcher-3] INFO primer.Primer - primer_server: Number 1001439 is not prime. [1/5]
[primer_system-akka.actor.default-dispatcher-3] INFO primer.Primer - primer_server: Number 21098598 is not prime. [2/5]
[primer_system-akka.actor.default-dispatcher-3] INFO primer.Primer - primer_server: Number 47257026 is not prime. [2/5]
[primer_system-akka.actor.default-dispatcher-3] INFO primer.Primer - primer_server: Number 47257026 is not prime. [3/5]
[primer_system-akka.actor.default-dispatcher-3] INFO primer.Primer - primer_server: Number 40857223 is not prime. [4/5]
[primer_system-akka.actor.default-dispatcher-3] INFO primer.Primer - primer_server: Number 40857223 is not prime. [5/5]
```

How would you change the system to print the results in the same order as they arrived?

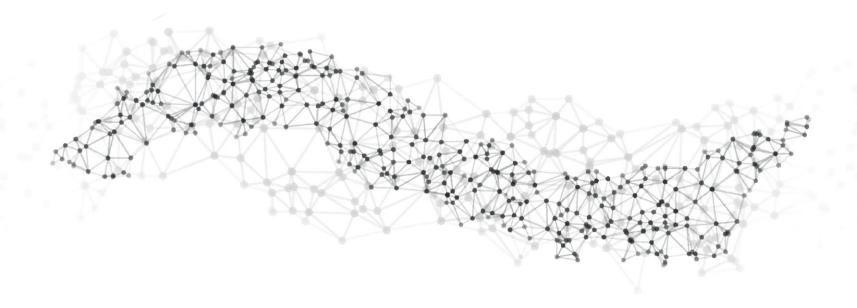




# Actors systems with <u>dynamic</u> topology



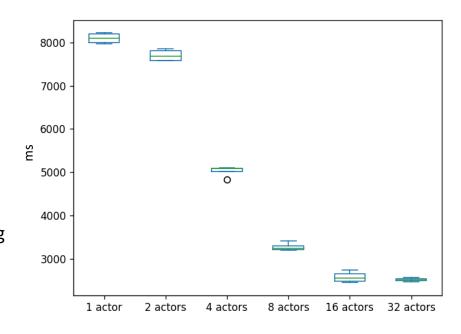
 The Actors model encourages creating many actors that perform small tasks and communicate with each other



## Do more actors improve performance?



- As usual, performance depends on the hardware
- These are the results of running the primer to check 1 million numbers between 1 billion and Integer.MAX\_VALUE
  - Not very strong statistics (4 runs for each number of actors)
- Akka implements actors systems using ForkJoinPools
- However, actor systems can be distributed among many JVMs and computers
  - We are not limited to a single computer throughput
  - See Akka cluster

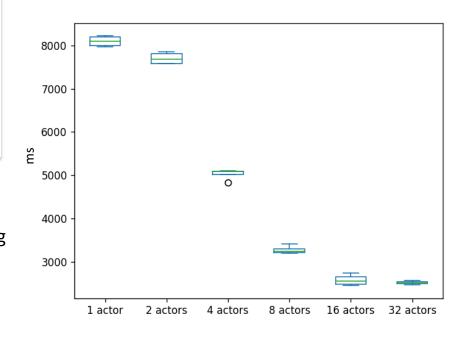


## Do more actors improve performance?



That said, distributing computation among actors makes it easy to implement fault-tolerant systems and adaptive load-balancing

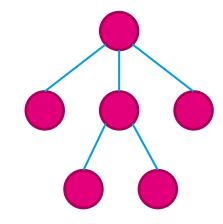
- Akka implements actors systems using ForkJoinPools
- However, actor systems can be distributed among many JVMs and computers
  - We are not limited to a single computer throughput
  - See Akka cluster



# Topology (Actor hierarchy)

•15

- We use the term topology to refer to the parental structure of actors in the system
  - In Akka, this structure is a tree, and it is called a hierarchy



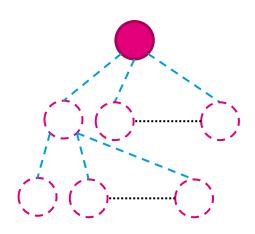
- The systems we have seen so far feature a static topology
  - All the actors in the system are spawned during initialization

Solid lines and actors represent elements that are created during initialization and never change

## Dynamic topology

•16

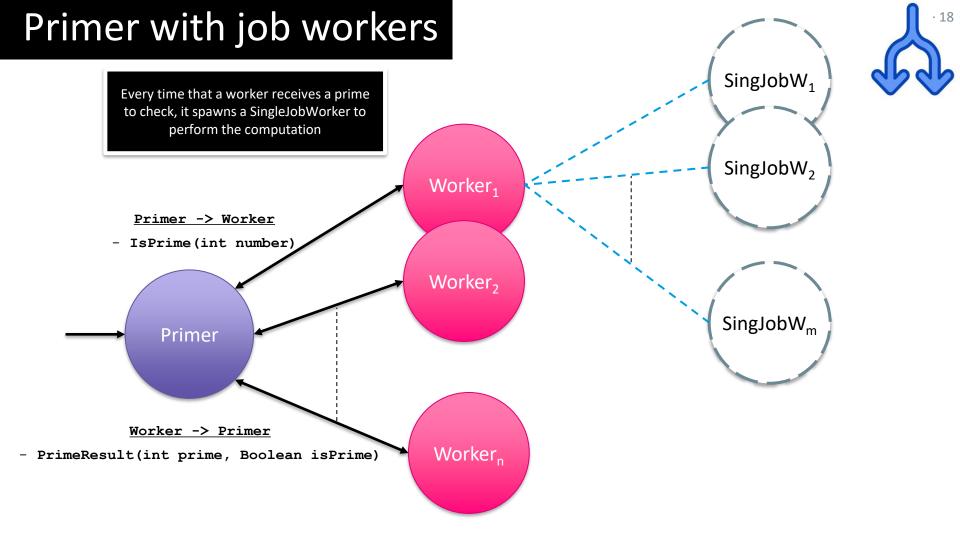
- Actor systems with static topology may not exploit computational resources effectively
  - As we saw, the system may slow down if some actors are consuming excessive computational resources
  - Actors may also crash, and the system should be able to recover from this (faulttolerance)
- The advantages of the actors model are better exploited when the system can adaptively decide the number of workers
- Actors should be seen as nice co-workers
  - A group of computational resources that collaborate to achieve a common goal

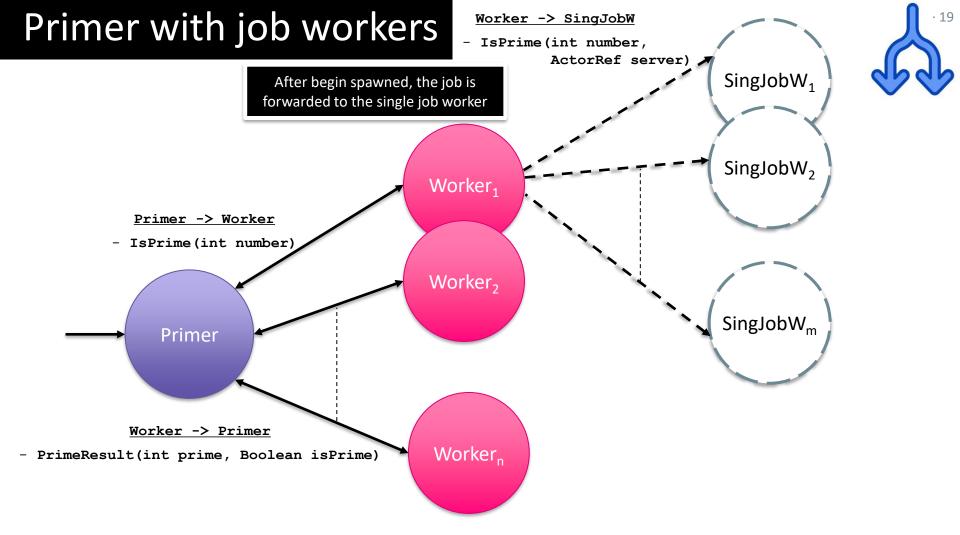


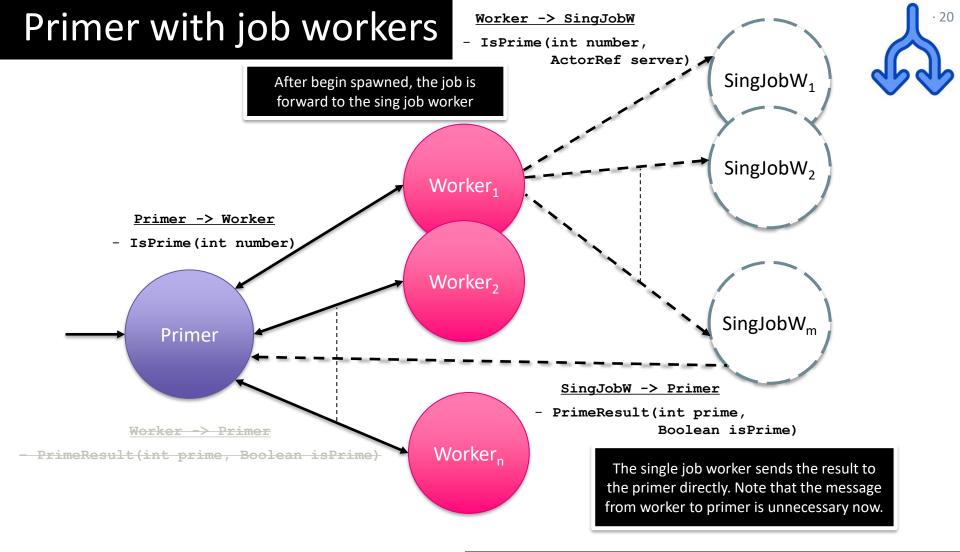
Dashed lines and actors represent elements that may be created dynamically (on-demand, after initialization)

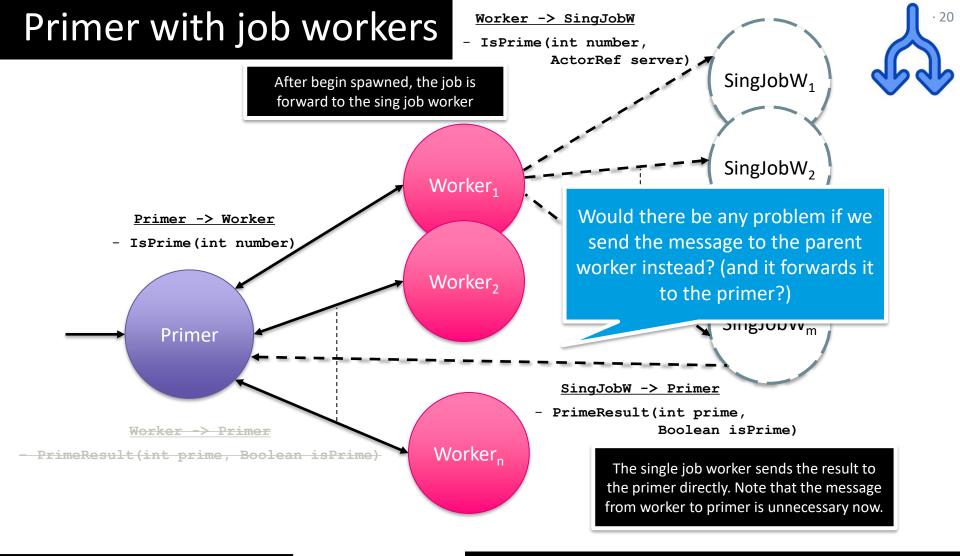


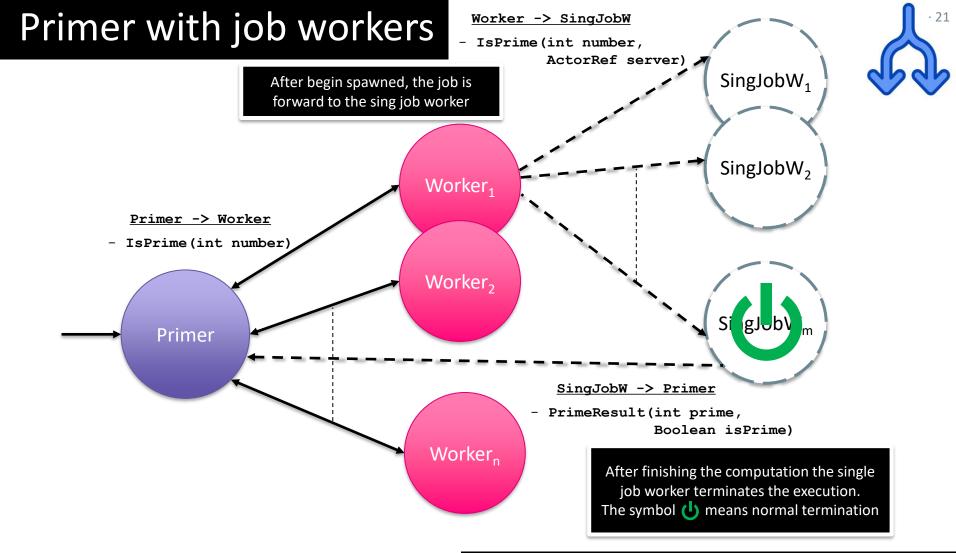
- To avoid excessive delays by primes that are difficult to check, we extend the system with dedicated actors whose only task is to check the prime
- After these dedicated actors have finished the computation they report the result and terminate the execution







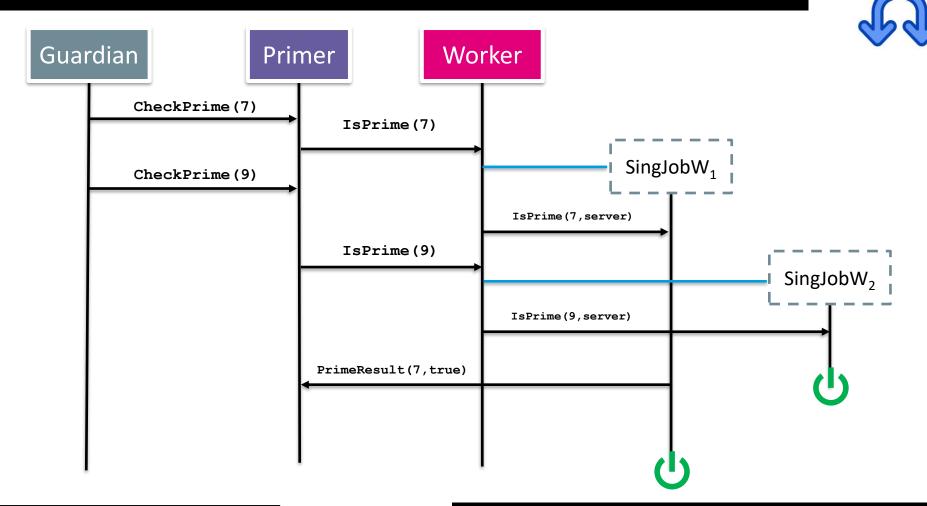


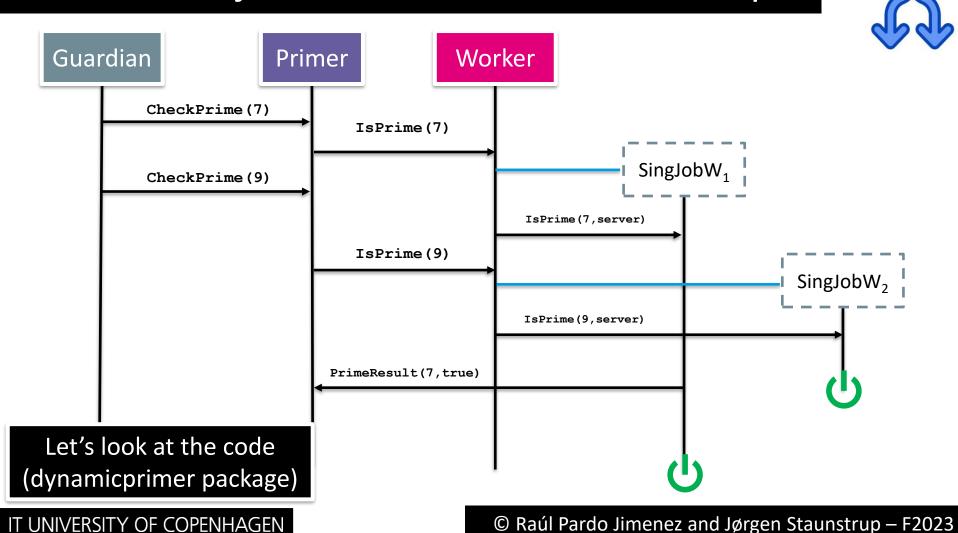


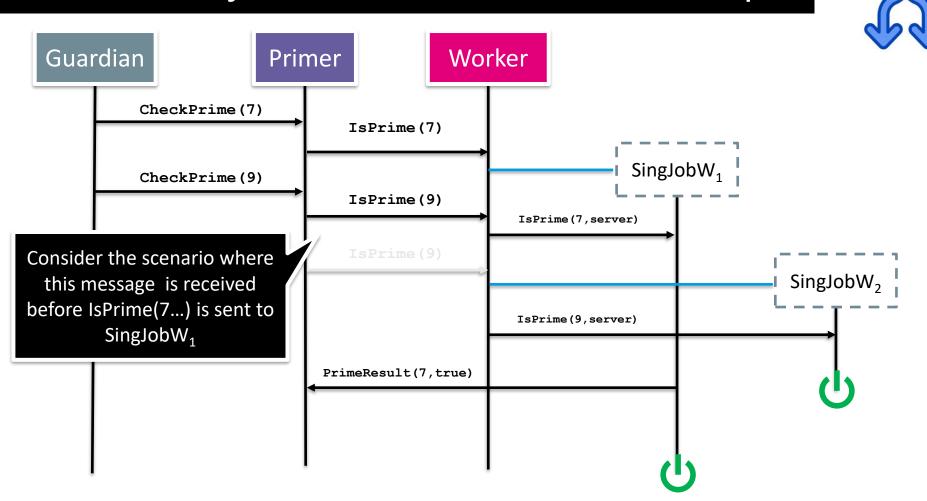


- The Akka API has a specific method to shutdown an actor
  - Behaviours.stopped()
- This behaviour can be used as the return object of a message handler

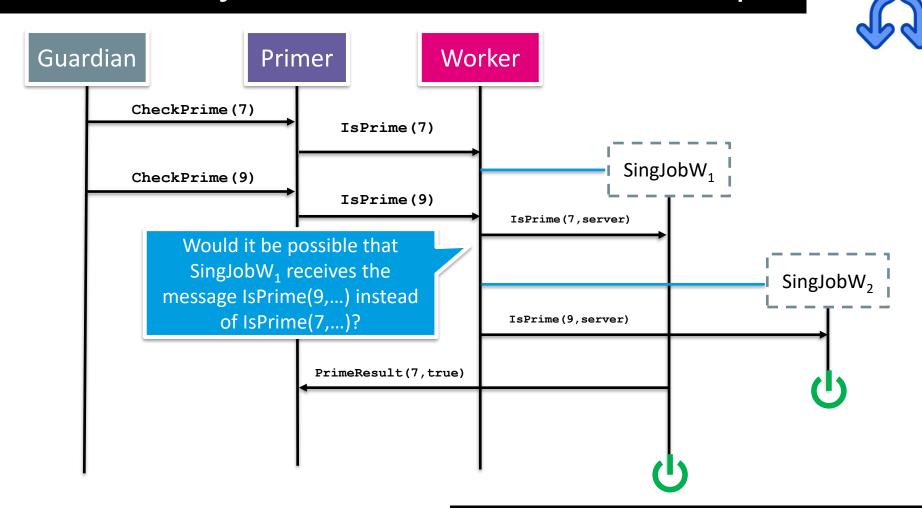
```
public Behavior<IsPrime> onIsPrime(IsPrime msg) {
    msg.server.tell(new Primer.PrimeResult(msg.number, isPrime(msg.number)));
    return Behaviors.stopped();
}
```

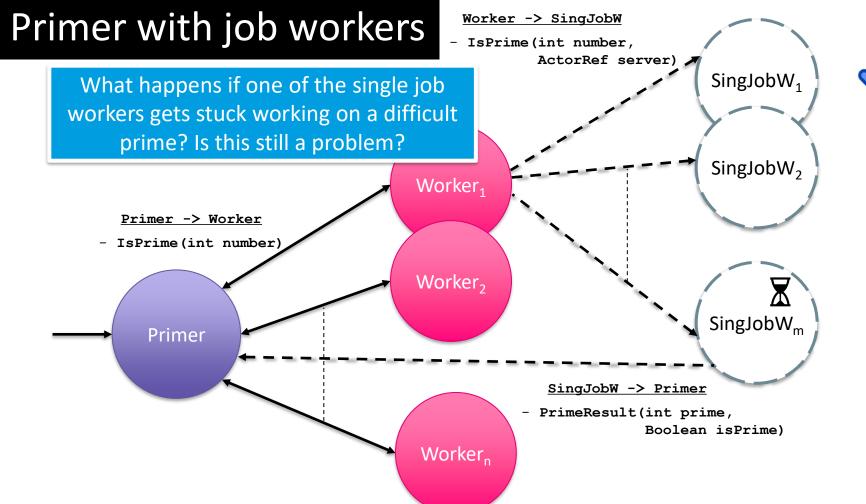


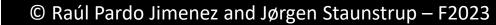




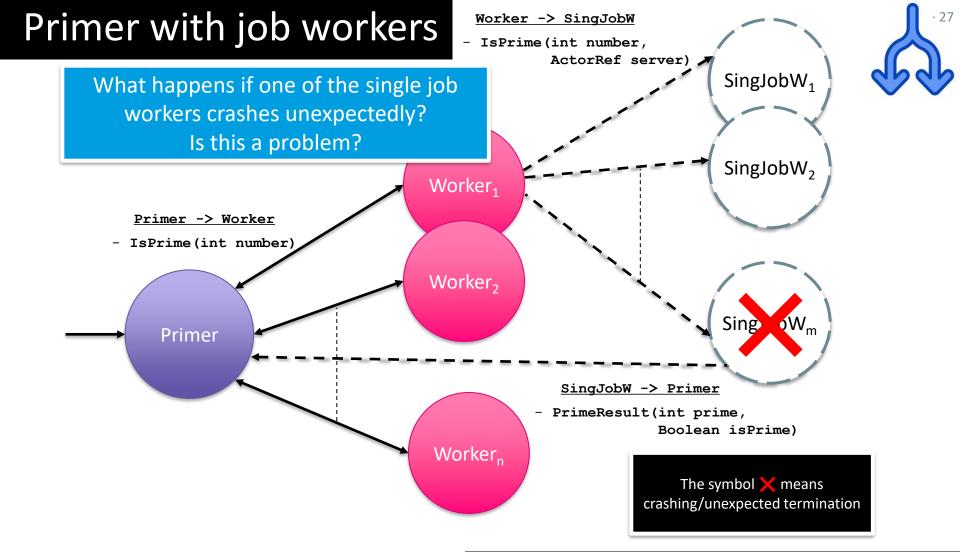
. 24







. 26

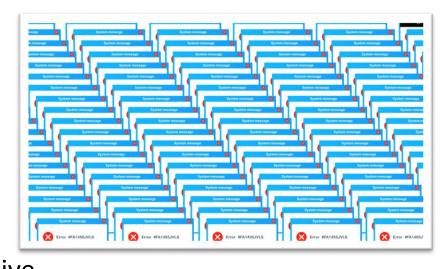




# Fault-tolerance in Akka



- Actor libraries and programming languages encourage a <u>let it crash</u> programming model
- Do not put a lot of effort ensuring that actors never crash
  - Assume that things will fail
- Develop actors systems ensuring that if an actors crashes the system can recover
- Specially, useful in distributed systems when you cannot predict what type of message you will receive





- Akka implements supervision mechanisms to react to failures
- Children may inform their parents when they terminate or fail
- Actors may use the function watch (ActorRef<T> actor) to supervise their children
- If an actor is being supervised by a (parent) watcher, it sends to the watcher
  - A ChildFailed signal, if it crashed due to an exception
  - A Terminated signal, if it terminates normally

But ChilFailed extends from Terminated



- A signal can be seen as a message that is automatically sent by Akka
- For a watcher to handle signals, it must have a message handler with onsignal (Signal.class, Function f)
- The message send by the signal contains a reference to the sender actor. It can be accessed with getRef()

```
/* --- Message handling -----
@Override
public Receive<T> createReceive() {
    return newReceiveBuilder()
     .onMessage(Message.class, this::onMessage)
     // Here order matters `ChildFailed extends Terminated`
     .onSignal(ChildFailed.class, this::onChildFailed)
     .onSignal(Terminated.class, this::onTerminated)
     .build();
```

Note I: When processing a message/signal, the message handler picks the handler that first matches the class of the message

Note II: Since ChildFailed extends Terminated, if onSignal(Terminated,...) appears before onSignal(ChildFailed,...), when a ChildFailed signal arrives, the latter on Signal will not be triggered

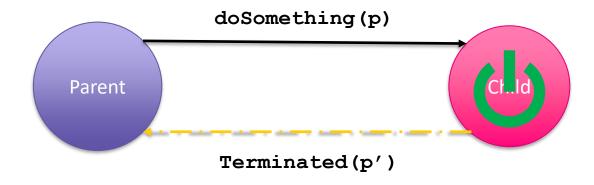
# Actor supervision (graphically)





# Actor supervision (graphically)

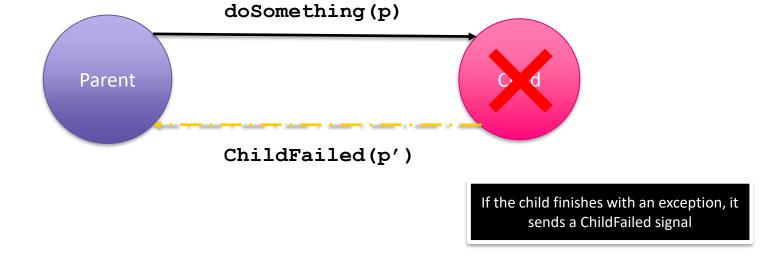


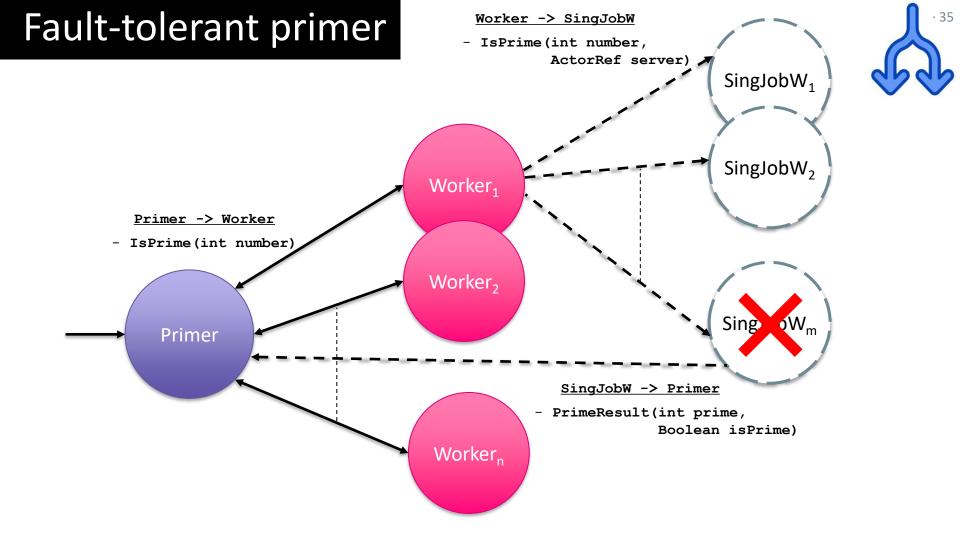


The dashed-dotted yellow arrow indicates the sending of a signal. These are sent automatically by Akka as part of the supervision functionality. If the child finishes normally, it sends a Terminated signal.

# Actor supervision (graphically)

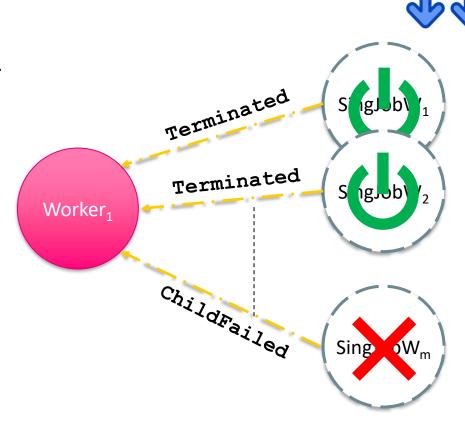






# Fault-tolerant primer

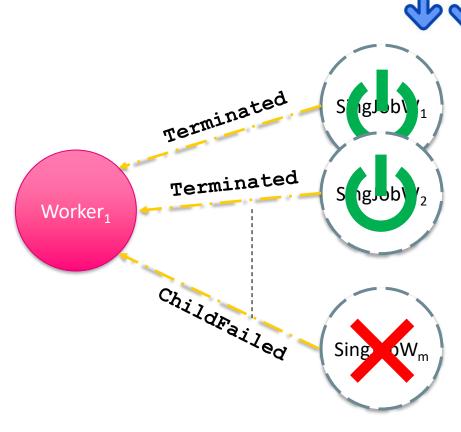
- We extend the primer to handle the case when a single job worker fails
- To this end, the worker needs to:
- 1. Watch all the actors it spawns
- 2. Handle ChildFailed signals
  - The handler spawns a new worker and sends the number again to check whether it is prime
- 3. Handle Terminated signals
  - No more computation needed, we can mark the number as checked



# Fault-tolerant primer

- We extend the primer to handle the case when a single job worker fails
- To this end, the worker needs to:
- Watch all the actors it spawns
- 2. Handle ChildFailed signals
  - The handler spawns a new worker and sends the number again to check whether it is prime
- 3. Han e Terminated signals
  - more computation needed, we
     mark the number as checked

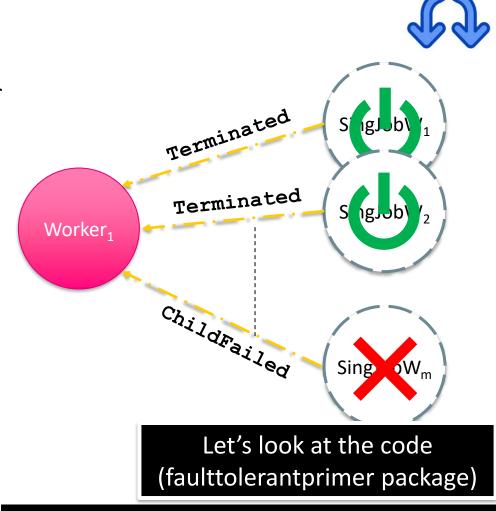
We must extend the actor's state to keep track of what number is computed by what actor



# Fault-tolerant primer

- We extend the primer to handle the case when a single job worker fails
- To this end, the worker needs to:
- 1. Watch all the actors it spawns
- 2. Handle ChildFailed signals
  - The handler spawns a new worker and sends the number again to check whether it is prime
- 3. Han e Terminated signals
  - nore computation needed, we mark the number as checked

We must extend the actor's state to keep track of what number is computed by what actor



IT UNIVERSITY OF COPENHAGEN



# Adaptive load balancing



- <u>Load balancing</u> refers to the process of distributing a set of tasks over a set of resources (computing units), with the aim of making their overall processing more efficient. [Wikipedia]
- In the (static) primer system, we indiscriminately spawned processes to perform tasks
  - This may cause sending tasks to busy workers while other idle workers could be processing them
- There exists some patterns that aim at distributing computation fairly among actors.
  - For instance, the scatter-gather pattern



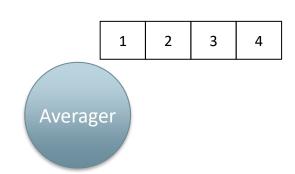
- Scatter-Gather is a common design patter in distributed systems that can be easily implemented with actors
- Typically, the level of scattering (i.e., number of spawned actors) depends on the size of the problem to solve (dynamic load balancing)
  - But it can also be limited by other factors, e.g., CPU or memory usage
- A scatter-gather systems contains two main type of actors
  - <u>Scatterer</u>: if possible, it splits computation in smaller units. Otherwise, it may perform a processing step in the atomic piece of data and send it to a gatherer
  - <u>Gatherer</u>: Receives pieces of data from scatterers or gatherers and combines them into a single piece of data. Then, if it is not the highestlevel gatherer, it sends the result to a gatherer

### Average and scatter-gather



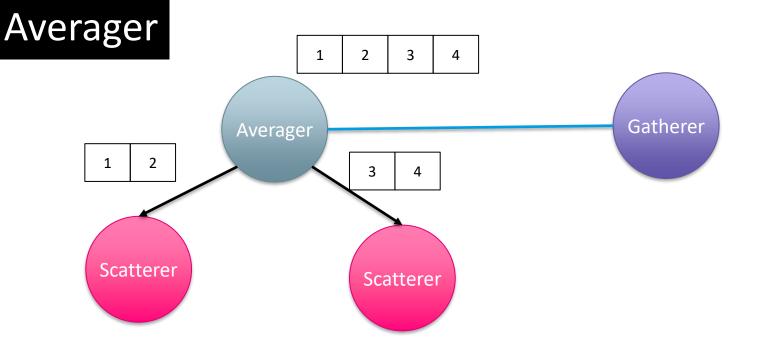
- A problem for which this pattern is suitable is computing the average of a list of numbers
- Given a set of natural numbers  $a_1$ ,  $a_2$ , ...,  $a_n$ , the average is  $\frac{1}{n}\sum_i a_i$ 
  - Note that this is equivalent to  $\sum_{i} \frac{a_i}{n}$
- In a nutshell, we can have scatterer actors splitting computation and computing each factor  $\frac{a_i}{n}$ , and gatherers summing up the results

Averager





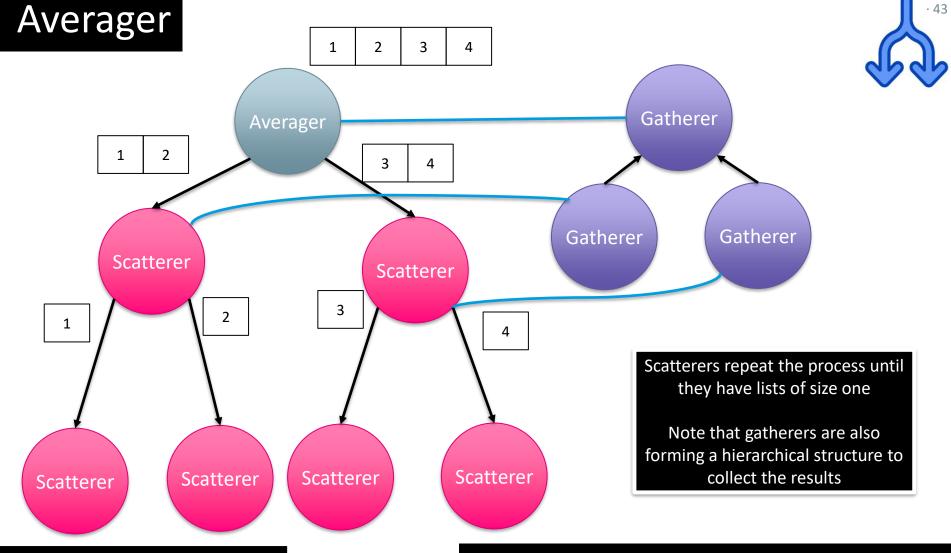
Consider a system that computes the average of a list of numbers



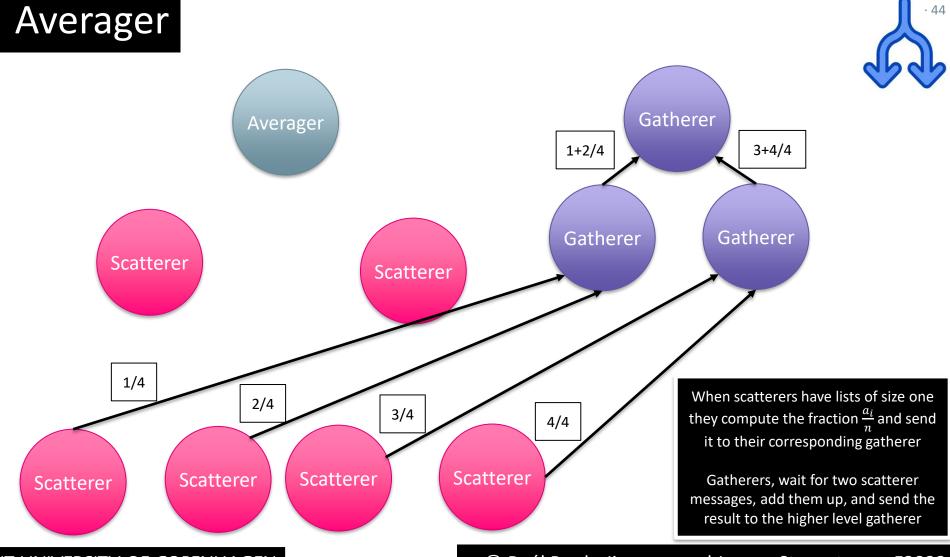


In the first step, we split the computation into two sublists, and assign them to separate scatterer workers

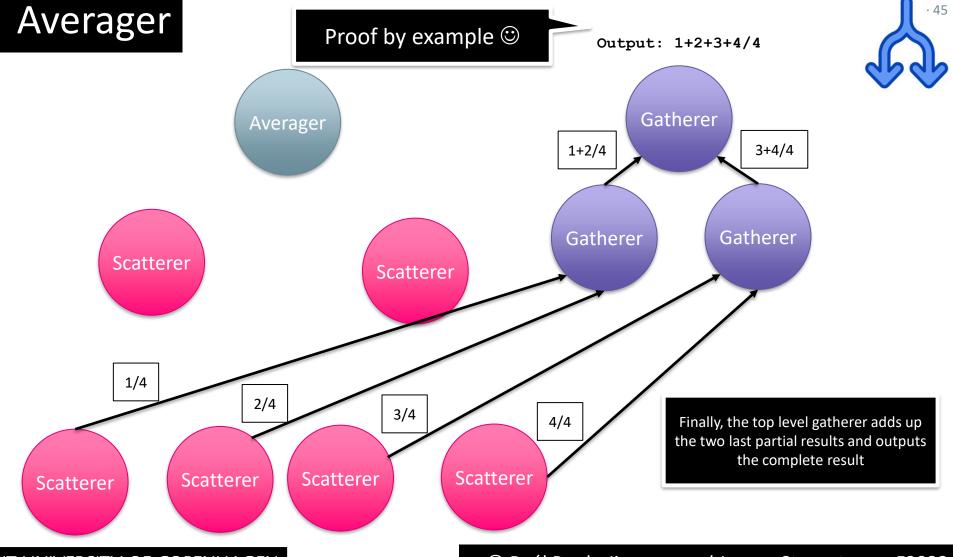
Also, we spawn a gatherer worker that will receive merge the average of each sublist



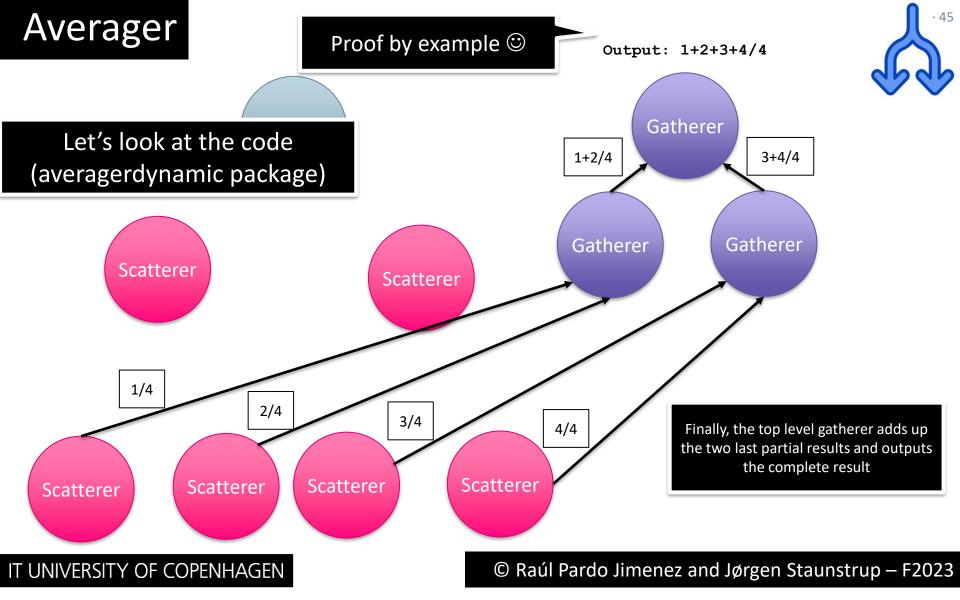
IT UNIVERSITY OF COPENHAGEN

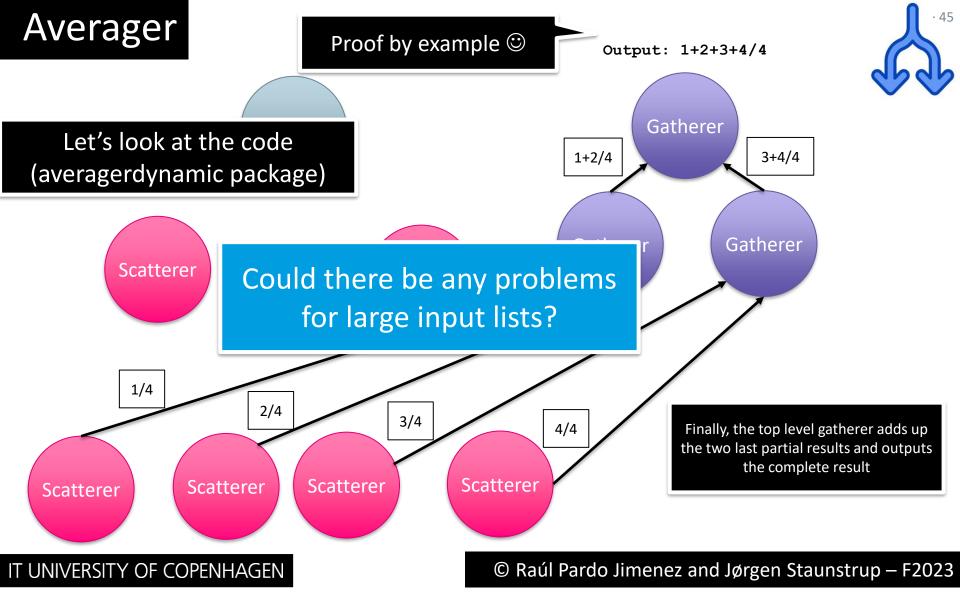


IT UNIVERSITY OF COPENHAGEN



IT UNIVERSITY OF COPENHAGEN







- The size of the problem does not necessarily need to determine the distribution of computation
- One may have HW restrictions
  - As we saw, actors systems running in a single machine may not scale well beyond the number of processors
- Another example of adaptive load balancing are elastic systems
  - Elastic systems try to keep a number of active actors proportional to the workload
  - Several exercises for this week target implementing an elastic server

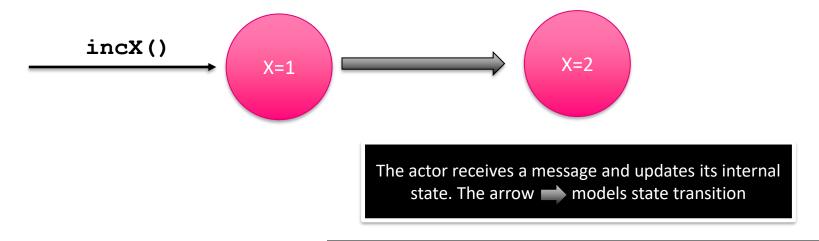


# Changing behaviour

# Actors with changing behaviour



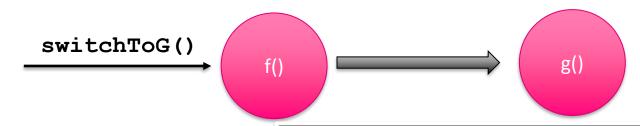
- The actors model states that an actor can "Change its behaviour (local state and/or message handlers)"
- So far we have considered change in behaviour as change in state



#### Actors with changing behaviour



- The actors model states that an actor can "Change its behaviour (local state and/or message handlers)"
- However, actors may also change the functions to process messages (i.e., message handlers)



In this example, the actor that was executing f() when a message comes, now executes g(). The new function g() could be completely different, e.g., changing how messages are processed or even waiting for different type of messages!



- In Akka, we can change the behaviour of an actor by defining a function that returns the new behaviour
  - Like in the behaviour defined in createReceive()
- In fact, we have already done this when we return Behaviors.stopped() to terminate an actor



- The packages averagerdynamic and averagerbehavior implement the same system, but the latter uses changing behaviour
- In averagerdynamic, we use a Boolean variable (receivedFirstNumber) to determine whether we have received the first or second GathererCommand message
- In averagebehavior, we define a new behaviour (waitForsecond) to which the actor transitions after receiving the first GathereCommand message
  - In this way we can do without the Boolean variable mentioned above

#### Averager with changing behaviour



- The packages averagerdynamic and averagerbehavior implement the same system, but the latter uses changing behaviour
- In averagerdynamic, we use a Boolean variable (receivedFirstNumber) to determine whether we have received the first or second GathererCommand message
- In averagebehavior, we define a new behaviour (waitForsecond) to which the actor transitions after receiving the first GathereCommand message
  - In this way we can do without the Boolean variable mentioned above

Let's look at the code (averagerbehaviors package)

#### Actors in distributed systems



#### The actors model has natural mapping in distributed systems

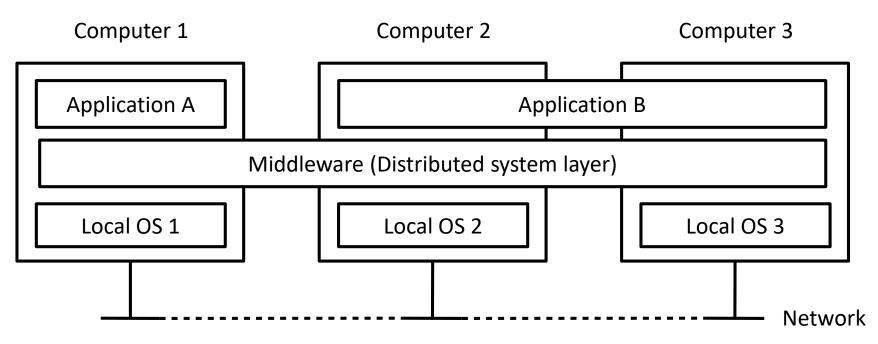
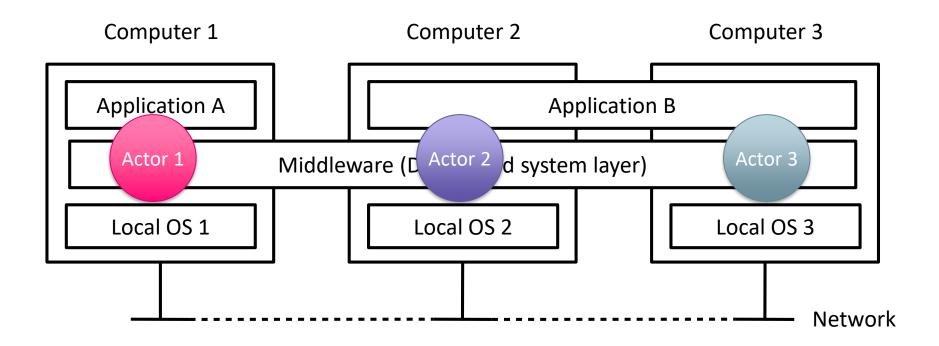


Figure taken from -> Distributed Systems: Principles and Paradigms. Andrew S. Tanenbaum and Maarten Van Steen. 2007.

#### Actors in distributed systems



The actors model has natural mapping in distributed systems

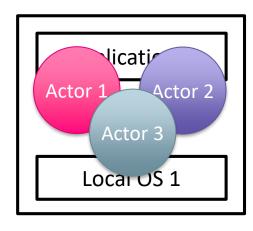


#### Actors in a single computer



The actors model is applicable in a single computer as well

Computer 1



In this course, we focus on this type of actor system

### Agenda



- Actors model (revisited)
  - Primer
- Dynamic topology
- Fault-tolerance
  - Supervision
- Adaptive load balancing
  - Scatter-Gather
- Changing behaviour



# BONUS: Revisiting Bounded Buffer

#### Producer-consumer problem | Intuition



Perhaps more intuitive example

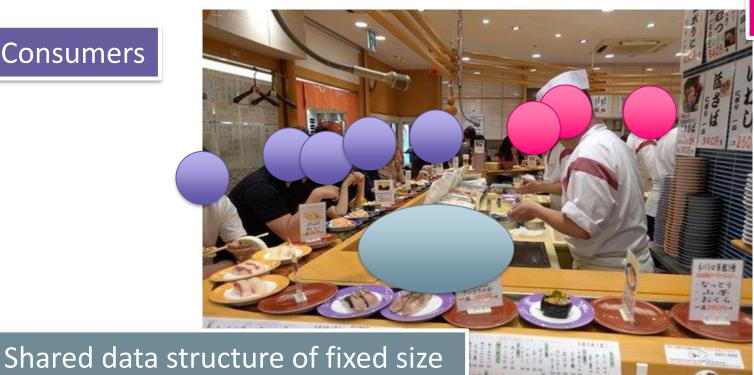
Consumers Shared data structure of fixed size

**Producers** 

#### Producer-consumer problem | Intuition



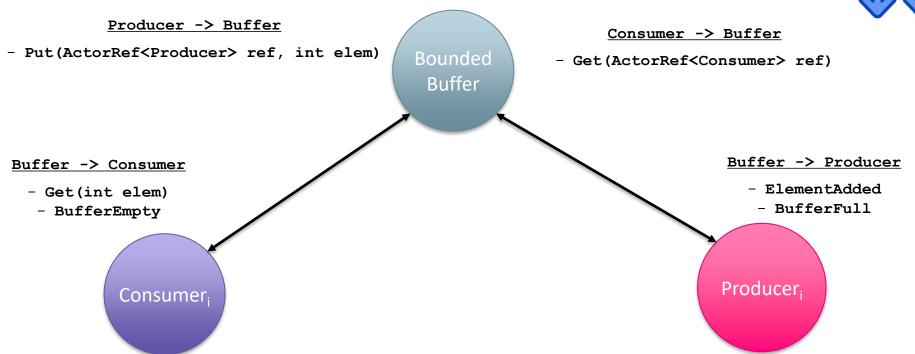
Consumers



**Producers** 

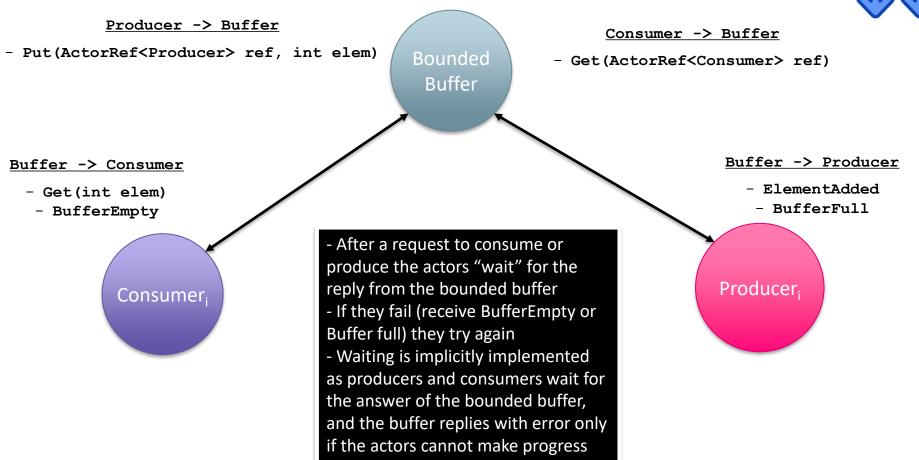
## **Bounded Buffer with Actors**





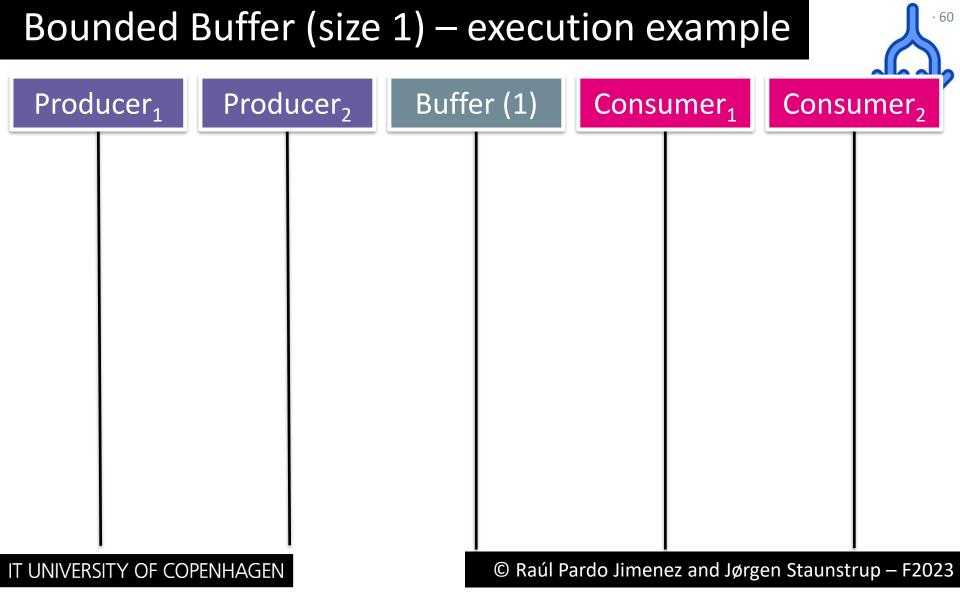
## **Bounded Buffer with Actors**

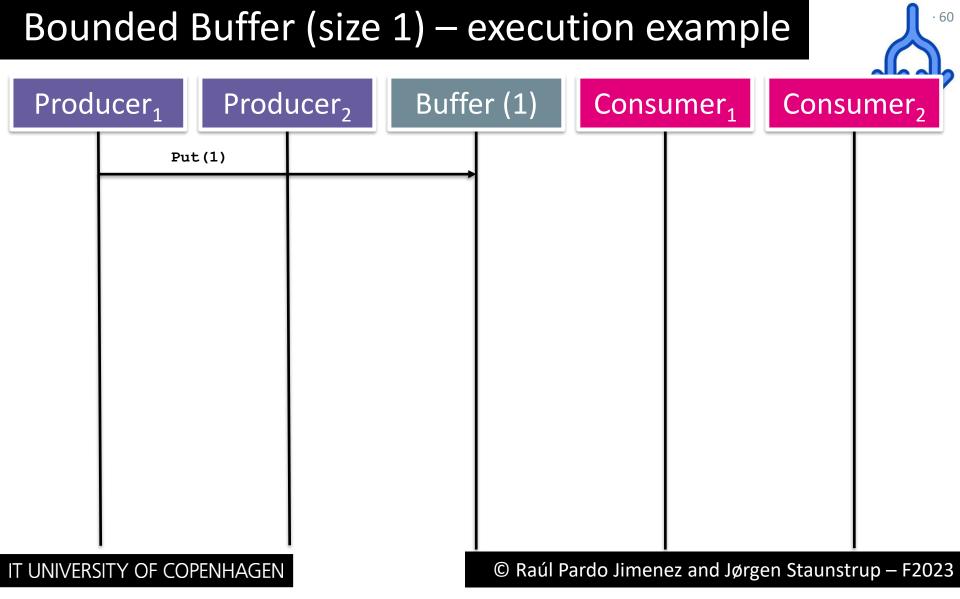


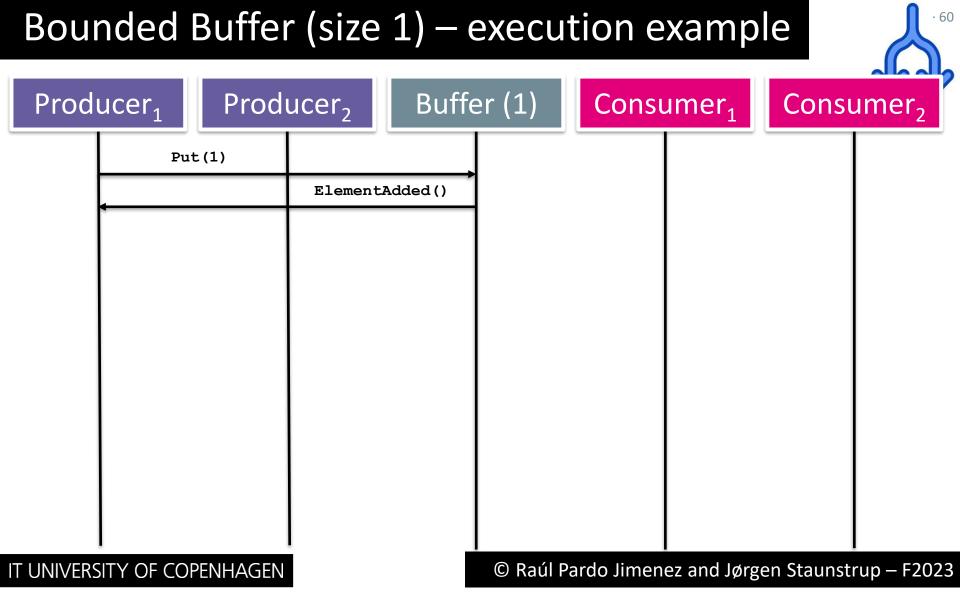


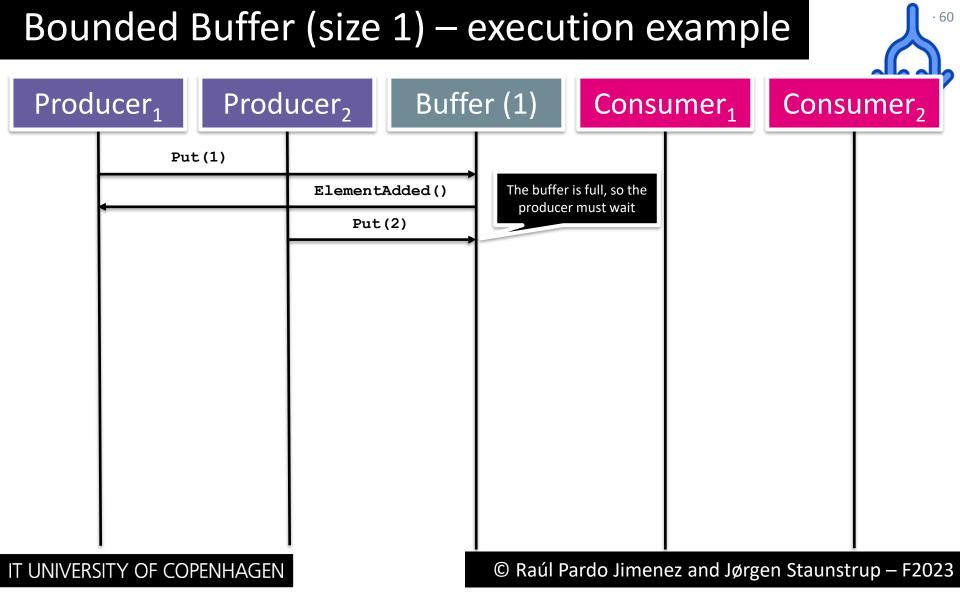
IT UNIVERSITY OF COPENHAGEN

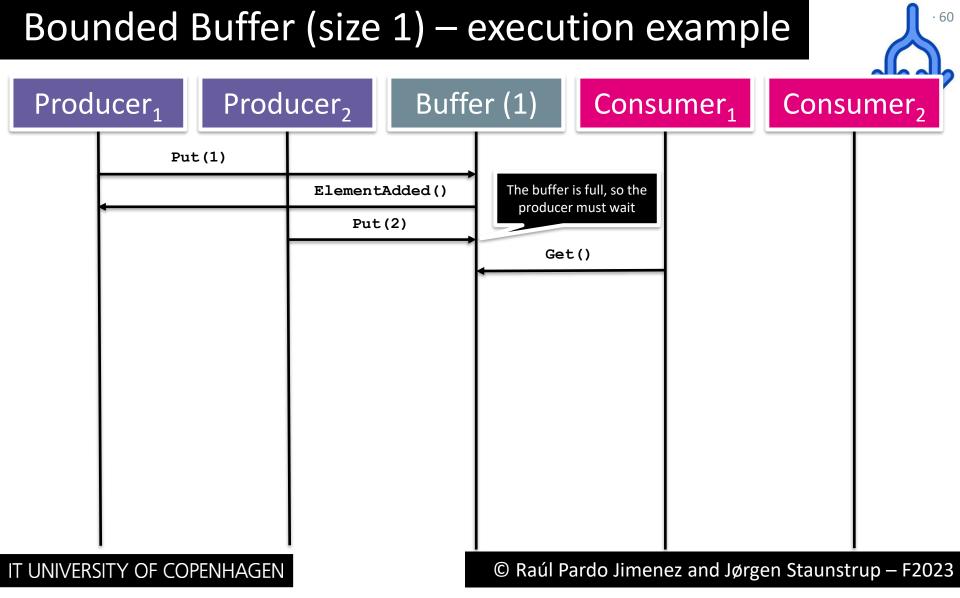
© Raúl Pardo Jimenez and Jørgen Staunstrup – F2023

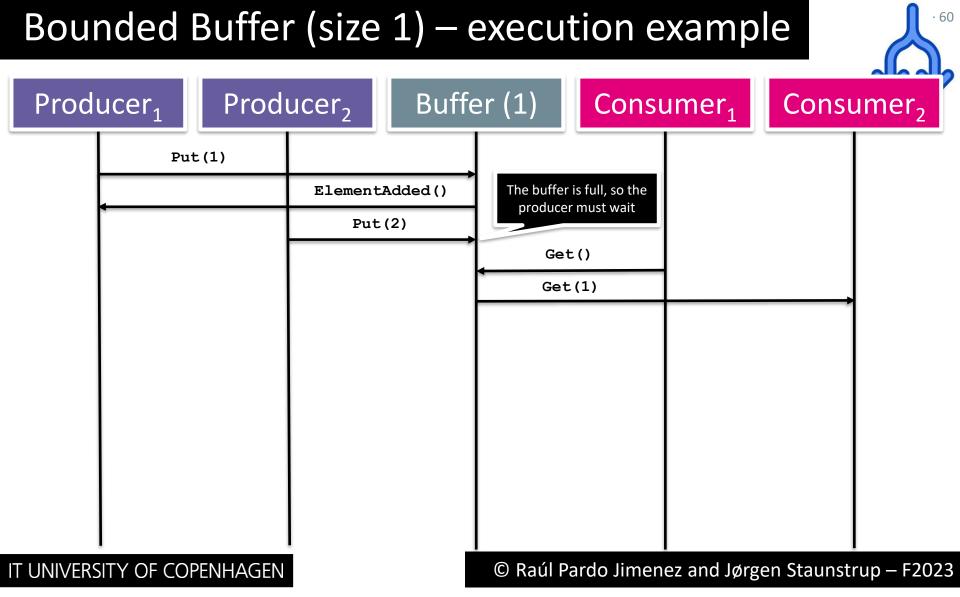


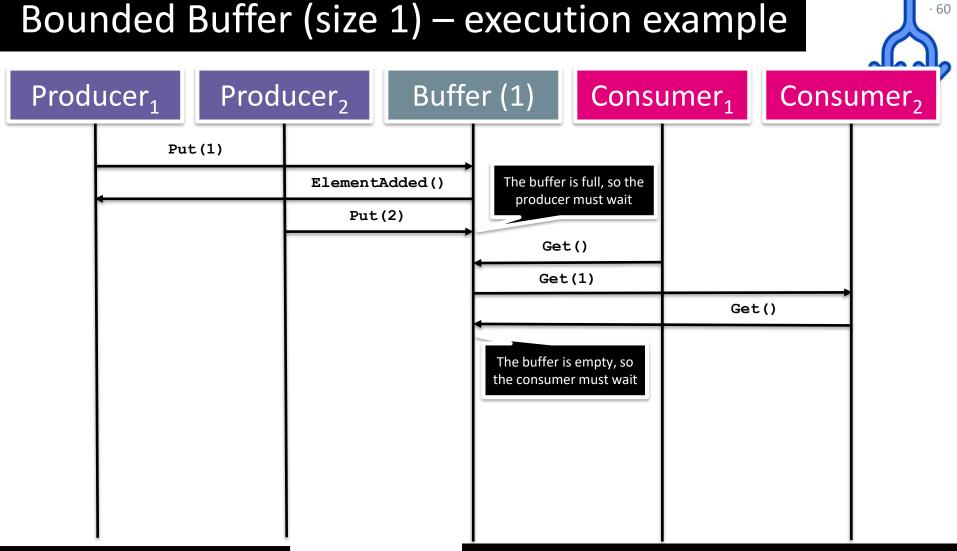




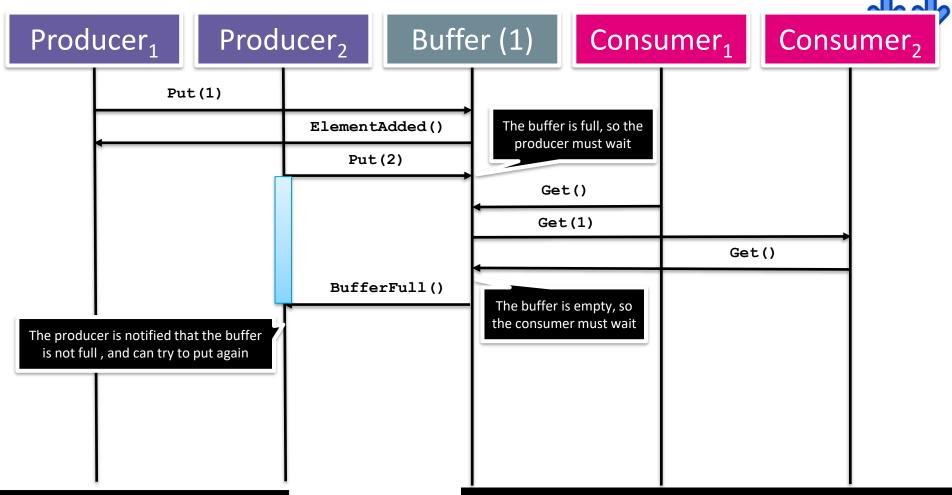










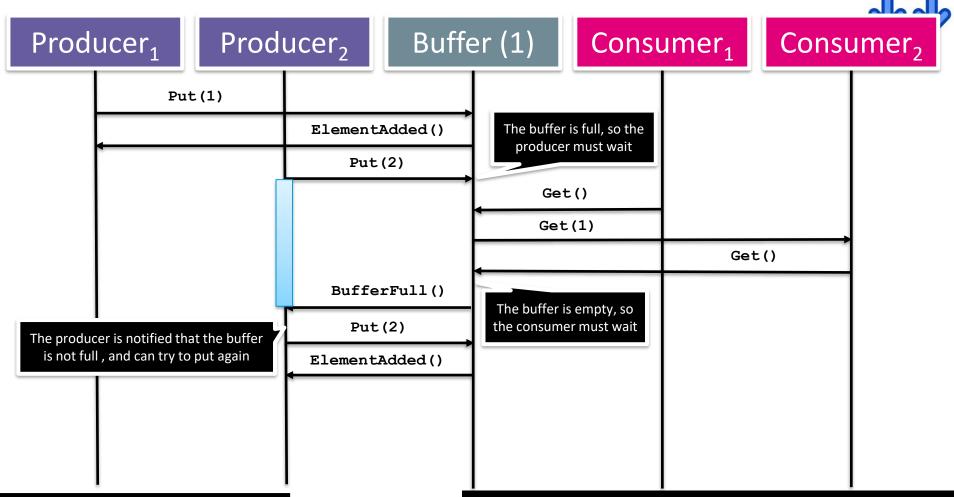


IT UNIVERSITY OF COPENHAGEN

© Raúl Pardo Jimenez and Jørgen Staunstrup – F2023

. 60

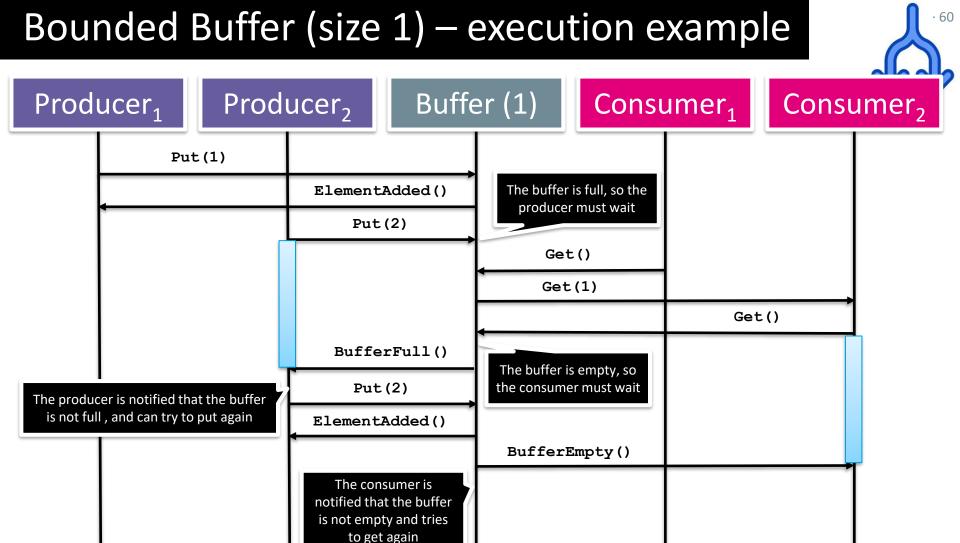
## Bounded Buffer (size 1) – execution example

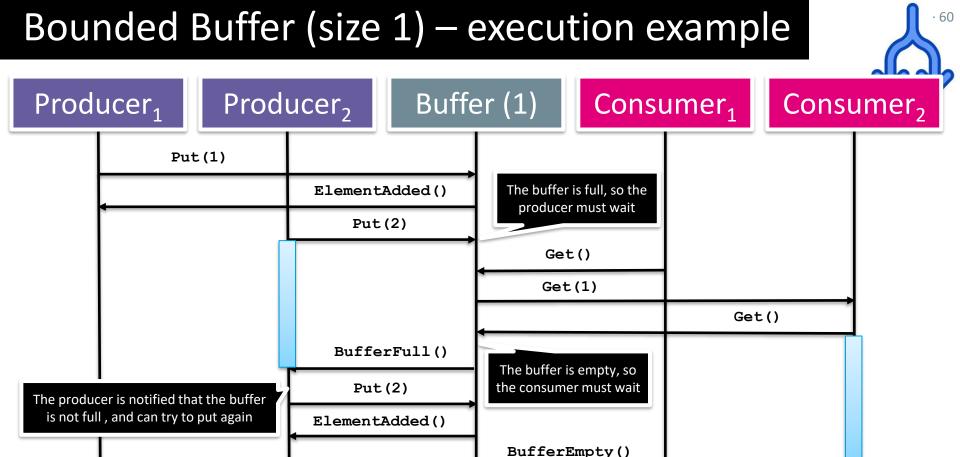


IT UNIVERSITY OF COPENHAGEN

© Raúl Pardo Jimenez and Jørgen Staunstrup – F2023

. 60





Get (2)

The consumer is

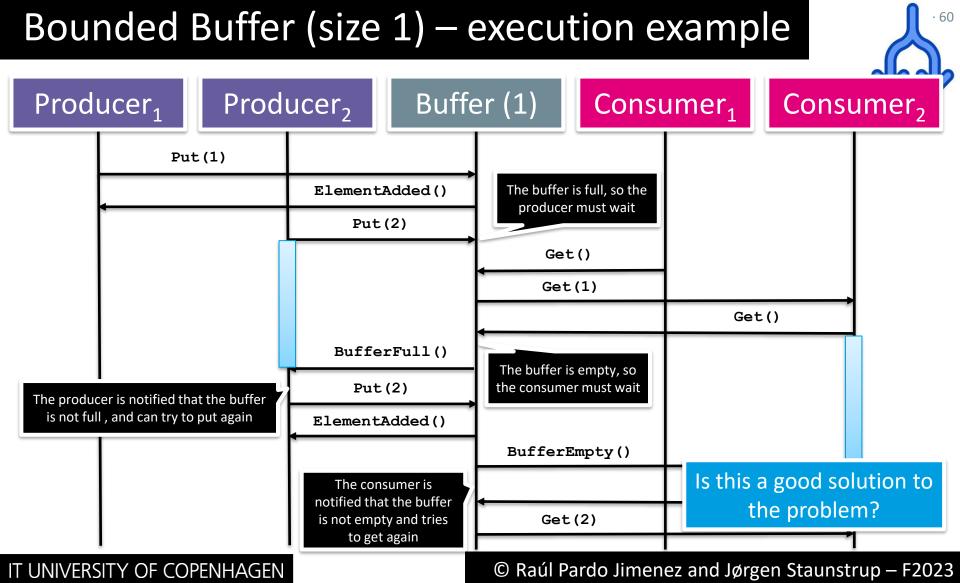
notified that the buffer is not empty and tries

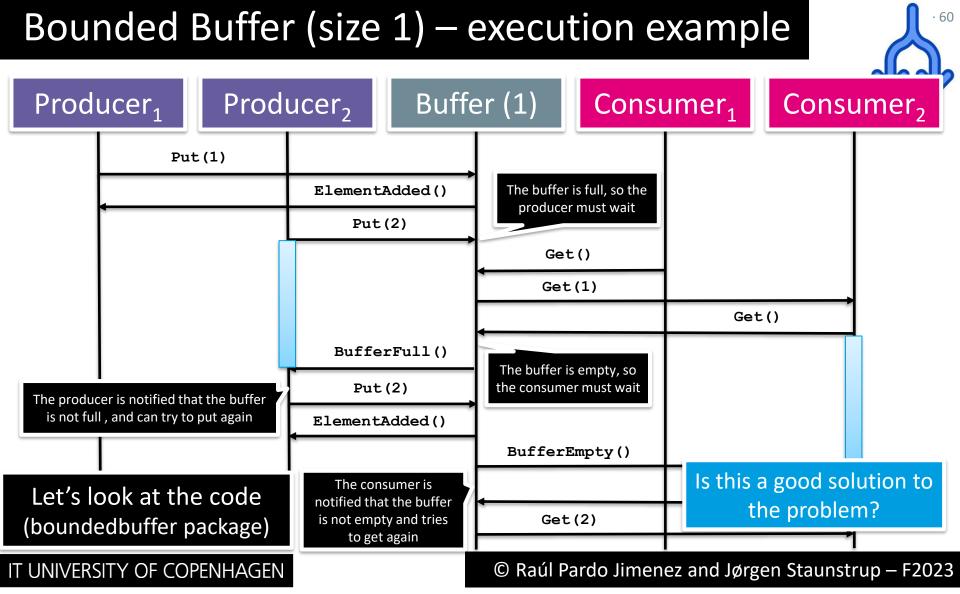
to get again

IT UNIVERSITY OF COPENHAGEN

© Raúl Pardo Jimenez and Jørgen Staunstrup – F2023

Get()



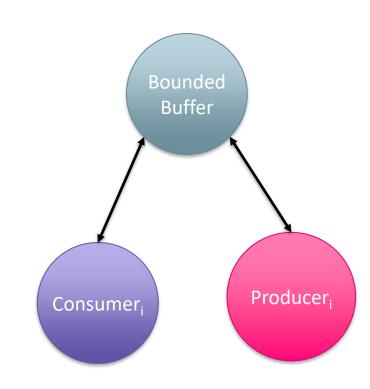


61

Assuming FIFO mailboxes (Akka's default)

- Consider this execution
- Producer1 sends put(1) to the buffer
- Consumer1 sends get() to the buffer
- 3. ...

Is it guaranteed that Consumer 1 will get the produced element?

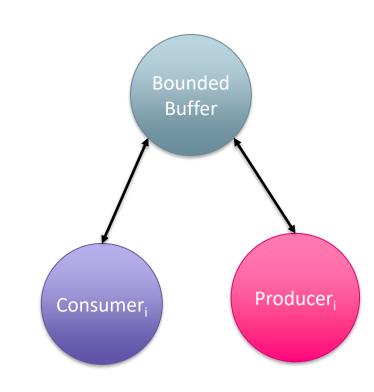




Assuming FIFO mailboxes (Akka's default)

- Consider this execution
- 1. Producer1 sends put(1) to the buffer
- Producer2 sends put(2) to the buffer
- 3. Consumer1 sends get() to the buffer
- 4. ...

Is it guaranteed that Consumer1 will get either 1 or 2?



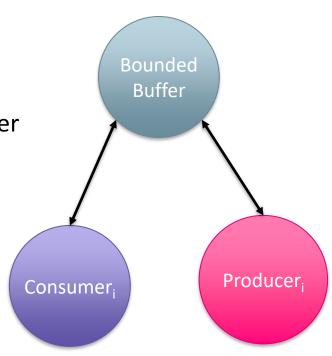
## Bounded Buffer (interesting) executions



Assuming FIFO mailboxes (Akka's default)

- Consider this execution
- 1. Producer1 sends put(1) to the buffer
- Producer2 sends put(2) to the buffer
- 3. Producer1 receives ElementAdded() from the buffer
- Consumer1 sends get() to the buffer
- 5. ..

Is it guaranteed that Consumer 1 will get the produced element?



IT UNIVERSITY OF COPENHAGEN

© Raúl Pardo Jimenez and Jørgen Staunstrup – F2023

- 64
- Note that in the previous questions the behaviour of the systems depends on the reception of messages
- Thus, the happened-before relation defined by Lamport is useful in reasoning about actor systems
  - An action a happens-before an action b
    if they belong to the same actor and
    a was executed before b
  - A send(m) action happens-before its corresponding receive(m)
- Note the similarity with the happens-before relation of the Java memory model
  - We reason about message exchange instead of locking (inherent coordination problems remain, i.e., "semantic" deadlock & starvation)
  - Visibility issues disappear as actors only access local memory